

Chapter 1

Introduction

A note on the use of these PowerPoint slides:

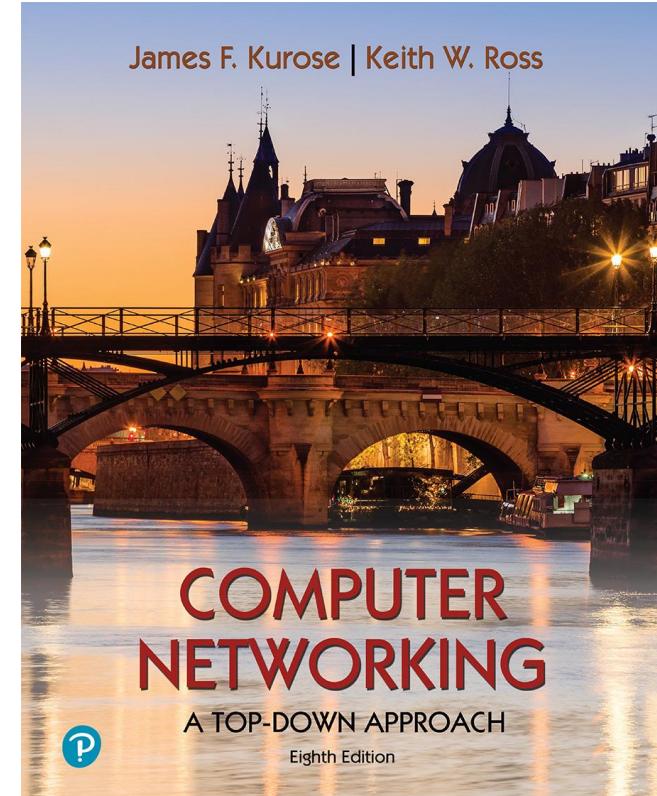
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2020
J.F Kurose and K.W. Ross, All Rights Reserved



*Computer Networking: A
Top-Down Approach*
8th edition
Jim Kurose, Keith Ross
Pearson, 2020

Chapter 1: introduction

Chapter goal:

- Get “feel,” “big picture,” introduction to terminology
 - more depth, detail *later* in course



Overview/roadmap:

- What *is* the Internet? What *is* a protocol?
- **Network edge:** hosts, access network, physical media
- **Network core:** packet/circuit switching, internet structure
- **Performance:** loss, delay, throughput
- Protocol layers, service models
- Security
- History

The Internet: a “nuts and bolts” view



Billions of connected computing *devices*:

- *hosts* = end systems
- running *network apps* at Internet's “edge”

Packet switches: forward packets (chunks of data)

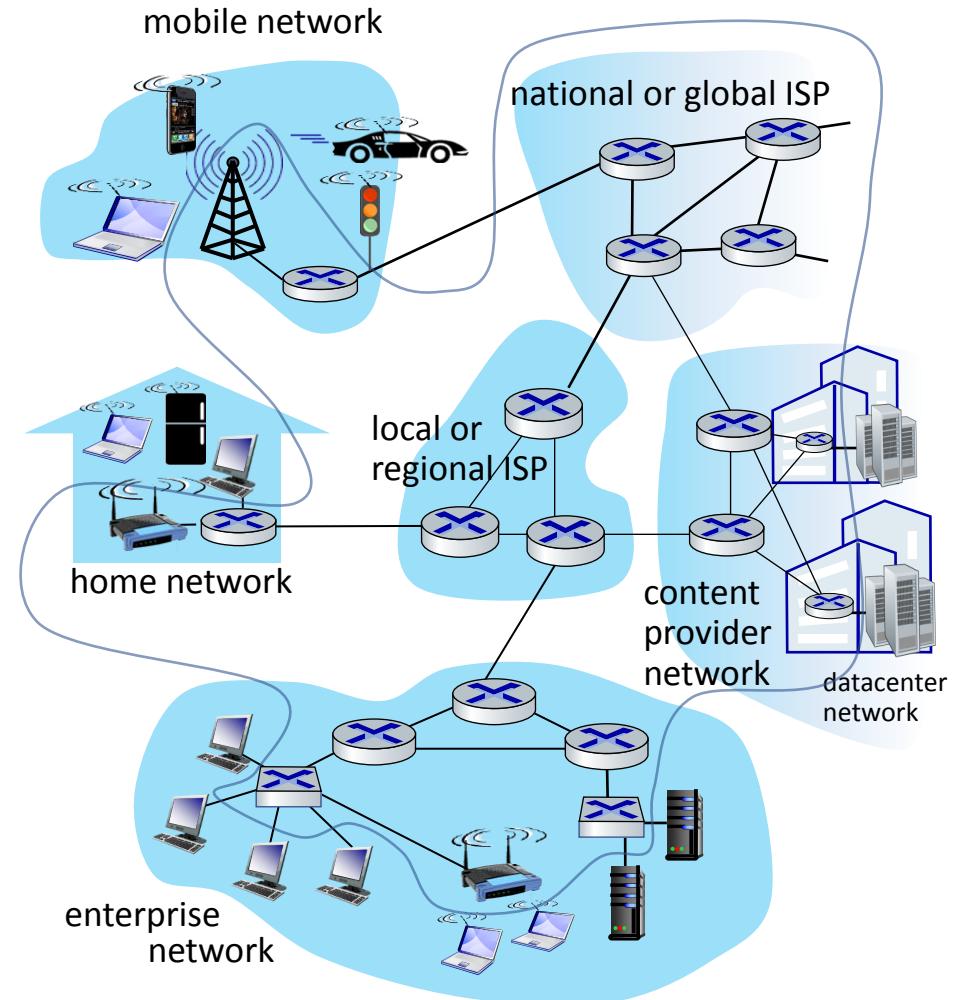
- routers, switches

Communication links

- fiber, copper, radio, satellite
- transmission rate: *bandwidth*

Networks

- collection of devices, routers, links: managed by an organization



“Fun” Internet-connected devices



Amazon Echo



Internet refrigerator



Security Camera



Internet phones



IP picture frame



Slingbox: remote control cable TV



Gaming devices



Pacemaker & Monitor



Web-enabled toaster + weather forecaster



sensorized, bed mattress



Fitbit



Tweet-a-watt:
monitor energy use

bikes



cars

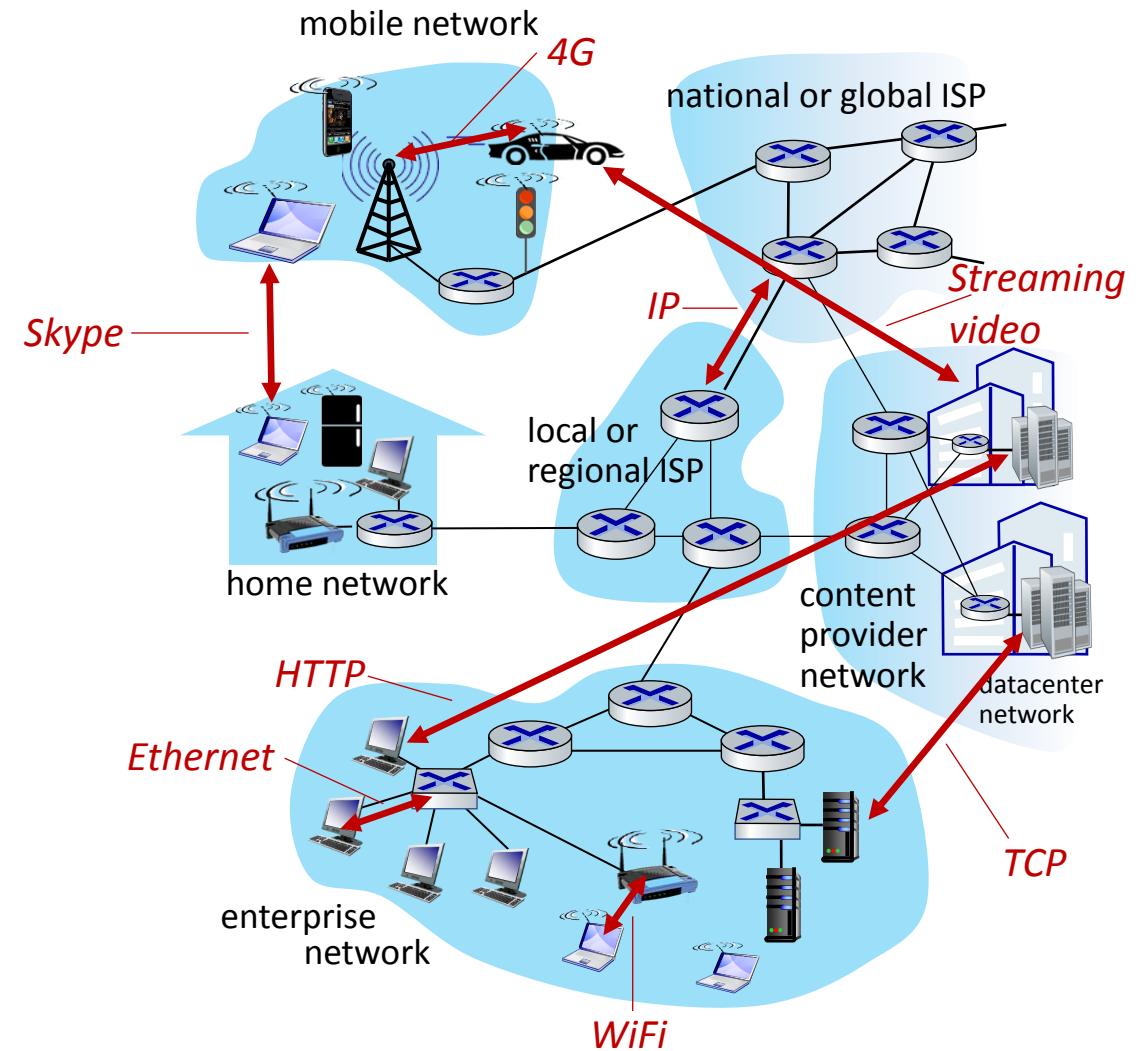


scooters

Others?

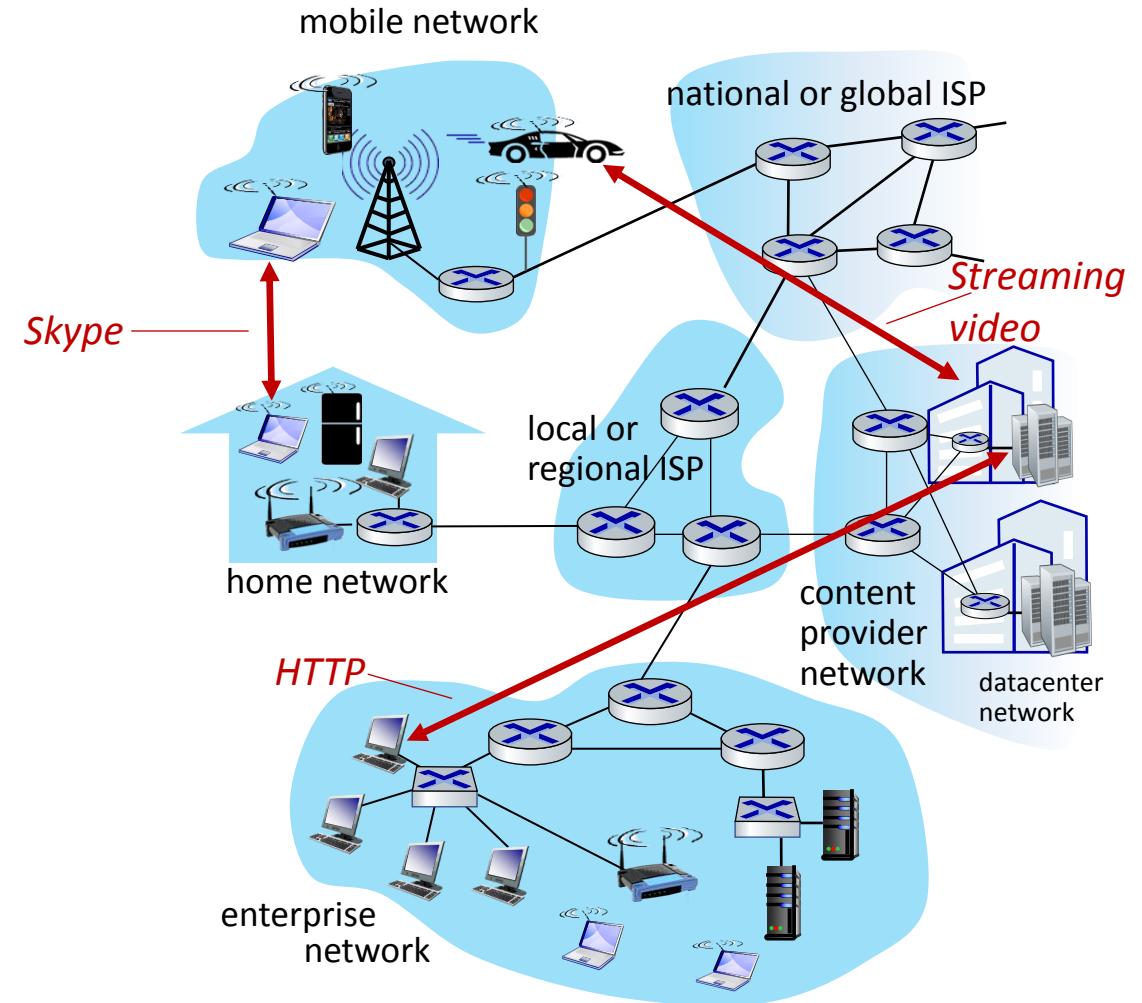
The Internet: a “nuts and bolts” view

- *Internet: “network of networks”*
 - Interconnected ISPs
- *protocols are everywhere*
 - control sending, receiving of messages
 - e.g., HTTP (Web), streaming video, Skype, TCP, IP, WiFi, 4G, Ethernet
- *Internet standards*
 - RFC: Request for Comments
 - IETF: Internet Engineering Task Force



The Internet: a “services” view

- *Infrastructure* that provides services to applications:
 - Web, streaming video, multimedia teleconferencing, email, games, e-commerce, social media, interconnected appliances, ...
- provides *programming interface* to distributed applications:
 - “hooks” allowing sending/receiving apps to “connect” to, use Internet transport service
 - provides service options, analogous to postal service



What's a protocol?

Human protocols:

- “what’s the time?”
- “I have a question”
- introductions

Rules for:

- ... specific messages sent
- ... specific actions taken
when message received,
or other events

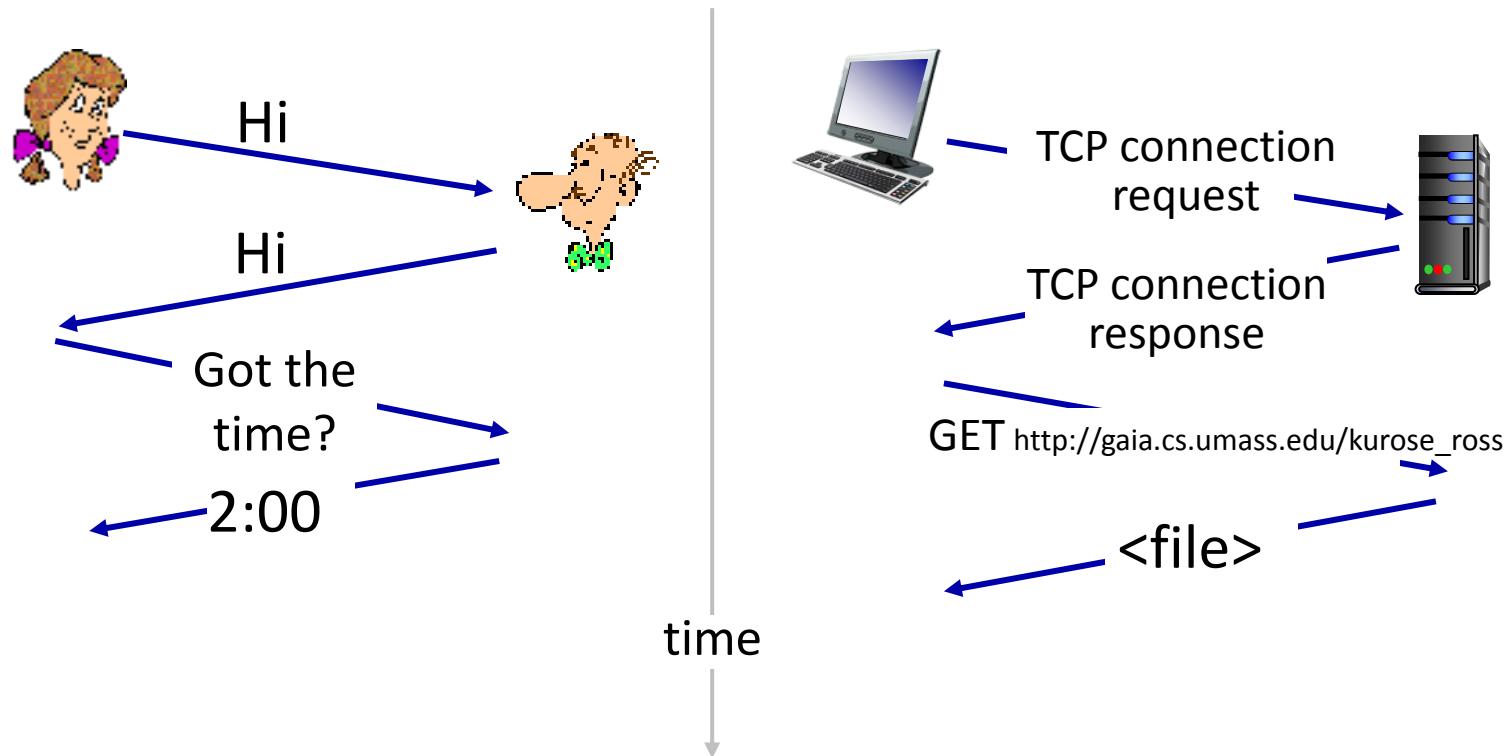
Network protocols:

- computers (devices) rather than humans
- all communication activity in Internet governed by protocols

*Protocols define the **format, order** of messages sent and received among network entities, and **actions taken** on message transmission, receipt*

What's a protocol?

A human protocol and a computer network protocol:



Q: other human protocols?

Chapter 1: roadmap

- What *is* the Internet?
- What *is* a protocol?
- **Network edge:** hosts, access network, physical media
- Network core: packet/circuit switching, internet structure
- Performance: loss, delay, throughput
- Security
- Protocol layers, service models
- History

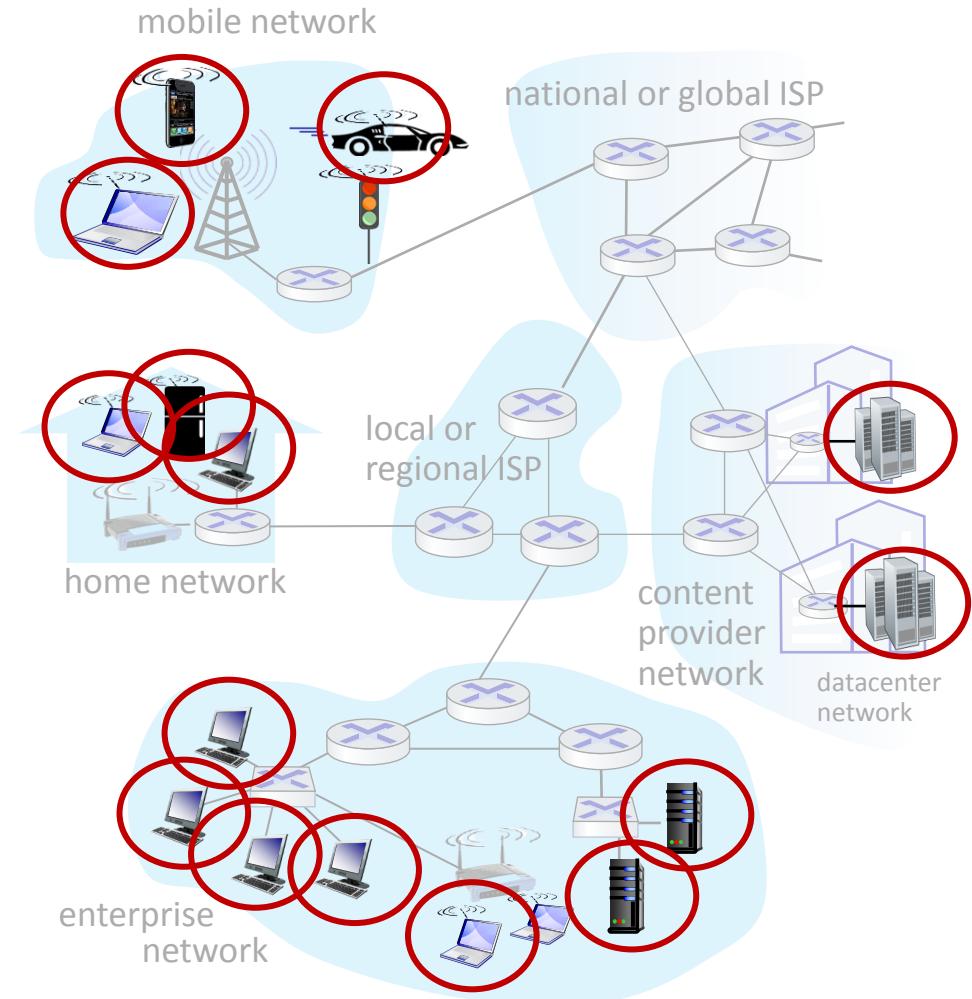


A closer look at Internet structure

Network edge:

End-systems/hosts:

- clients and servers
- servers often in data centers



A closer look at Internet structure

Network edge:

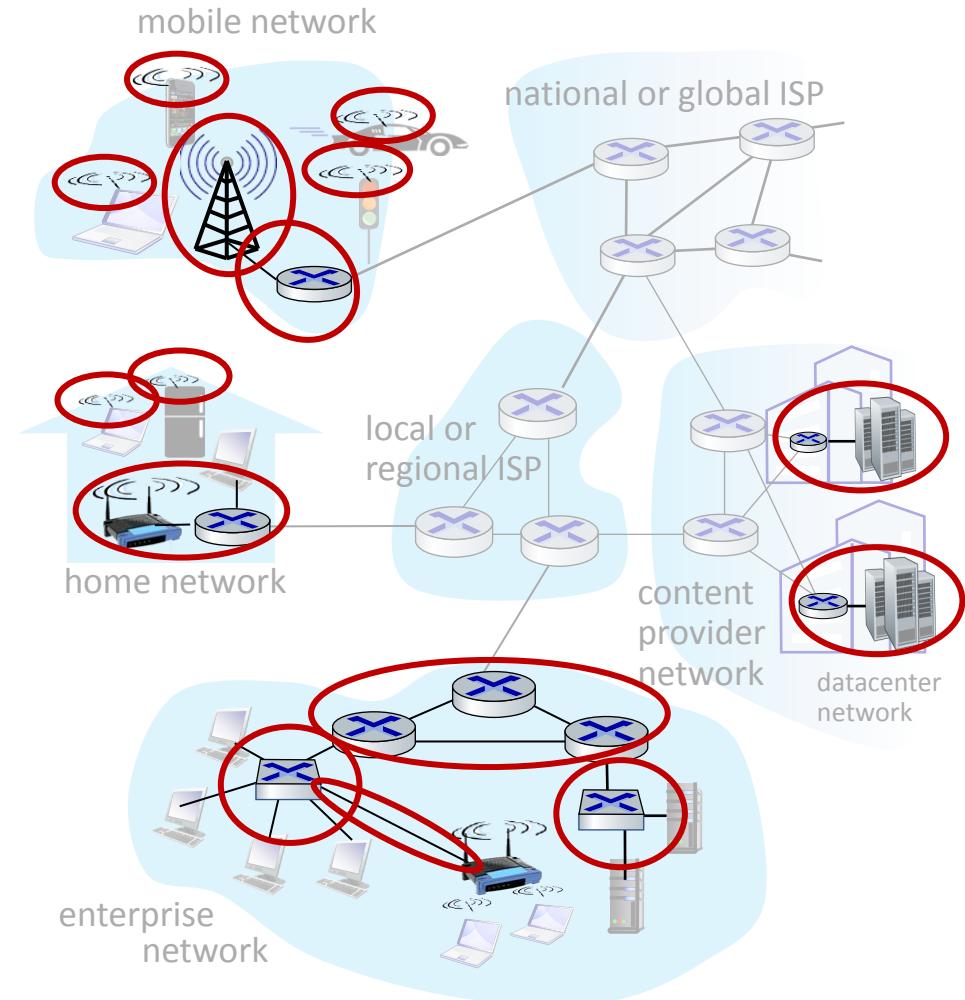
End-systems/hosts

- hosts: clients and servers
- servers often in data centers

Access networks: connects hosts to the “edge router”

wired, wireless communication links

Switches



A closer look at Internet structure

Network edge:

End-systems/hosts

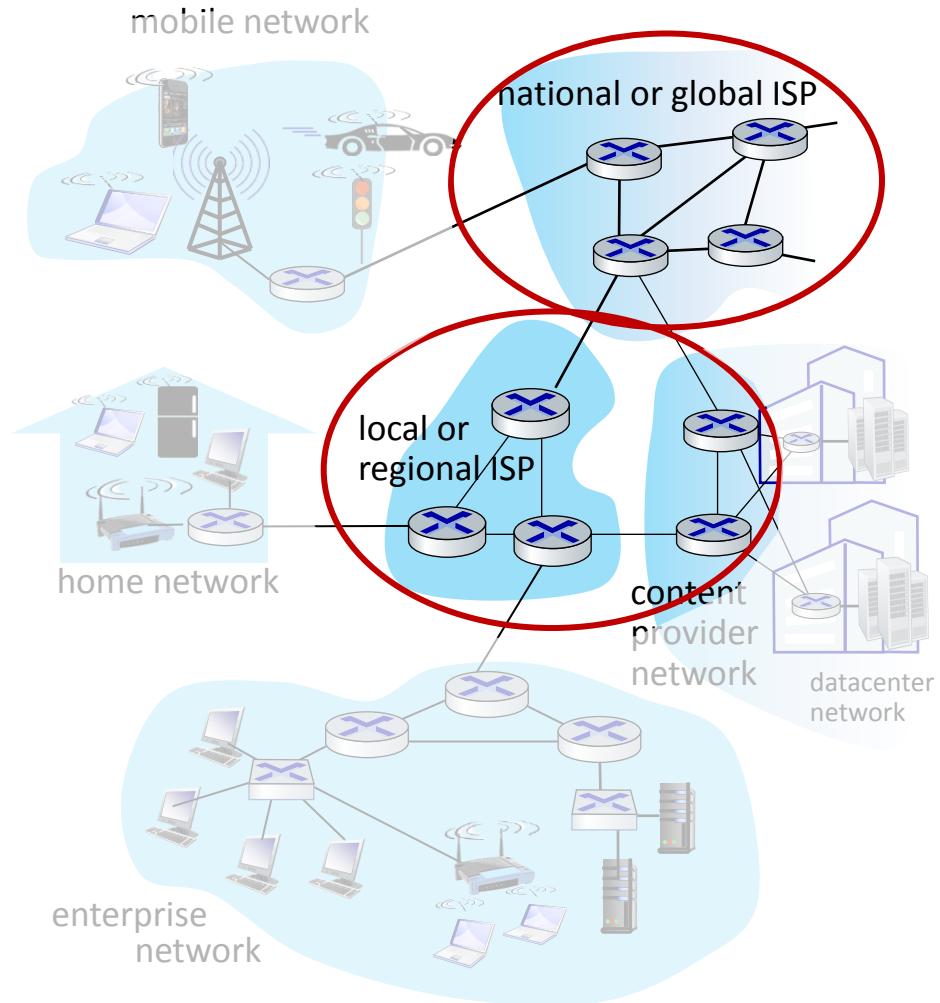
- hosts: clients and servers
- servers often in data centers

Access networks, physical media:

- wired, wireless communication links

Network core:

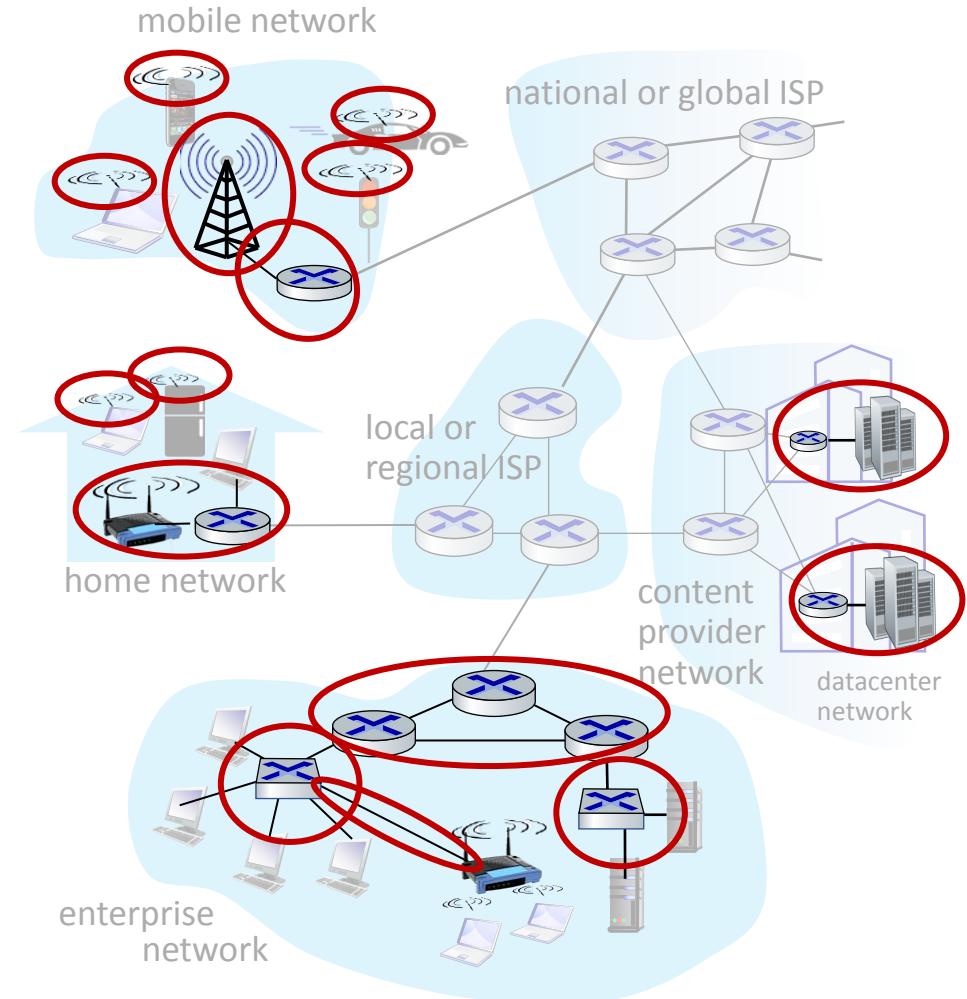
- interconnected routers
- network of networks



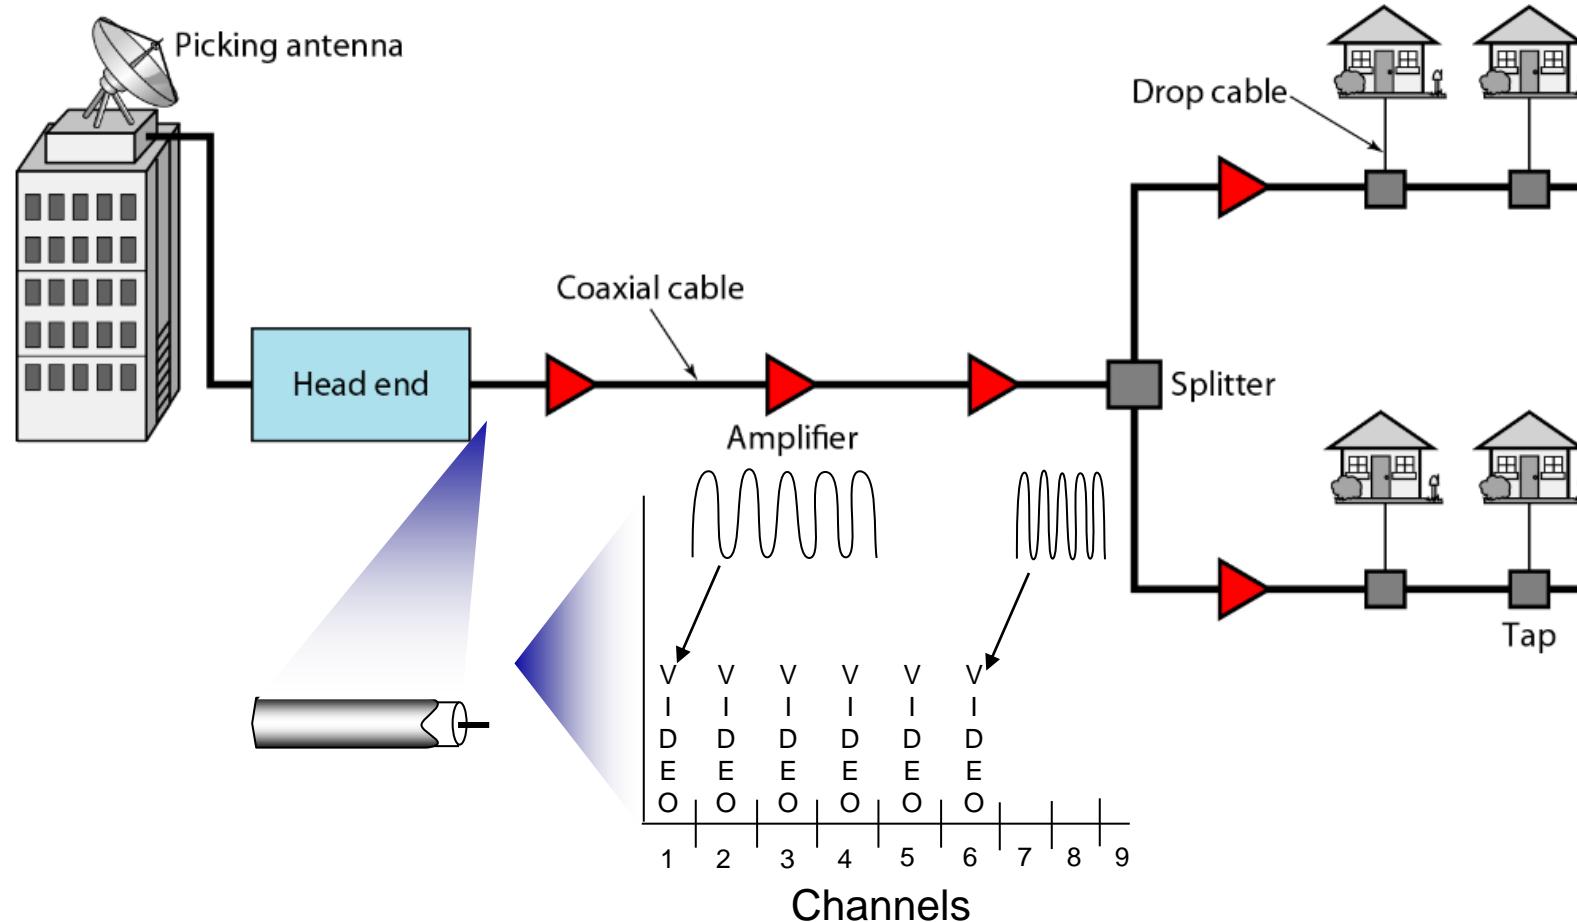
Access networks and physical media

*Q: How to connect end systems
to edge router?*

- residential access nets
- institutional access networks (school, company)
- mobile access networks (WiFi, 4G/5G)



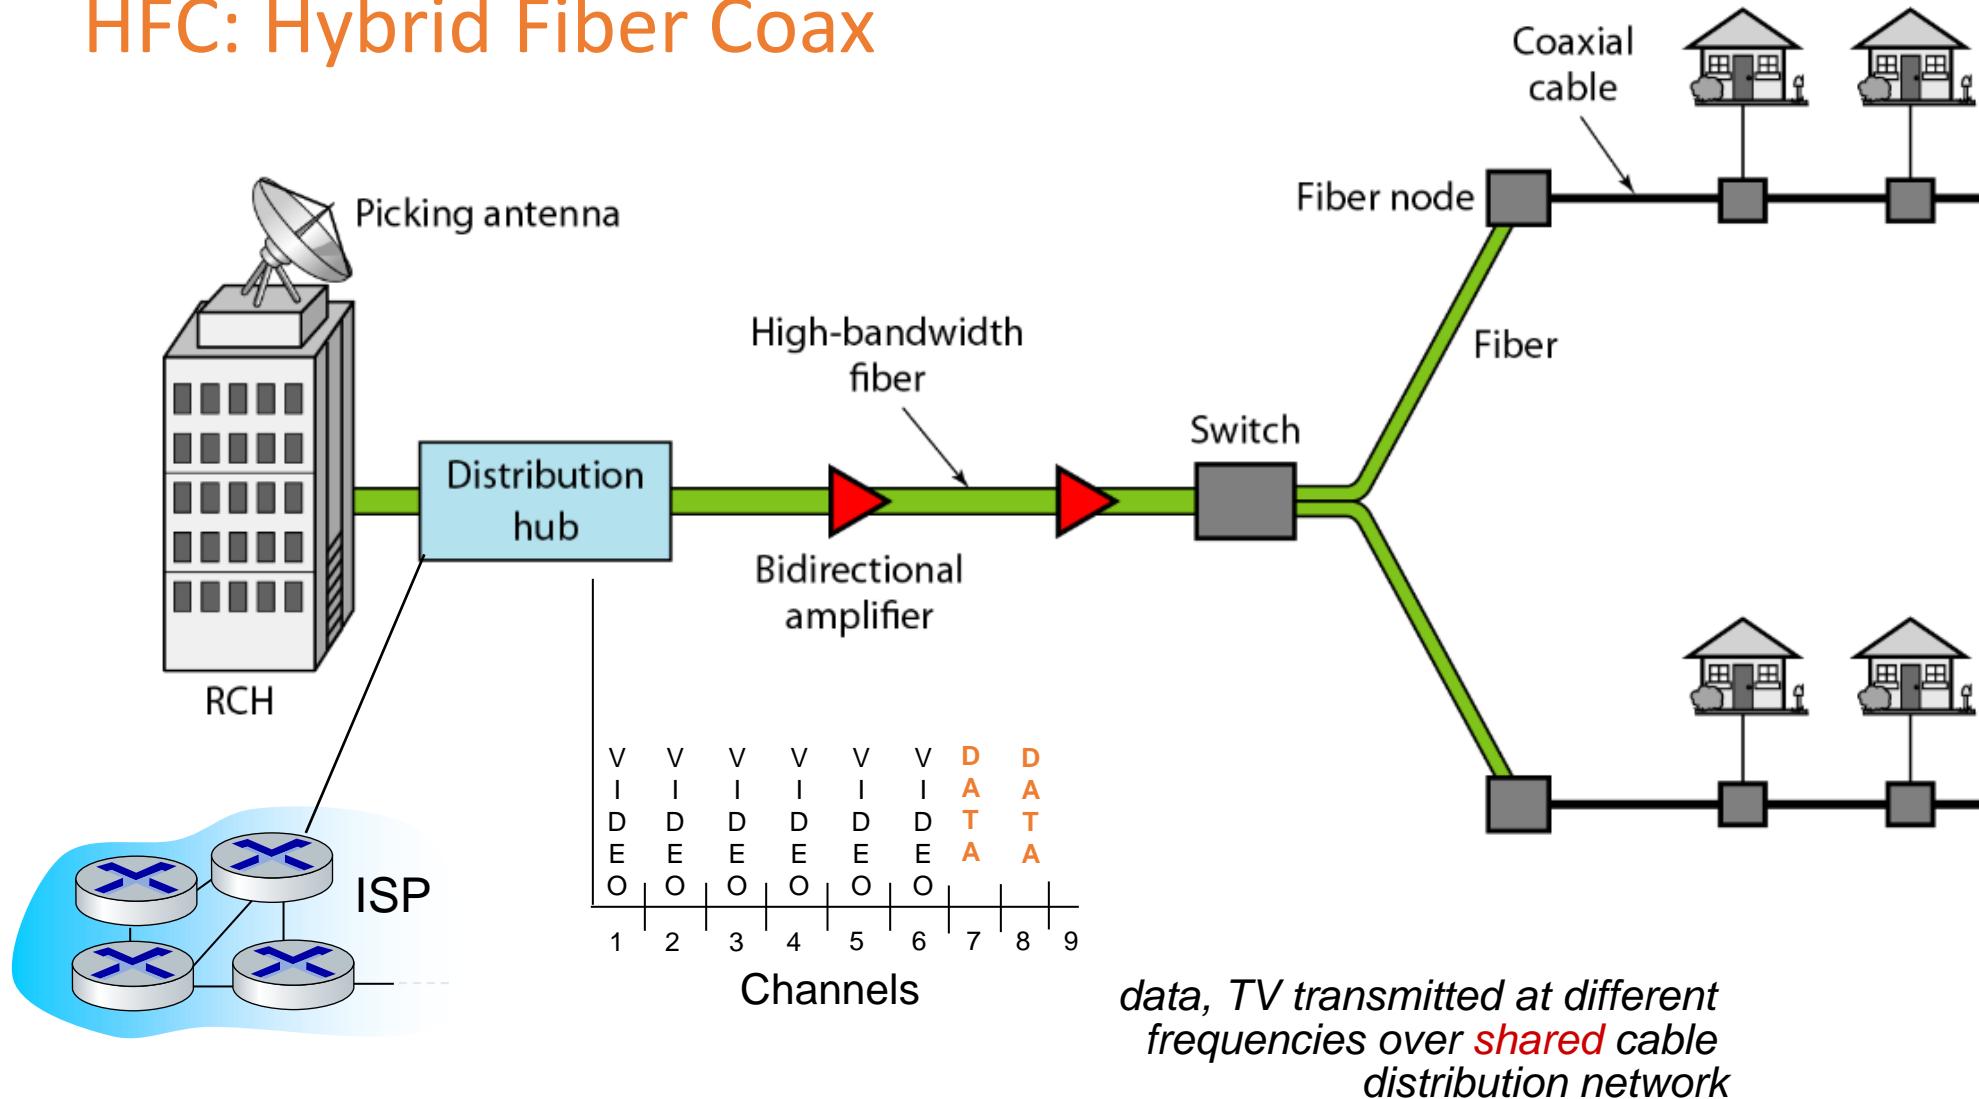
Access networks: cable-based access



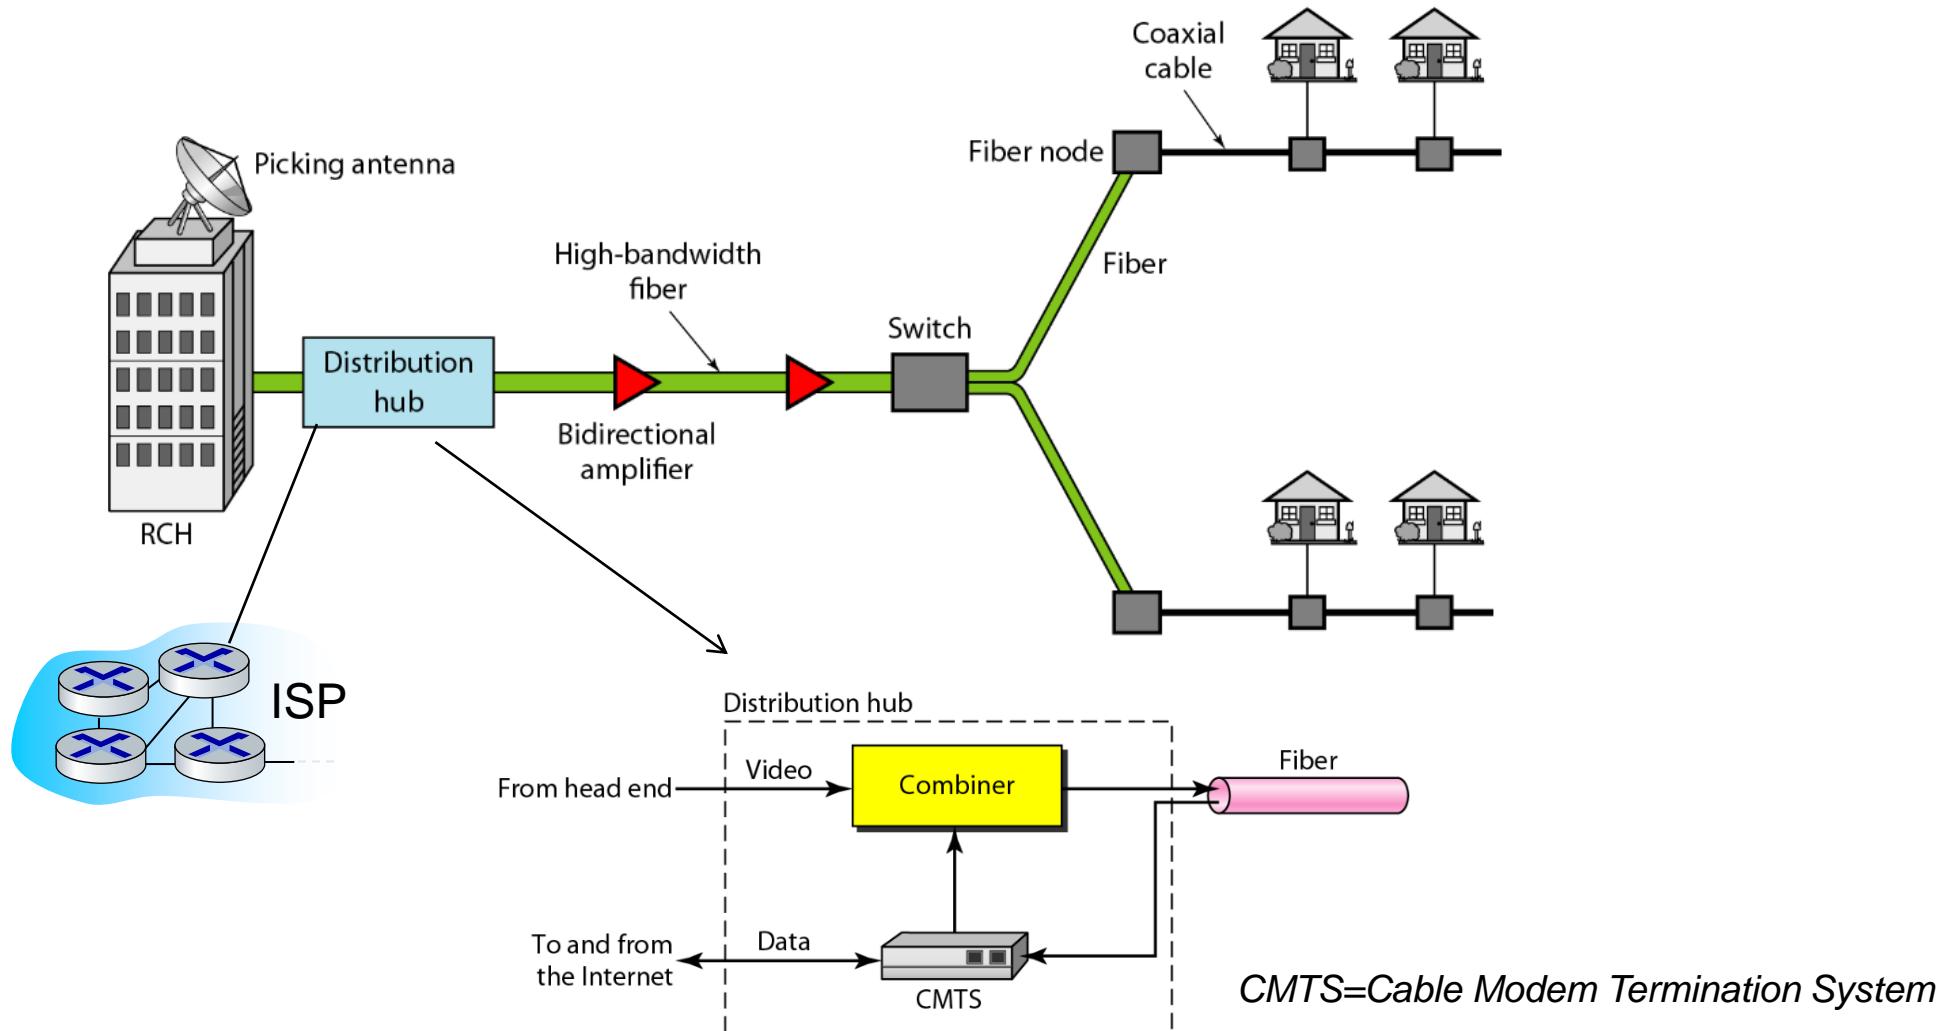
frequency division multiplexing (FDM): different channels transmitted in different frequency bands

Access networks: cable-based access

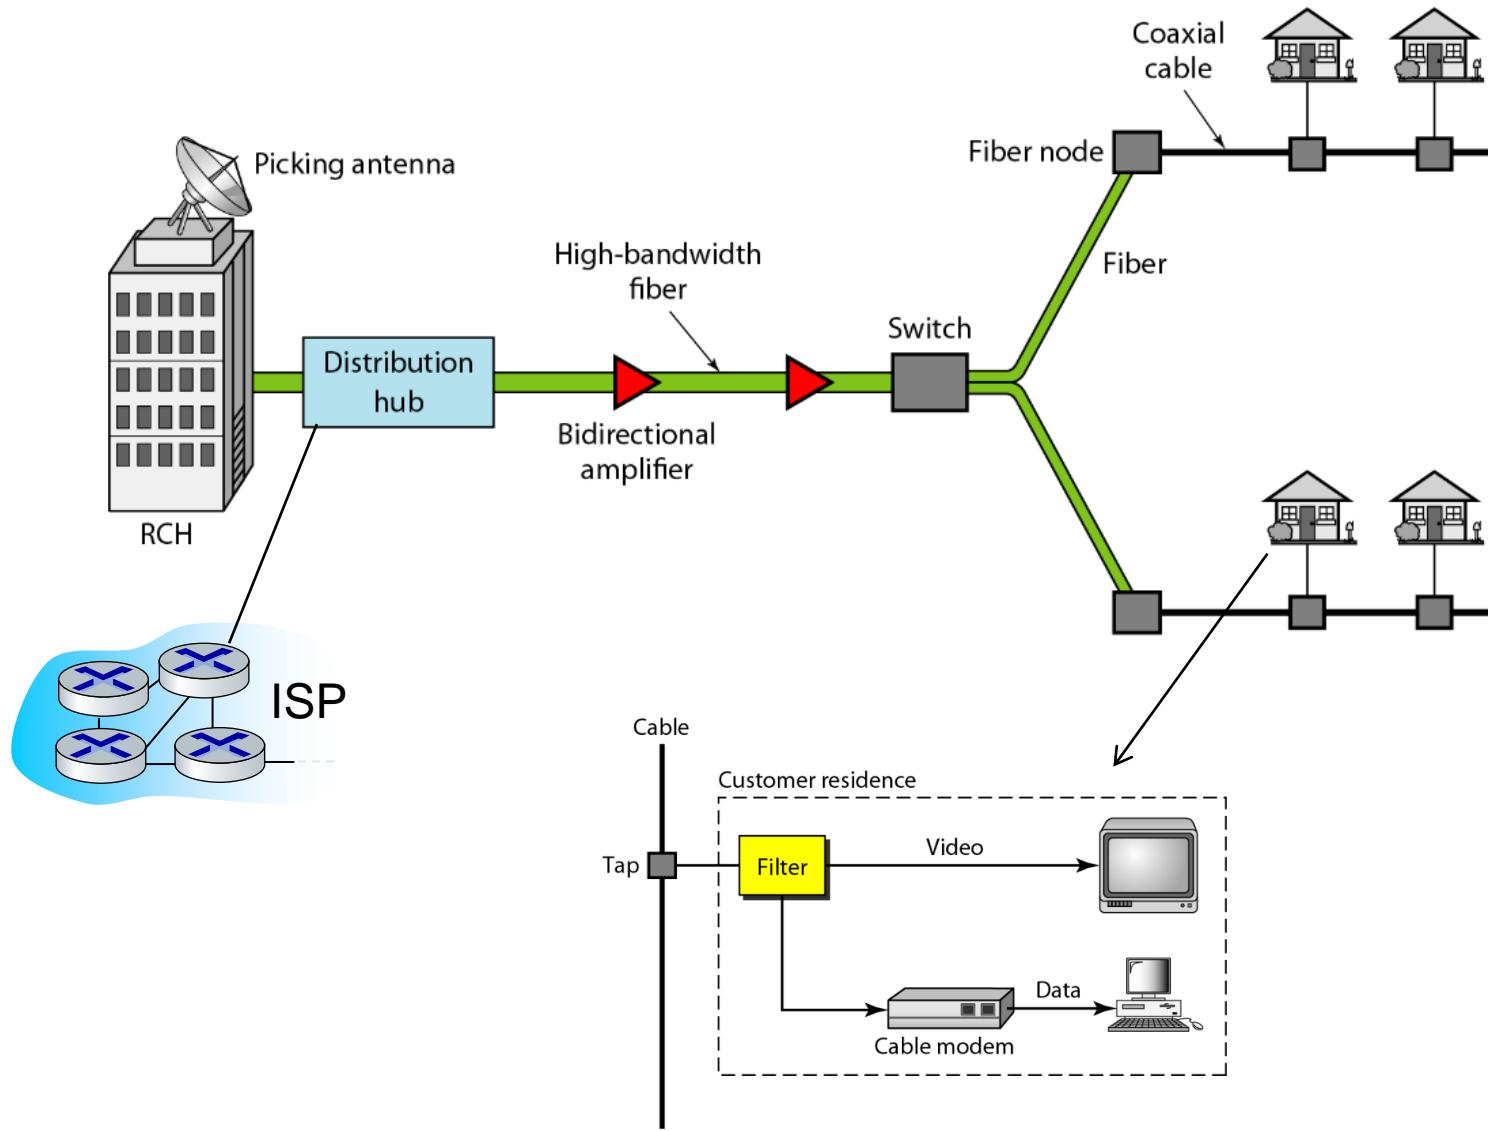
HFC: Hybrid Fiber Coax



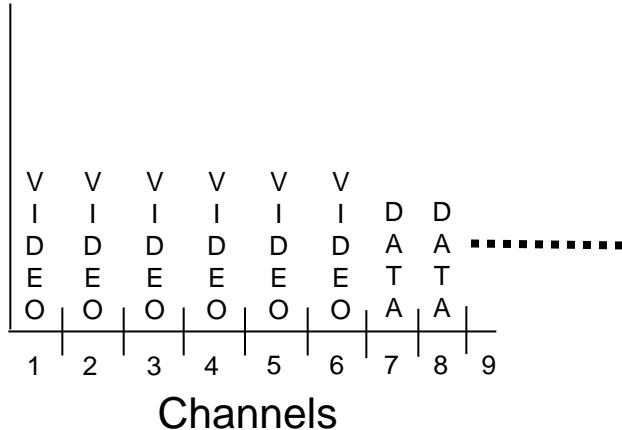
Access networks: cable-based access



Access networks: cable-based access



Access networks: cable-based access



Each channel is **shared** among a number of subscribers

Downlink: **multicasting**

Uplink: **multiple-access**

Downlink faster than uplink

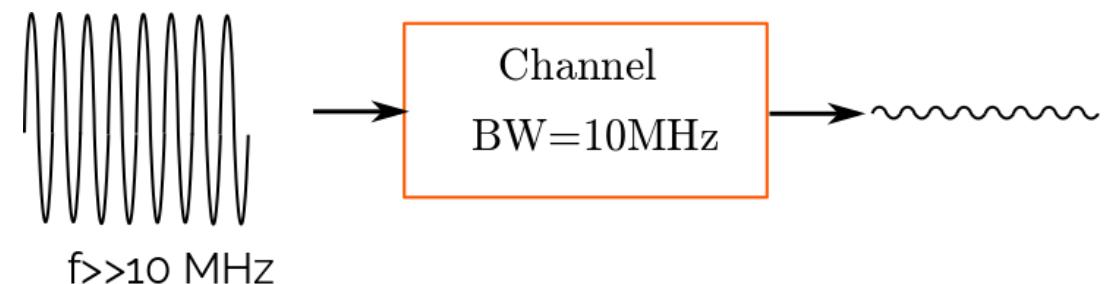
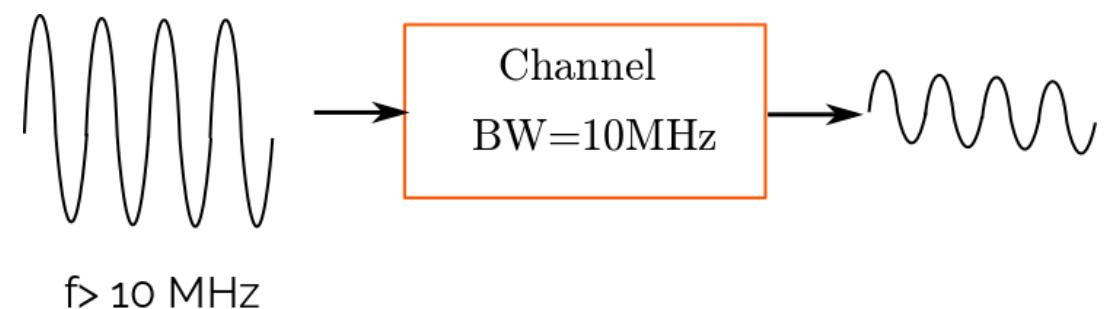
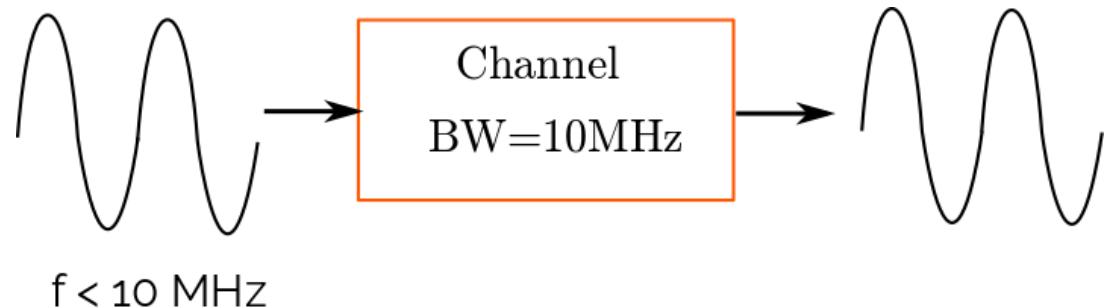
DL: 40 Mbps-1.2 Gbps, UL: 30-100 Mbps (DOCSIS 2.0, 3.0)
Because of different channel bandwidths/modulations

Aside: “Bandwidth”, “modulation”, and “bit rate”

Bandwidth:

Unit: Hertz (Hz)

Limits the max frequency of passing sines



Aside: “Bandwidth”, “modulation”, and “bit rate”

A pulse (carries a 0/1) can be written as the sum of infinite sines

- First harmony, second harmony, ...

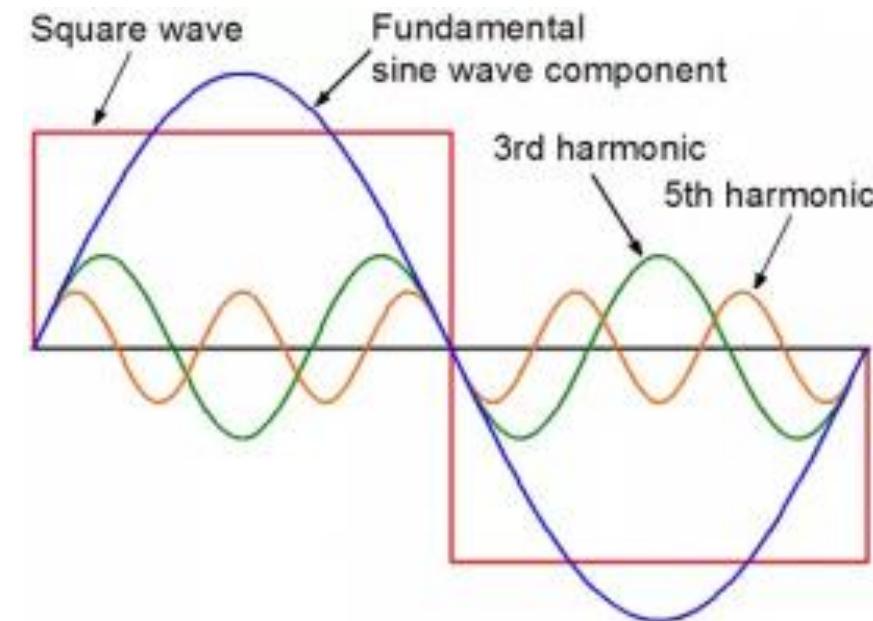
As the pulse's width decreases, the frequency of the first harmonic increases.

The BW of the channel limits the pulse rate

aka symbol rate

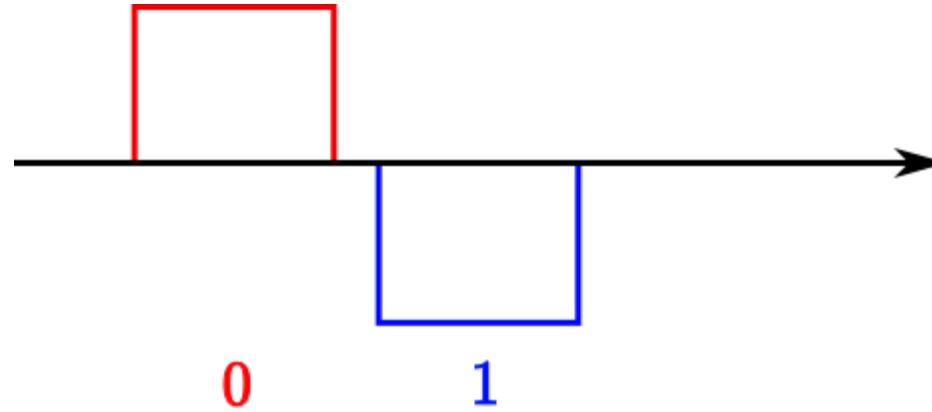
aka baud-rate

Baud rate \approx BW



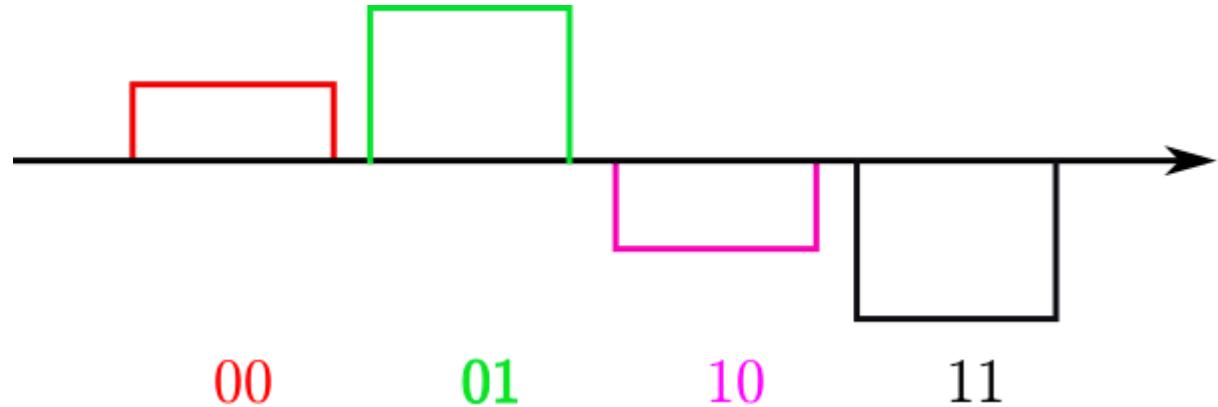
Aside: “Bandwidth”, “modulation”, and “bit rate”

A pulse can carry a single bit or more than one bit



Modulation: the number of levels

M= Number of bits per pulse = log (number of levels)



Aside: “Bandwidth”, “modulation”, and “bit rate”

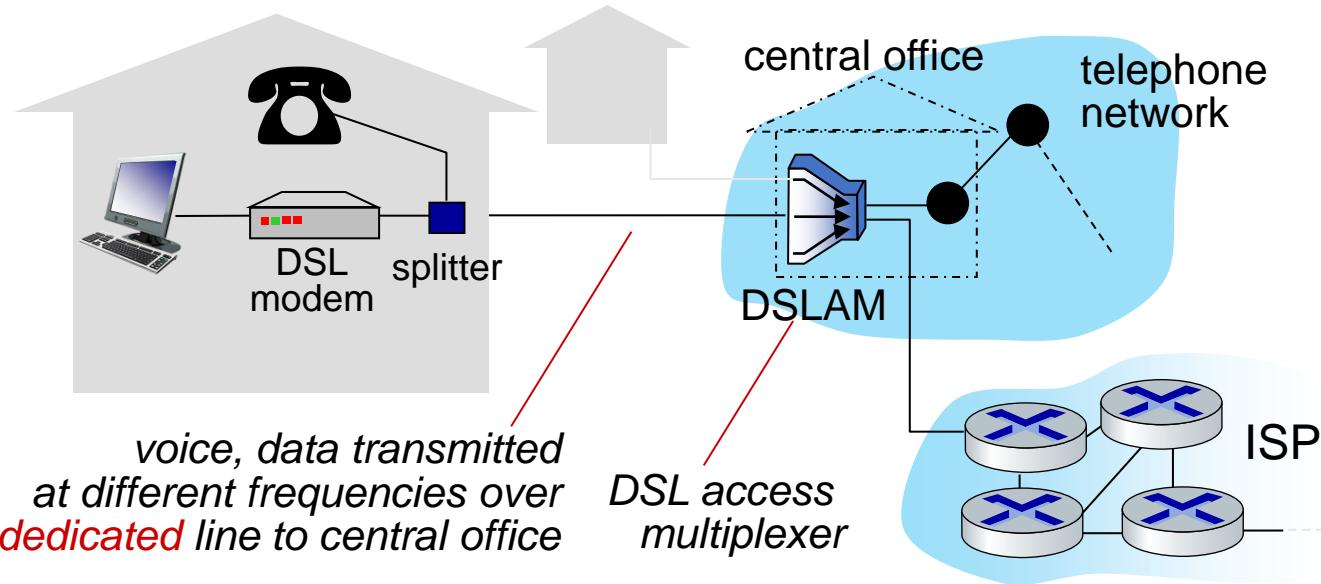
$$\text{Bit rate} = M \times \text{Baud-rate}$$

Aside: “Bandwidth”, “modulation”, and “bit rate”

But in Network:

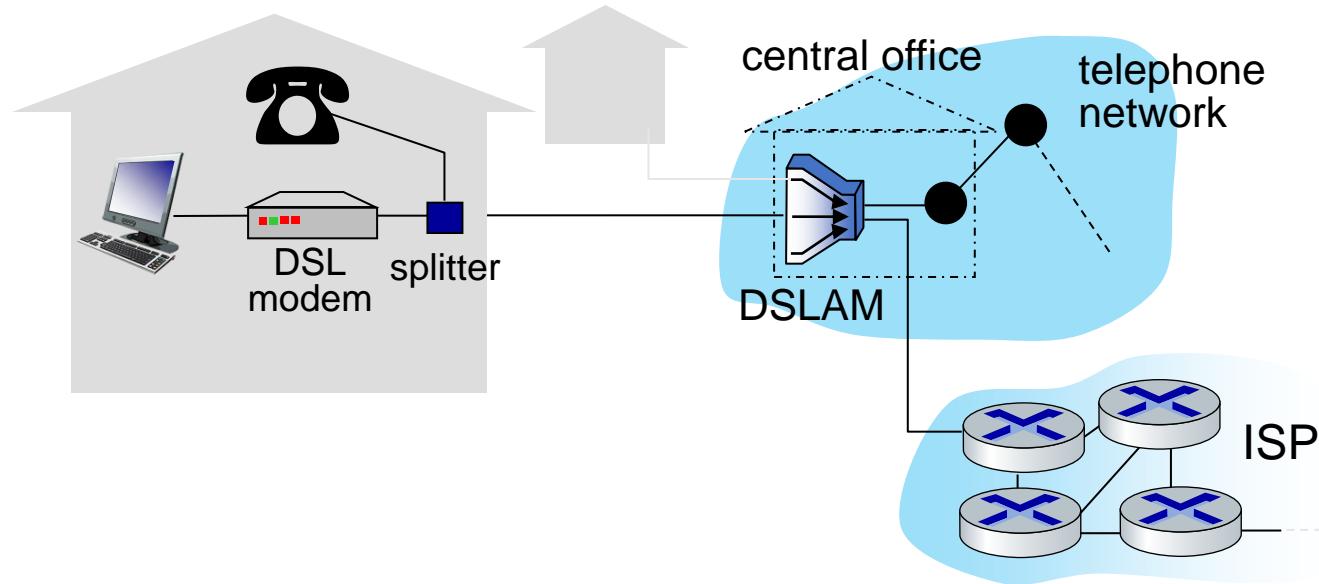
Bandwidth = Bit-rate

Access networks: digital subscriber line (DSL)



- use *existing* telephone line (local loop) to central office DSLAM
 - data over DSL phone line goes to Internet
 - voice over DSL phone line goes to telephone net
 - 24-52 Mbps **dedicated** downstream transmission rate
 - 3.5-16 Mbps **dedicated** upstream transmission rate
- } Asymmetric -> **ADSL**

Access networks: digital subscriber line (DSL)

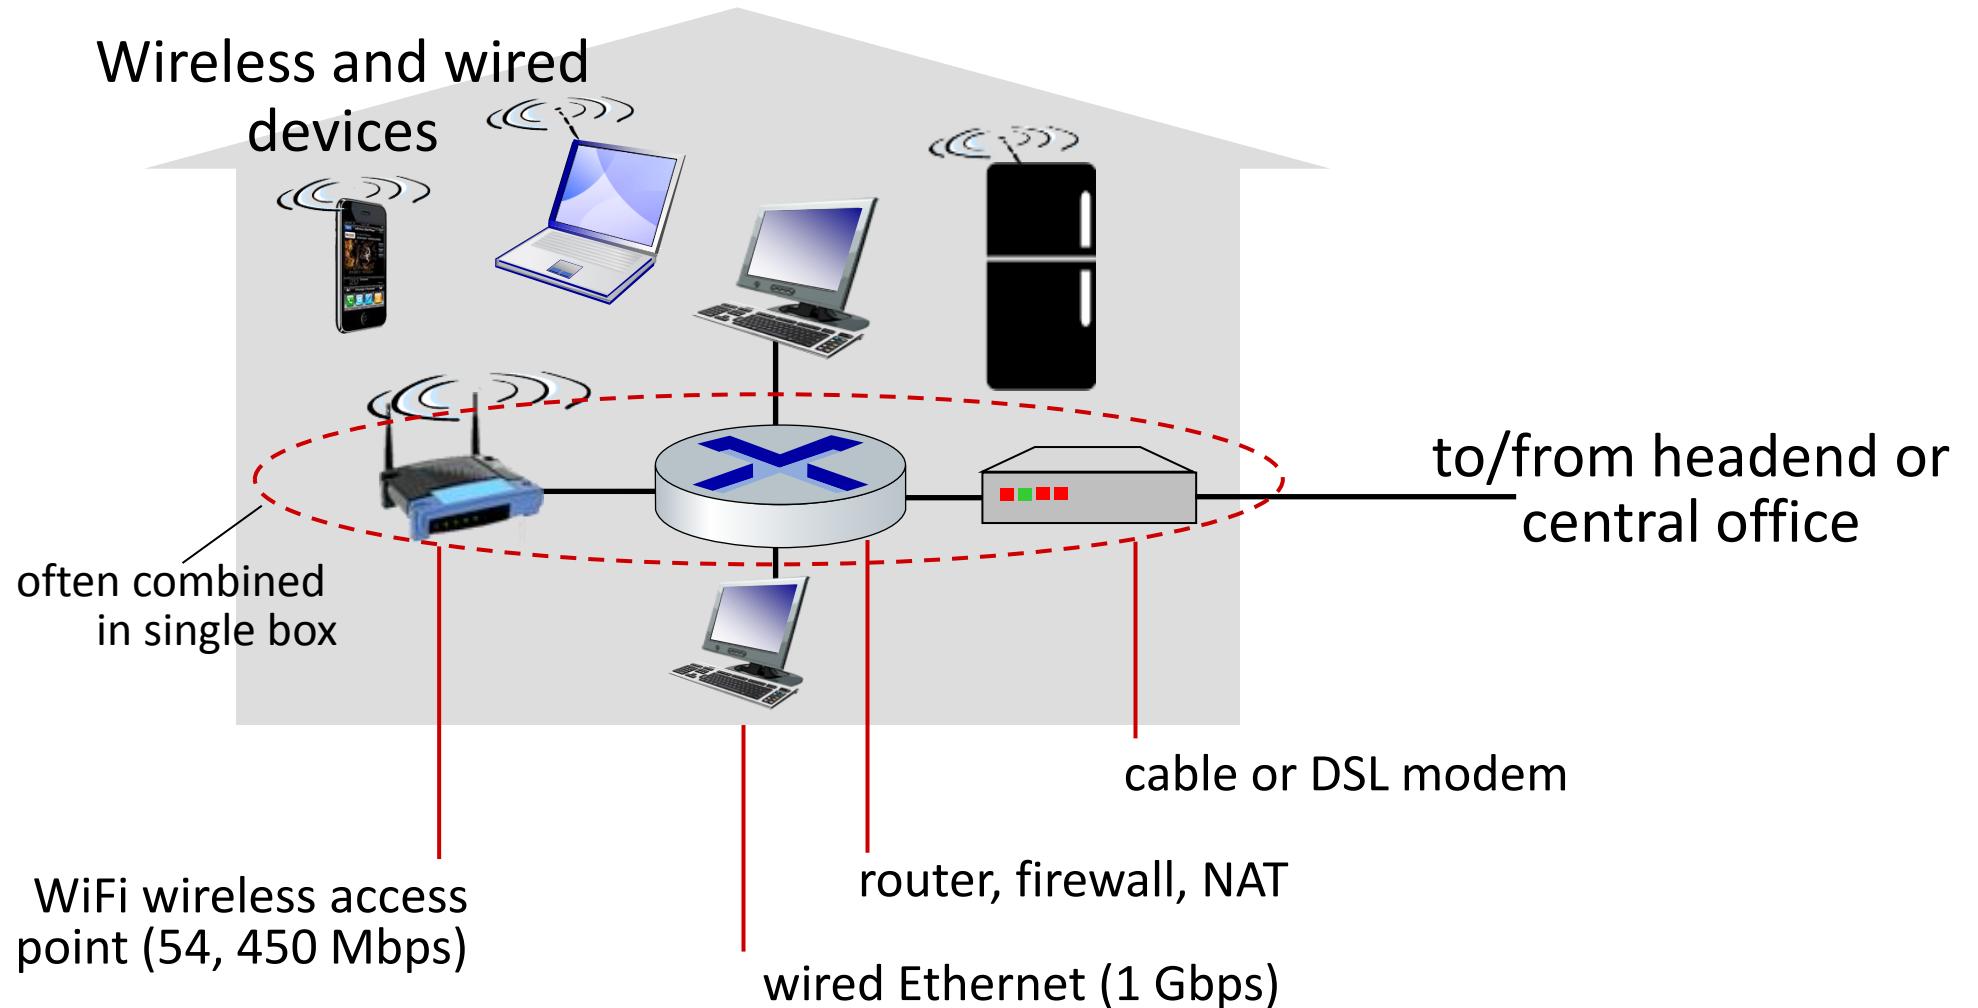


>> ADSL is dedicated but not independent!

DQ: Why?

>> ADSL channels vary -> ADSL modems work adaptively
time of the day
Distance to CO

Access networks: home networks



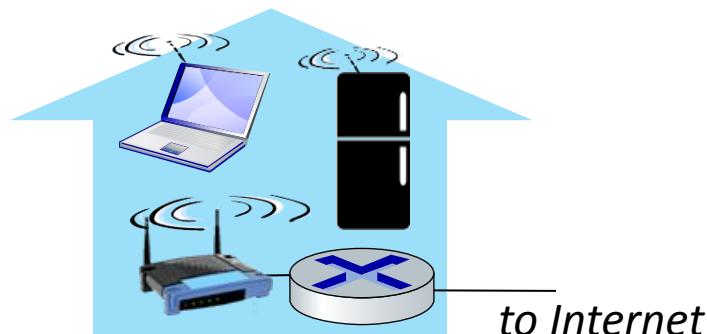
Wireless access networks

Shared *wireless* access network connects end system to router

- via base station aka “access point”

Wireless local area networks (WLANs)

- typically within or around building (~100 ft)
- 802.11b/g/n (WiFi): 11, 54, 450 Mbps transmission rate

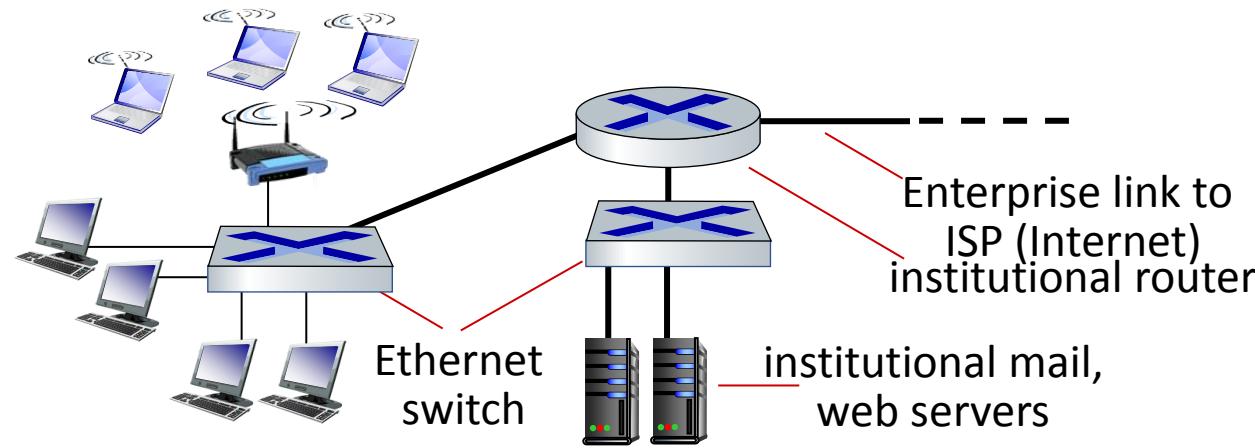


Wide-area cellular access networks

- provided by mobile, cellular network operator (10's km)
- 10's Mbps
- 4G cellular networks (5G coming)



Access networks: enterprise networks



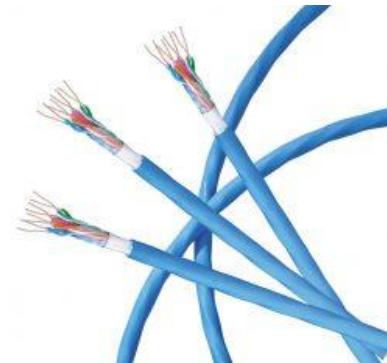
- companies, universities, etc.
- mix of wired, wireless link technologies, connecting a mix of switches and routers (we'll cover differences shortly)
 - Ethernet: wired access at 100Mbps, 1Gbps, 10Gbps
 - WiFi: wireless access points at 11, 54, 450 Mbps

Links: physical media

- **bit**: propagates between transmitter/receiver pairs
- **physical link**: what lies between transmitter & receiver
- **guided media**:
 - signals propagate in solid media: copper, fiber, coax
- **unguided media**:
 - signals propagate freely, e.g., radio

Twisted pair (TP)

- two insulated copper wires
 - Category 5: 100 Mbps, 1 Gbps Ethernet
 - Category 6: 10Gbps Ethernet



Links: physical media

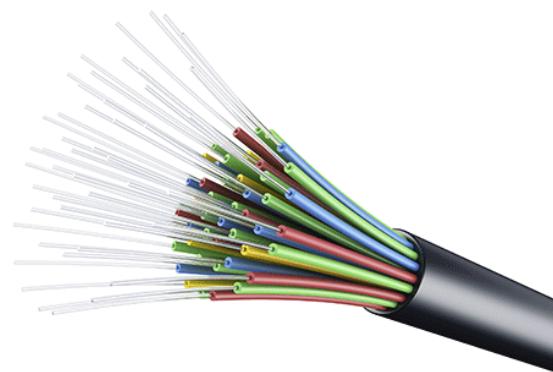
Coaxial cable:

- two concentric copper conductors
- bidirectional
- broadband:
 - multiple frequency channels on cable
 - 100's Mbps per channel



Fiber optic cable:

- glass fiber carrying light pulses, each pulse a bit
- high-speed operation:
 - high-speed point-to-point transmission (10's-100's Gbps)
- low error rate:
 - repeaters spaced far apart
 - immune to electromagnetic noise



Links: physical media

Wireless radio

- signal carried in various “bands” in electromagnetic spectrum
- no physical “wire”
- broadcast, “half-duplex” (sender to receiver)
- propagation environment effects:
 - reflection
 - obstruction by objects
 - Interference/noise

Radio link types:

- **Wireless LAN (WiFi)**
 - 10-100's Mbps; 10's of meters
- **wide-area** (e.g., 4G cellular)
 - 10's Mbps over ~10 Km
- **Bluetooth:** cable replacement
 - short distances, limited rates
- **terrestrial microwave**
 - point-to-point; 45 Mbps channels
- **satellite**
 - up to 45 Mbps per channel
 - 270 msec end-end delay

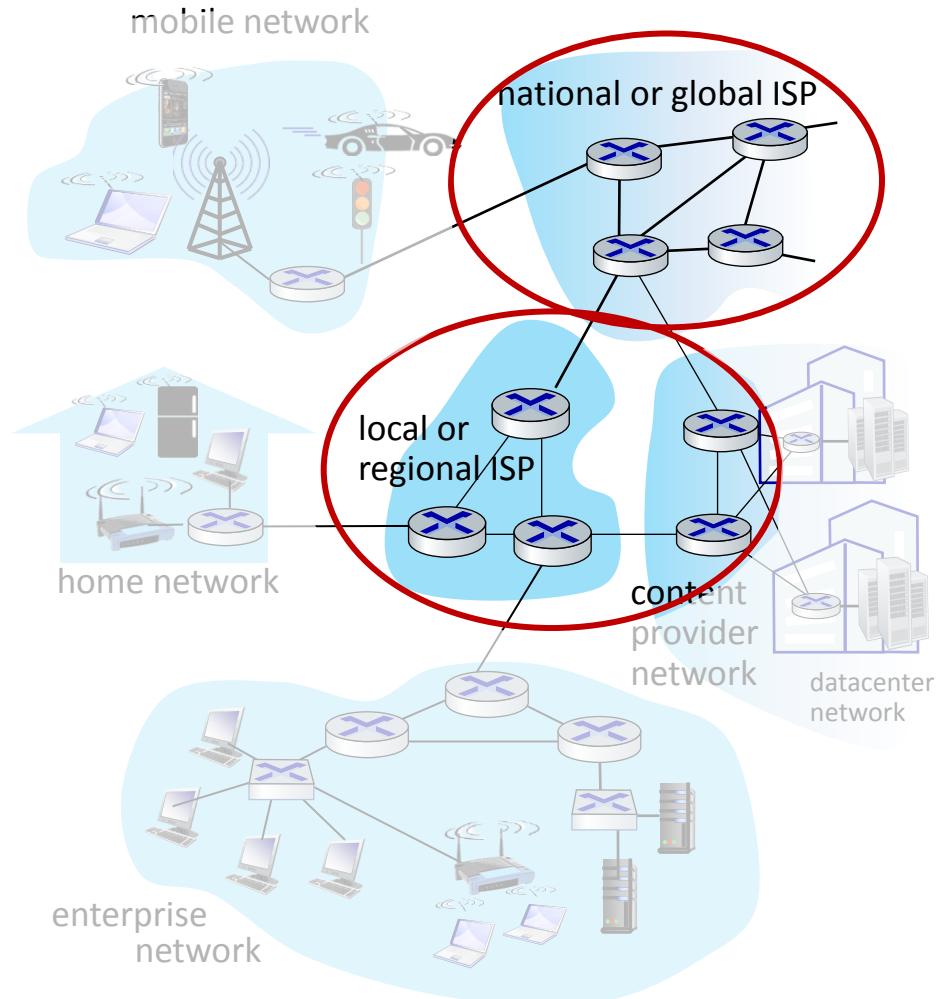
Chapter 1: roadmap

- What *is* the Internet?
- What *is* a protocol?
- Network edge: hosts, access network, physical media
- **Network core:** packet/circuit switching, internet structure
- Performance: loss, delay, throughput
- Security
- Protocol layers, service models
- History



The network core

- A network at core: mesh of interconnected routers/switches
- Operating based on **packet switching** or **circuit switching**



How networks Work ...

Packet switching

Circuit switching

How networks Work ...

Packet switching

Circuit switching

Packet-switching

packet-switching: hosts break application-layer messages into
packets

Packets = payload+header

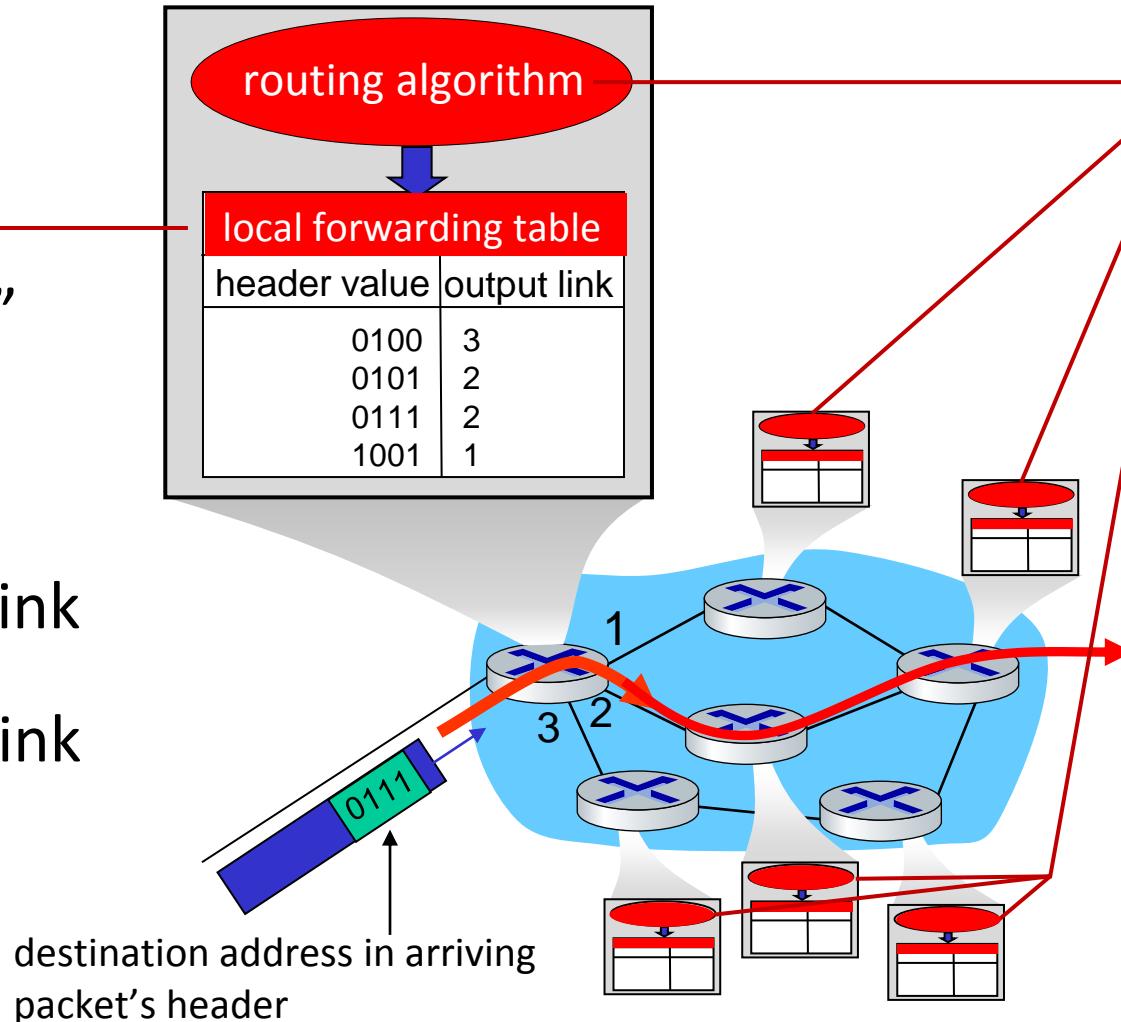
Packets have source/destination addresses inside their headers

network **forwards** packets from one router to the next, across
links on path from **source to destination**

Packet-switching network: functionalities

Forwarding:

- aka “switching”
- *local* action:
move arriving
packets from
router’s input link
to appropriate
router output link



Routing:

- *global* action:
determine source-
destination paths
taken by packets
- routing algorithms





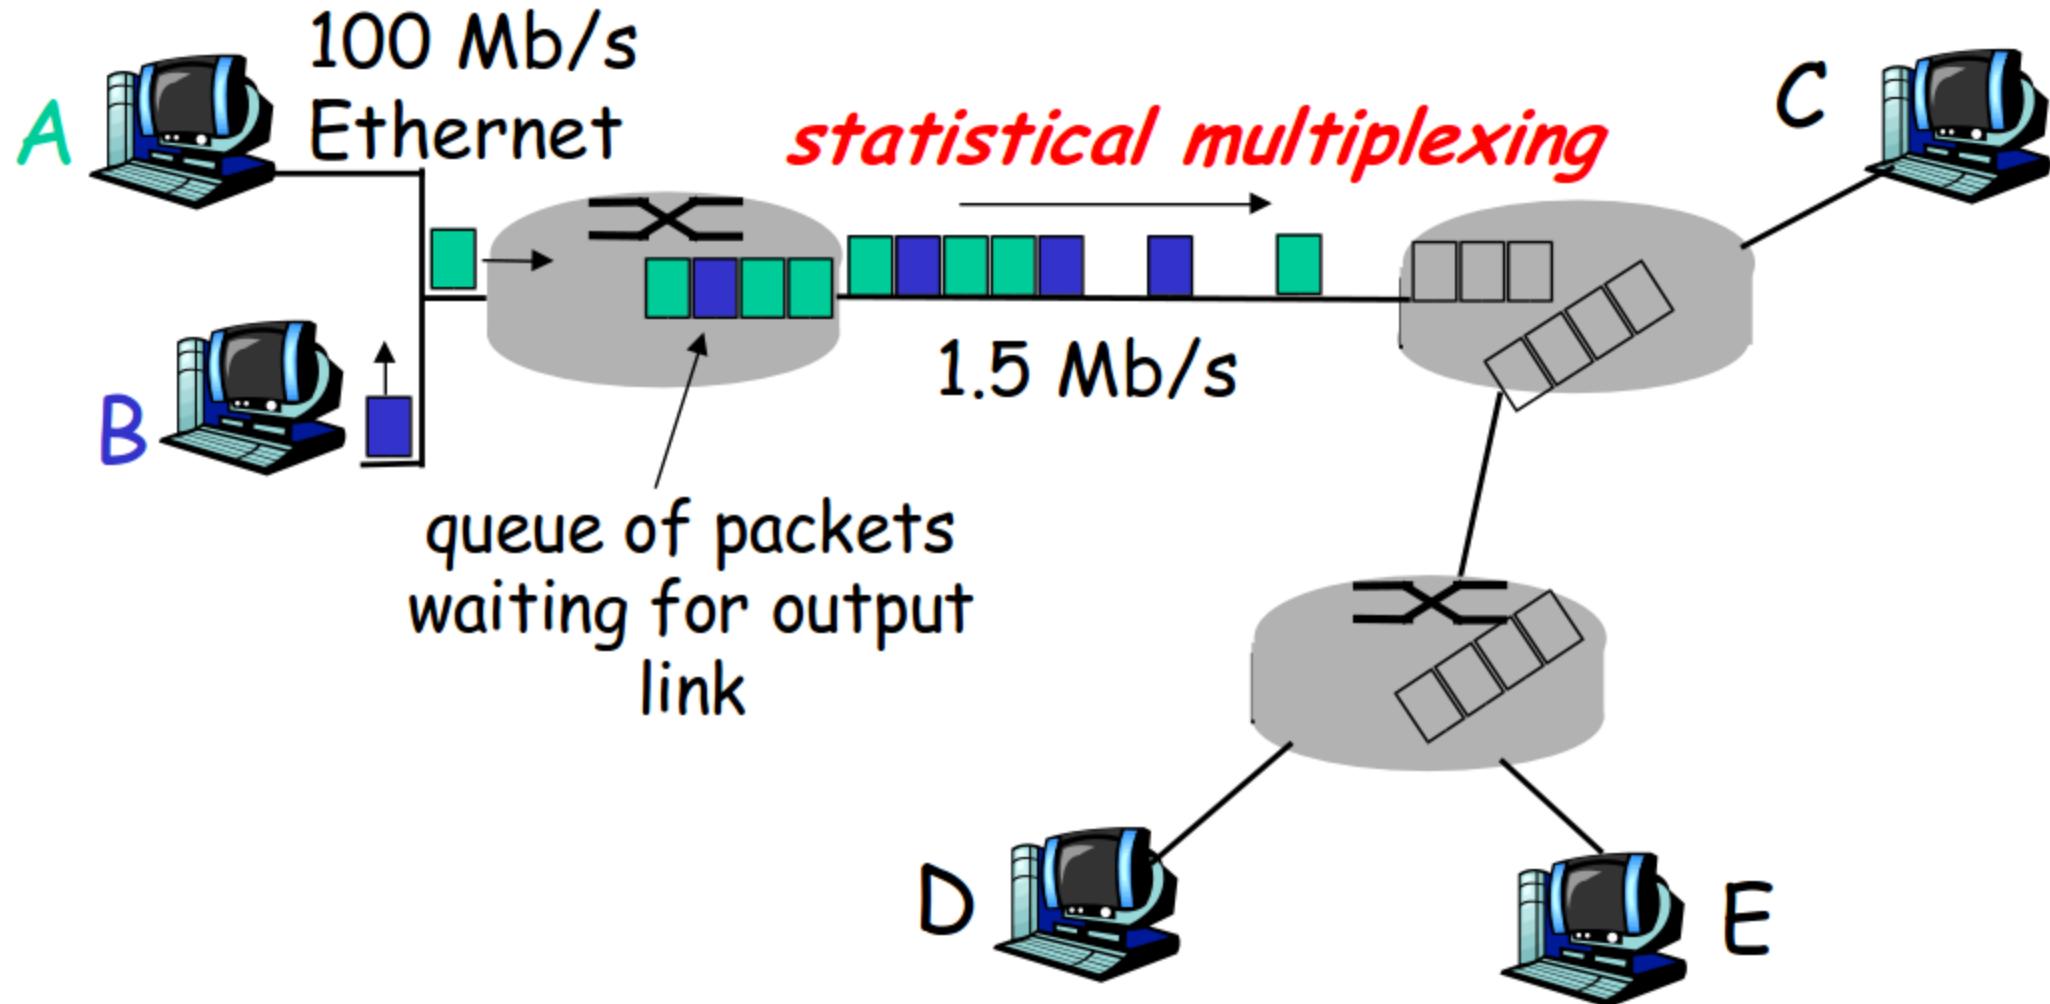
Packet-switching

packet-switching: hosts break application-layer messages into *packets*

>> No resource reservation (no call setup)

+ Resource sharing (statistical multiplexing)





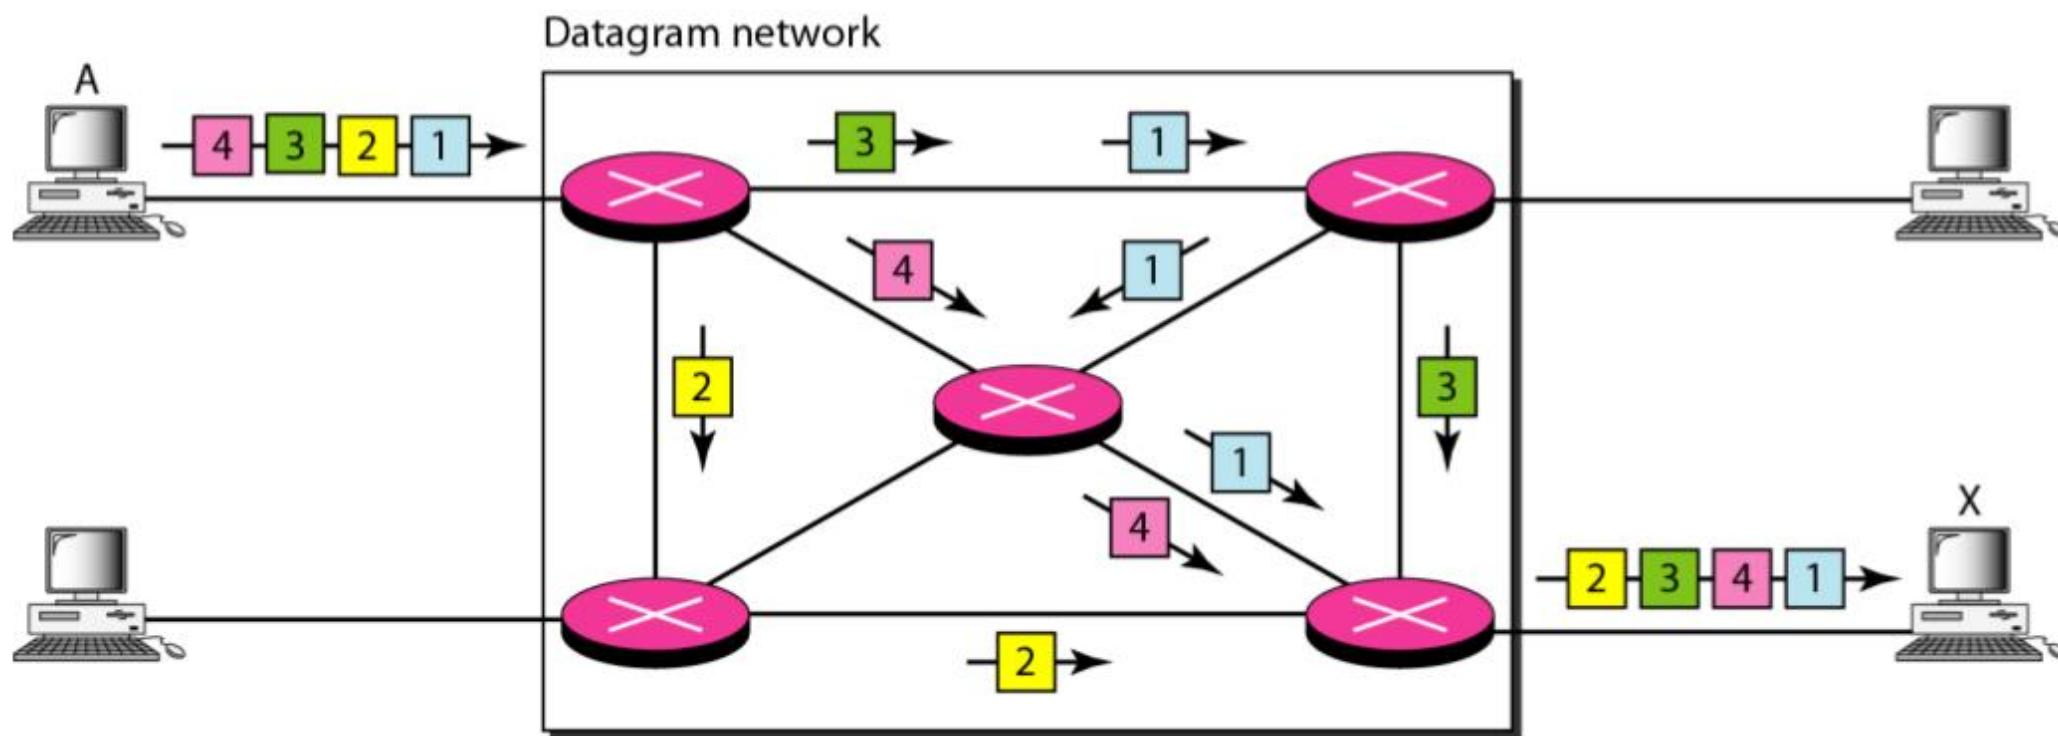
Packet-switching

packet-switching: hosts break application-layer messages into *packets*

>> No resource reservation (no call setup)

- + Resource sharing (statistical multiplexing)
- Different packets are routed independently, and may take different paths





Packet-switching

packet-switching: hosts break application-layer messages into *packets*

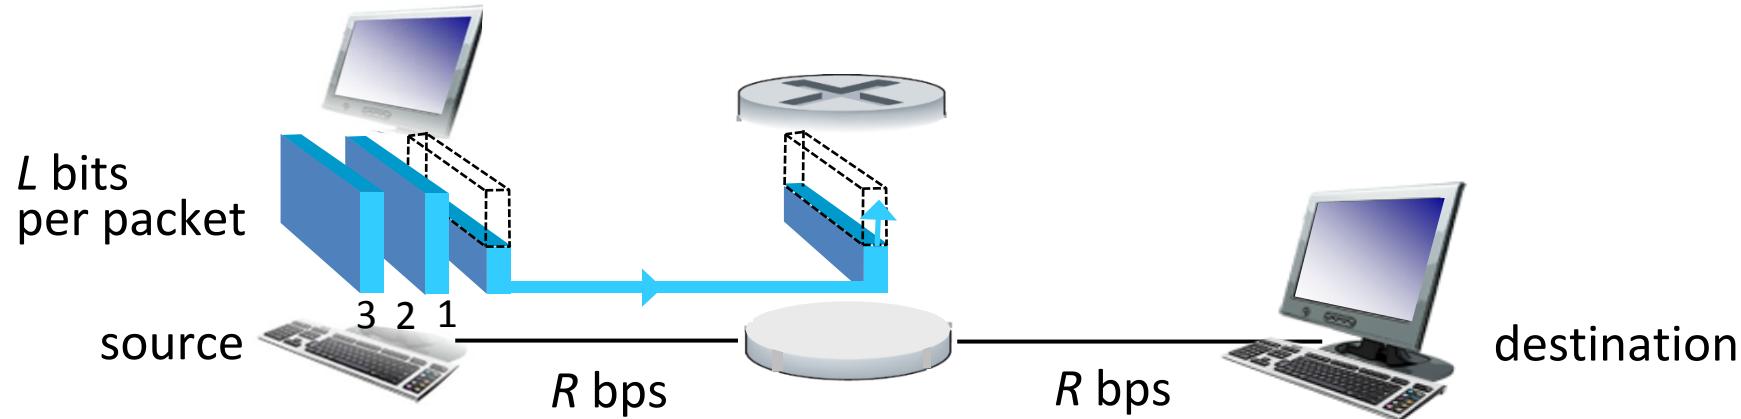
>> No resource reservation (no call setup)

- + Resource sharing (statistical multiplexing)
- Different packets are routed independently, and may
 - take different paths
 - have different delays (missing orders)
 - be dropped (lost)

BEST EFFORT!



Store and forward

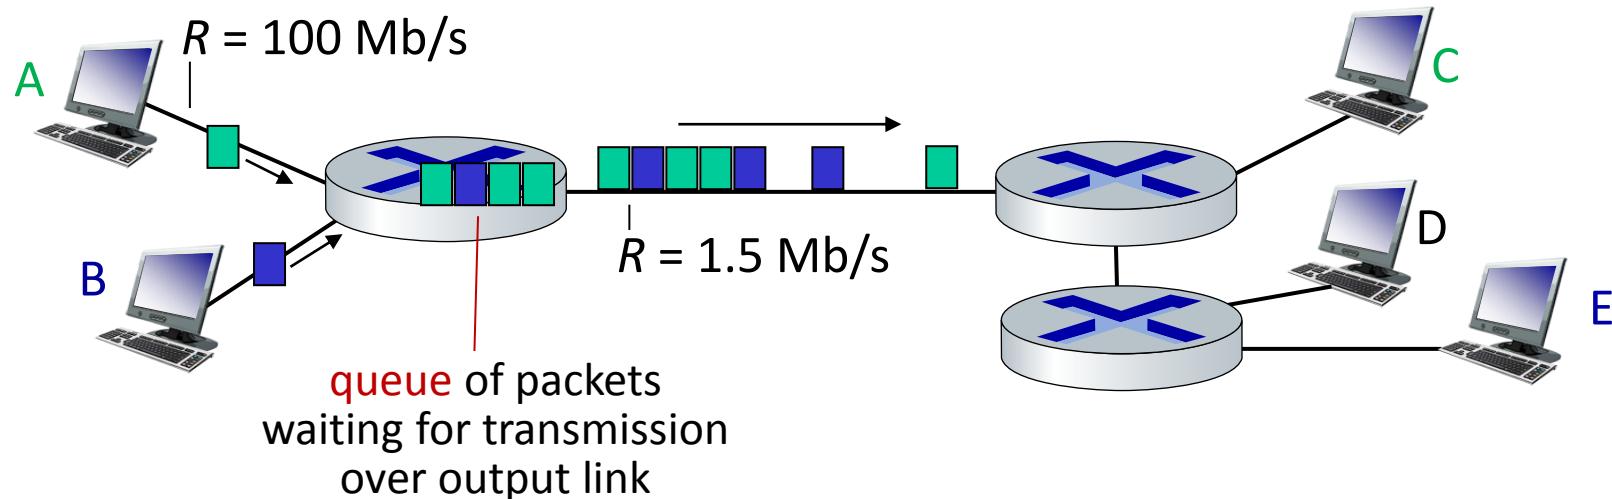


- **packet transmission delay:** takes L/R seconds to transmit (push out) L -bit packet into link at R bps
- ***store and forward:*** entire packet must arrive at router before it can be transmitted on next link

One-hop numerical example:

- $L = 10$ Kbits
- $R = 100$ Mbps
- one-hop transmission delay = 0.1 msec

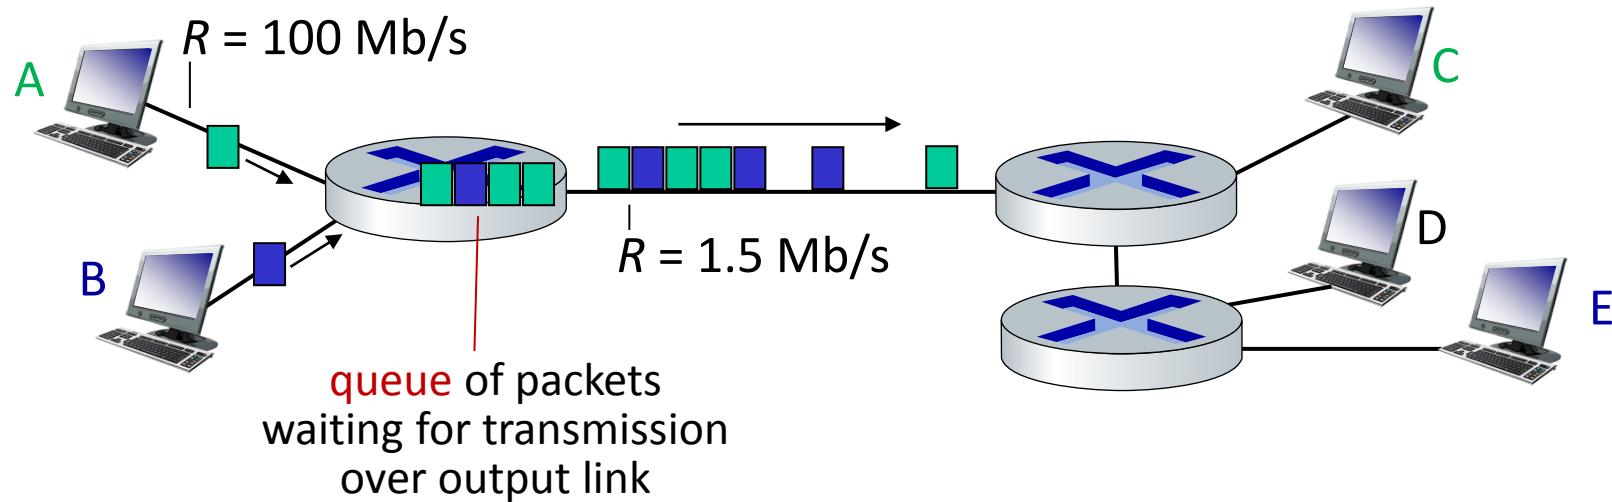
Queueing



Queueing occurs when work arrives faster than it can be serviced:



Queueing



Packet queuing and loss: if arrival rate (in bps) to link exceeds transmission rate (bps) of link for some period of time:

- packets will queue, waiting to be transmitted on output link
- packets can be dropped (lost) if memory (buffer) in router fills up

How networks Work ...

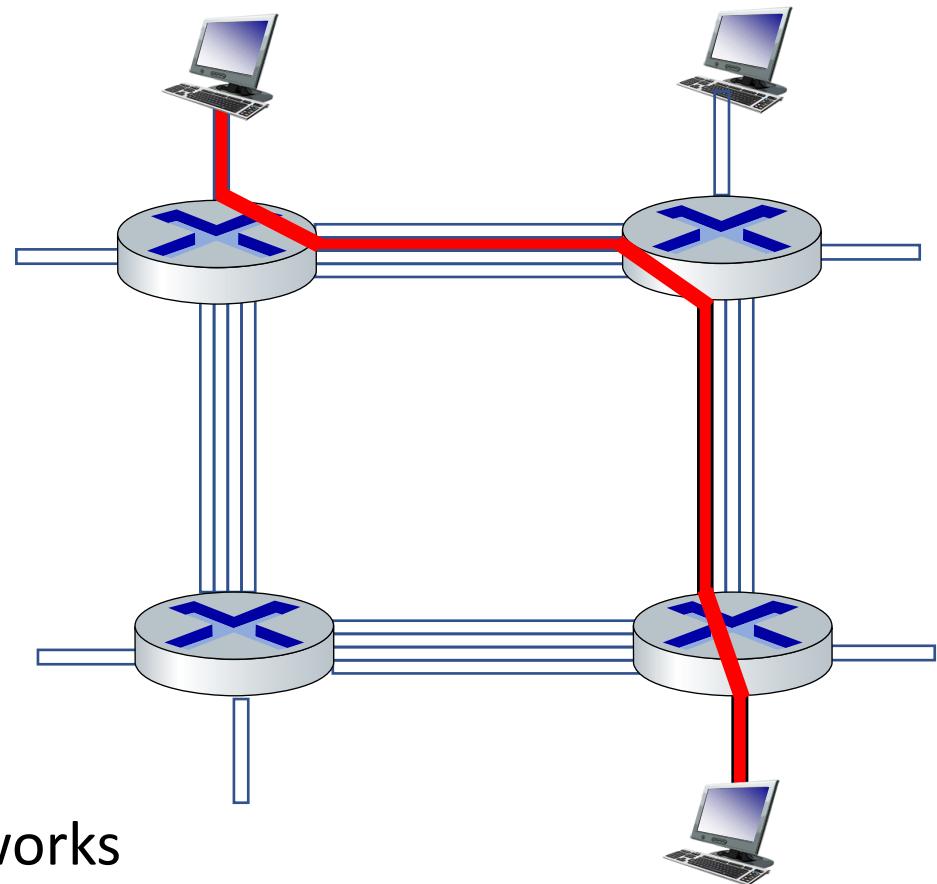
Packet switching

Circuit switching

Circuit switching

end-end resources allocated to,
reserved for “call” between source
and destination

- in diagram, each link has four circuits.
 - call gets 2nd circuit in top link and 1st circuit in right link.
- dedicated resources: no sharing
 - circuit-like (guaranteed) performance
- circuit segment idle if not used by call (**no sharing**)
- commonly used in traditional telephone networks



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive

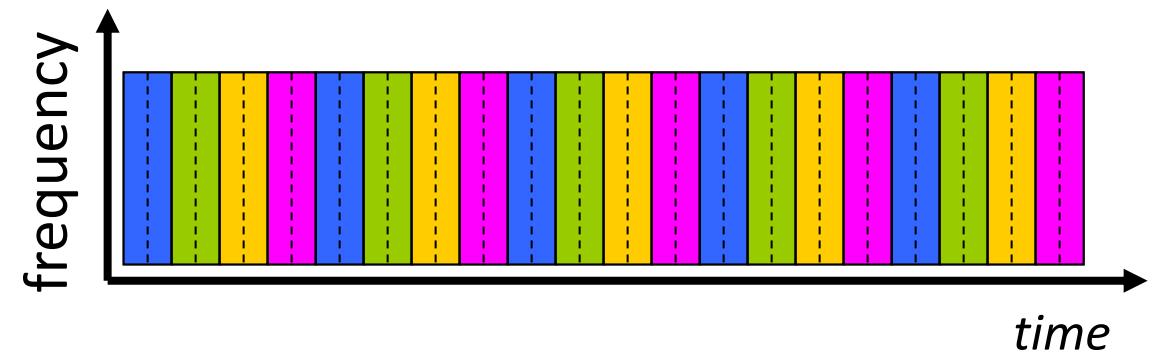
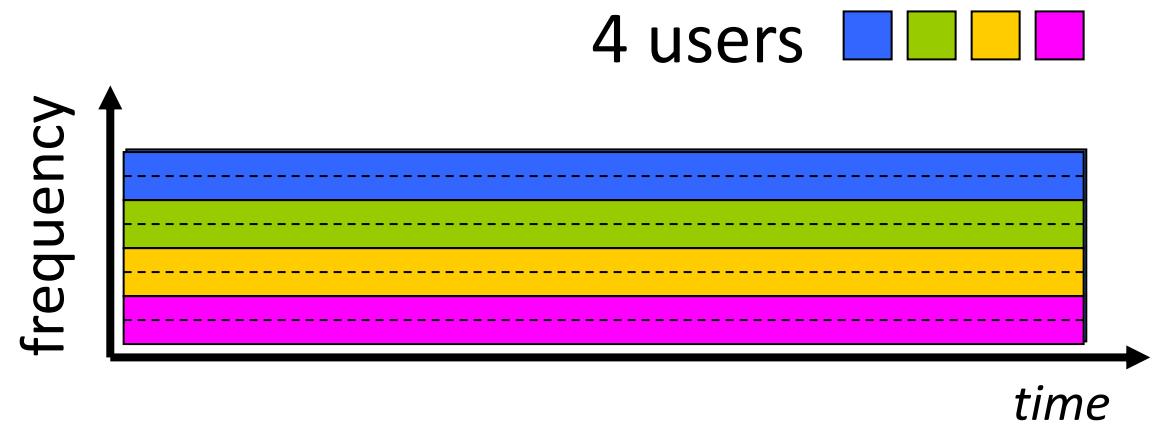
FDM and TDM

Frequency Division Multiplexing (FDM)

- optical, electromagnetic frequencies divided into (narrow) frequency bands
- each call allocated its own band, can transmit at max rate of that narrow band

Time Division Multiplexing (TDM)

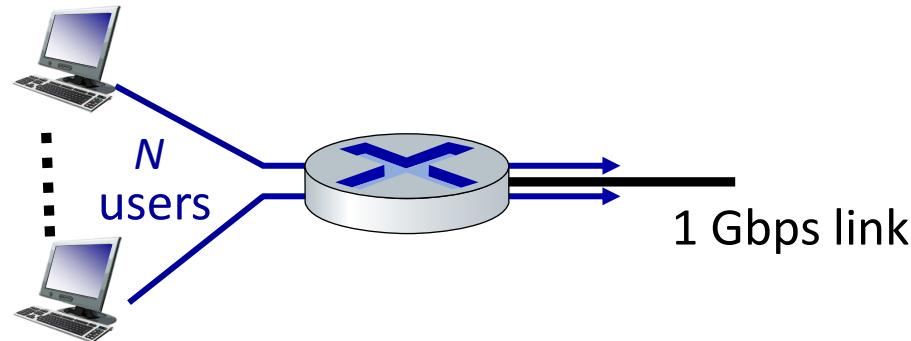
- time divided into slots
- each call allocated periodic slot(s), can transmit at maximum rate of (wider) frequency band (only) during its time slot(s)



Example: Packet switching vs. circuit switching

example:

- 1 Gb/s link
- each user:
 - 100 Mb/s when “active”
 - active 10% of time



Q: how many users can use this network under circuit-switching and packet switching?

- *circuit-switching:* 10 users
- *packet switching:* with 35 users,
probability > 10 active at same time
is less than .0004 *

Q: how did we get value 0.0004?
A: HW problem (for those with
course in probability only)

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive

Packet switching vs. circuit switching

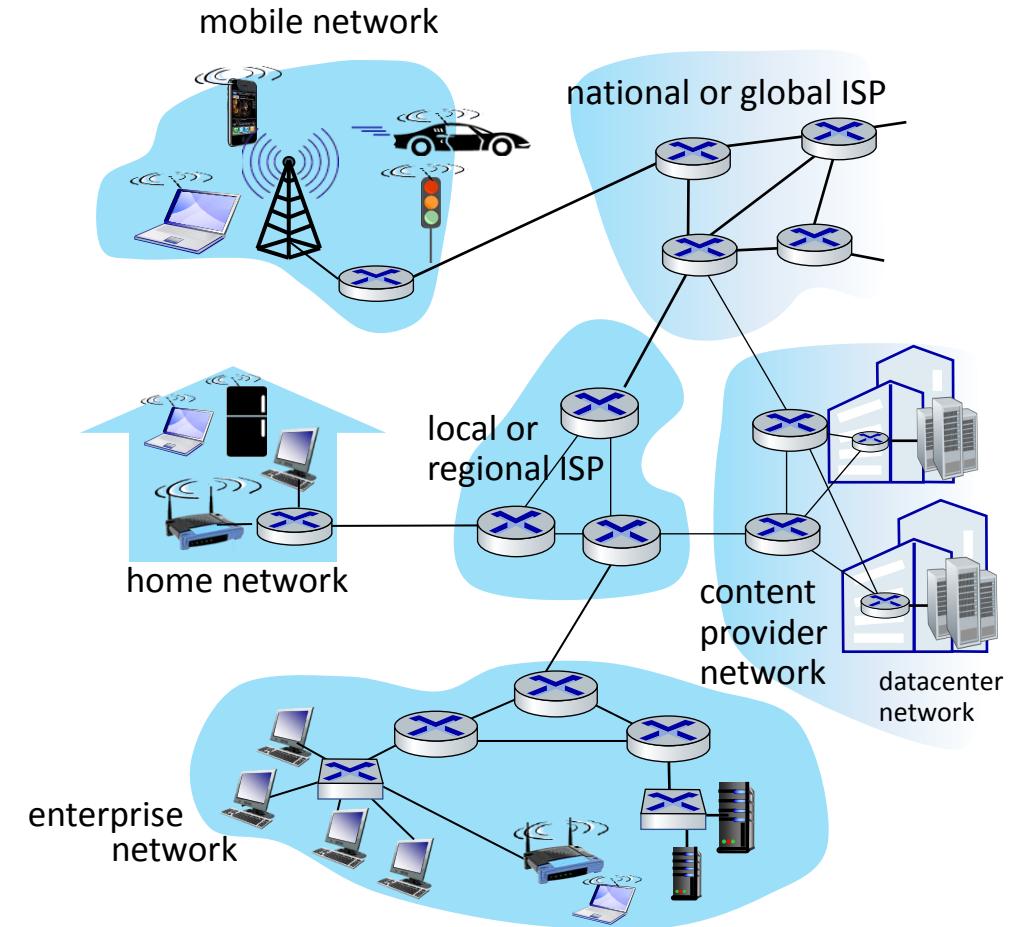
Is packet switching a “slam dunk winner”?

- great for “bursty” data – sometimes has data to send, but at other times not
 - resource sharing
 - simpler, no call setup
- **excessive congestion possible:** packet delay and loss due to buffer overflow
 - protocols needed for reliable data transfer, congestion control
- ***Q: How to provide circuit-like behavior with packet-switching?***
 - “It’s complicated.” We’ll study various techniques that try to make packet switching as “circuit-like” as possible.

Internet Structure

Internet structure: a “network of networks”

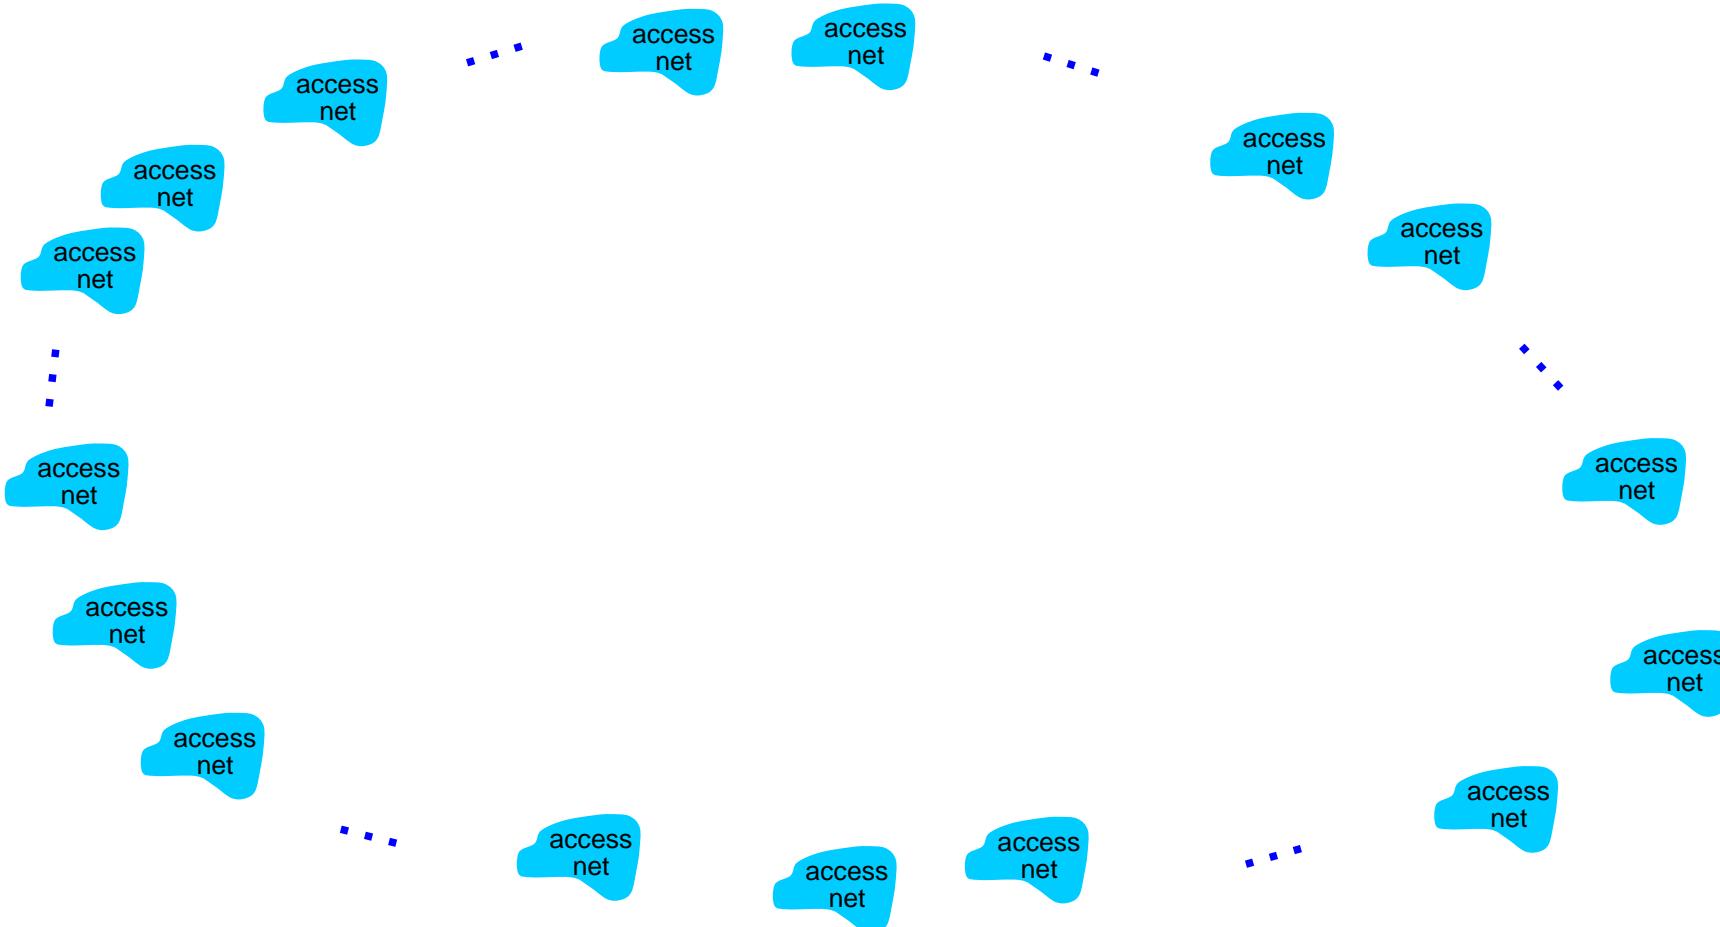
- hosts connect to Internet via **access** Internet Service Providers (ISPs)
- access ISPs in turn must be interconnected
 - so that *any* two hosts (*anywhere!*) can send packets to each other
- resulting network of networks is very complex
 - evolution driven by **economics, national policies**



Let's take a stepwise approach to describe current Internet structure

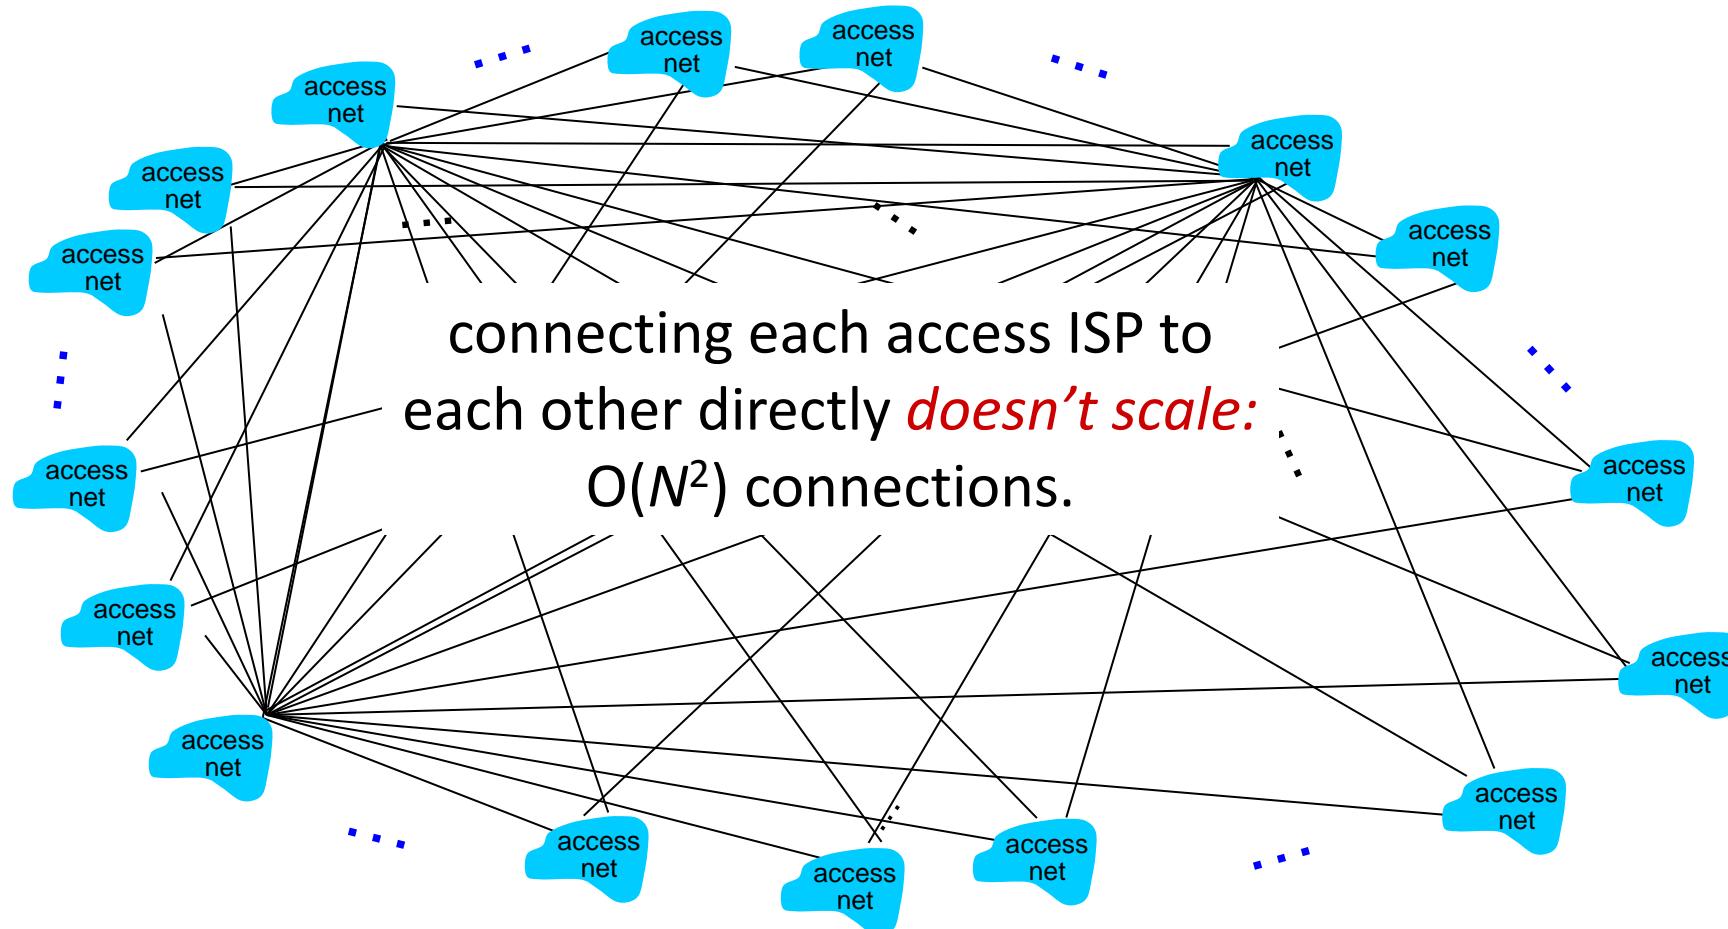
Internet structure: a “network of networks”

Question: given *millions* of access ISPs, how to connect them together?



Internet structure: a “network of networks”

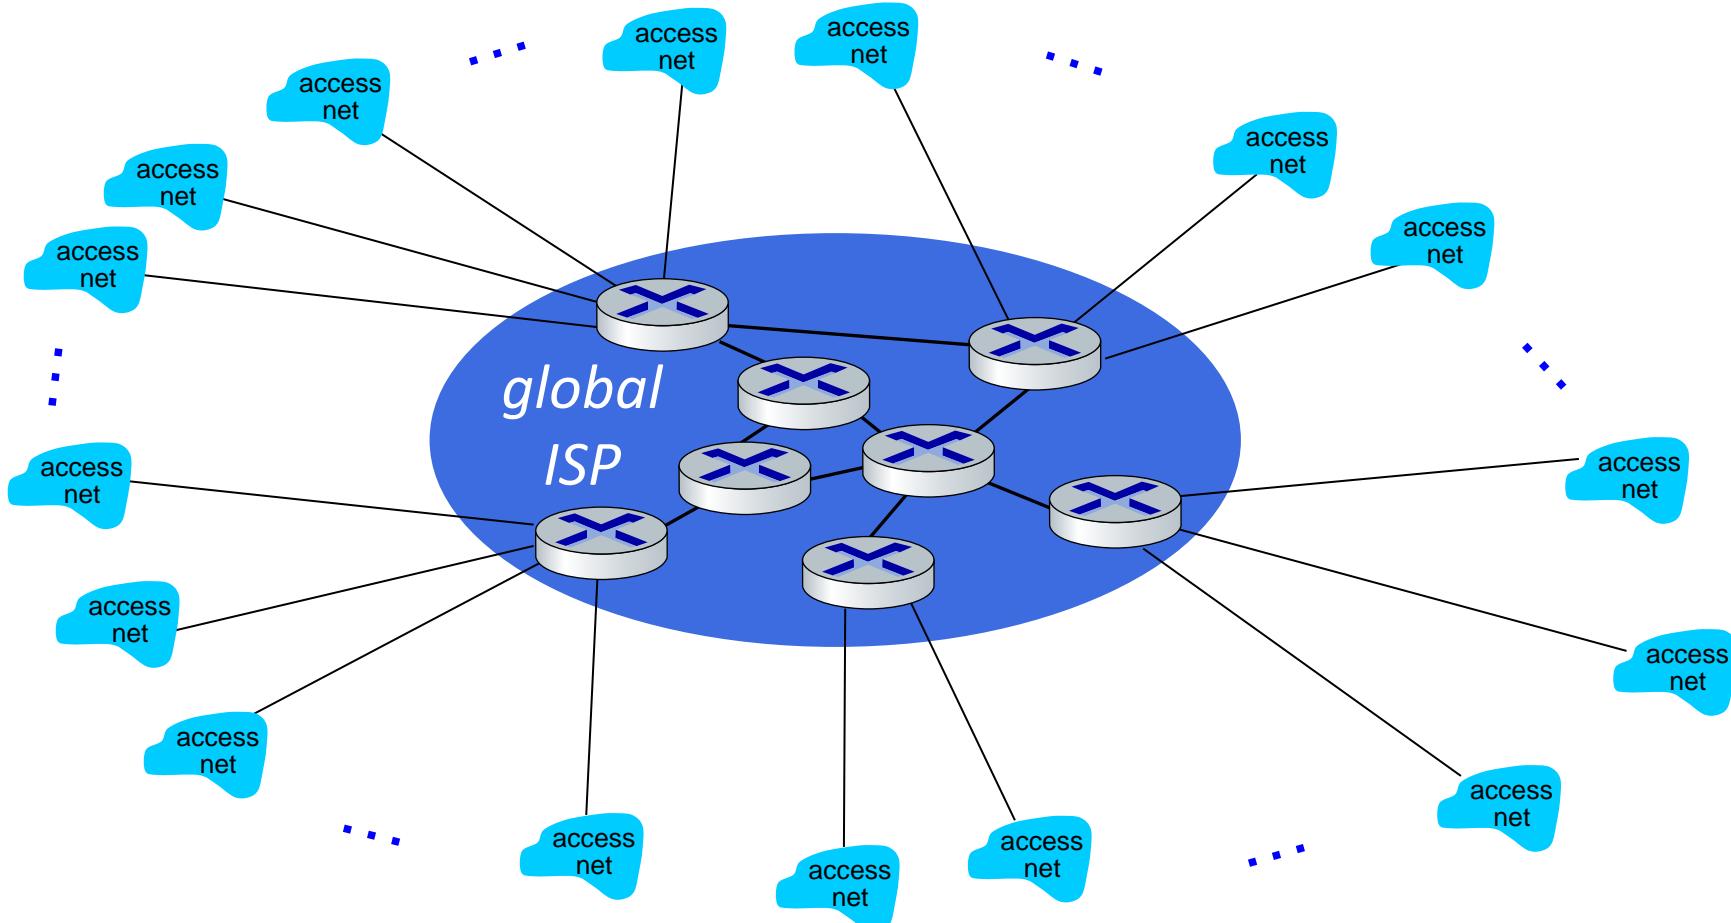
Question: given *millions* of access ISPs, how to connect them together?



Internet structure: a “network of networks”

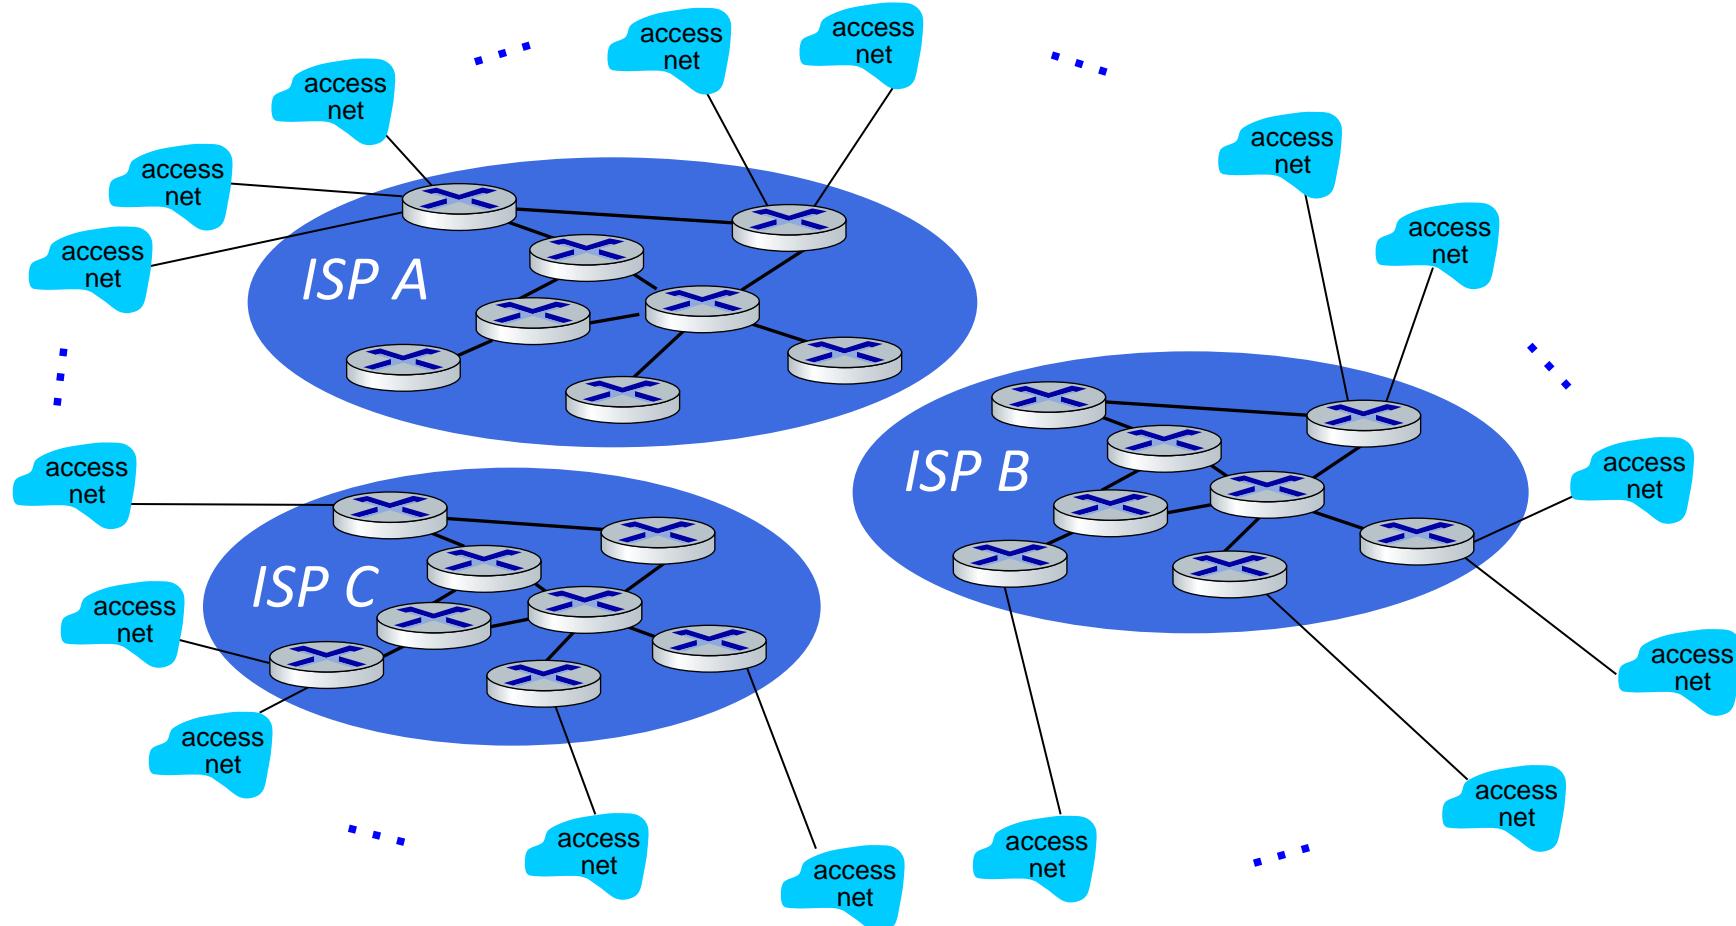
Option: connect each access ISP to one global transit ISP?

Customer and provider ISPs have economic agreement.



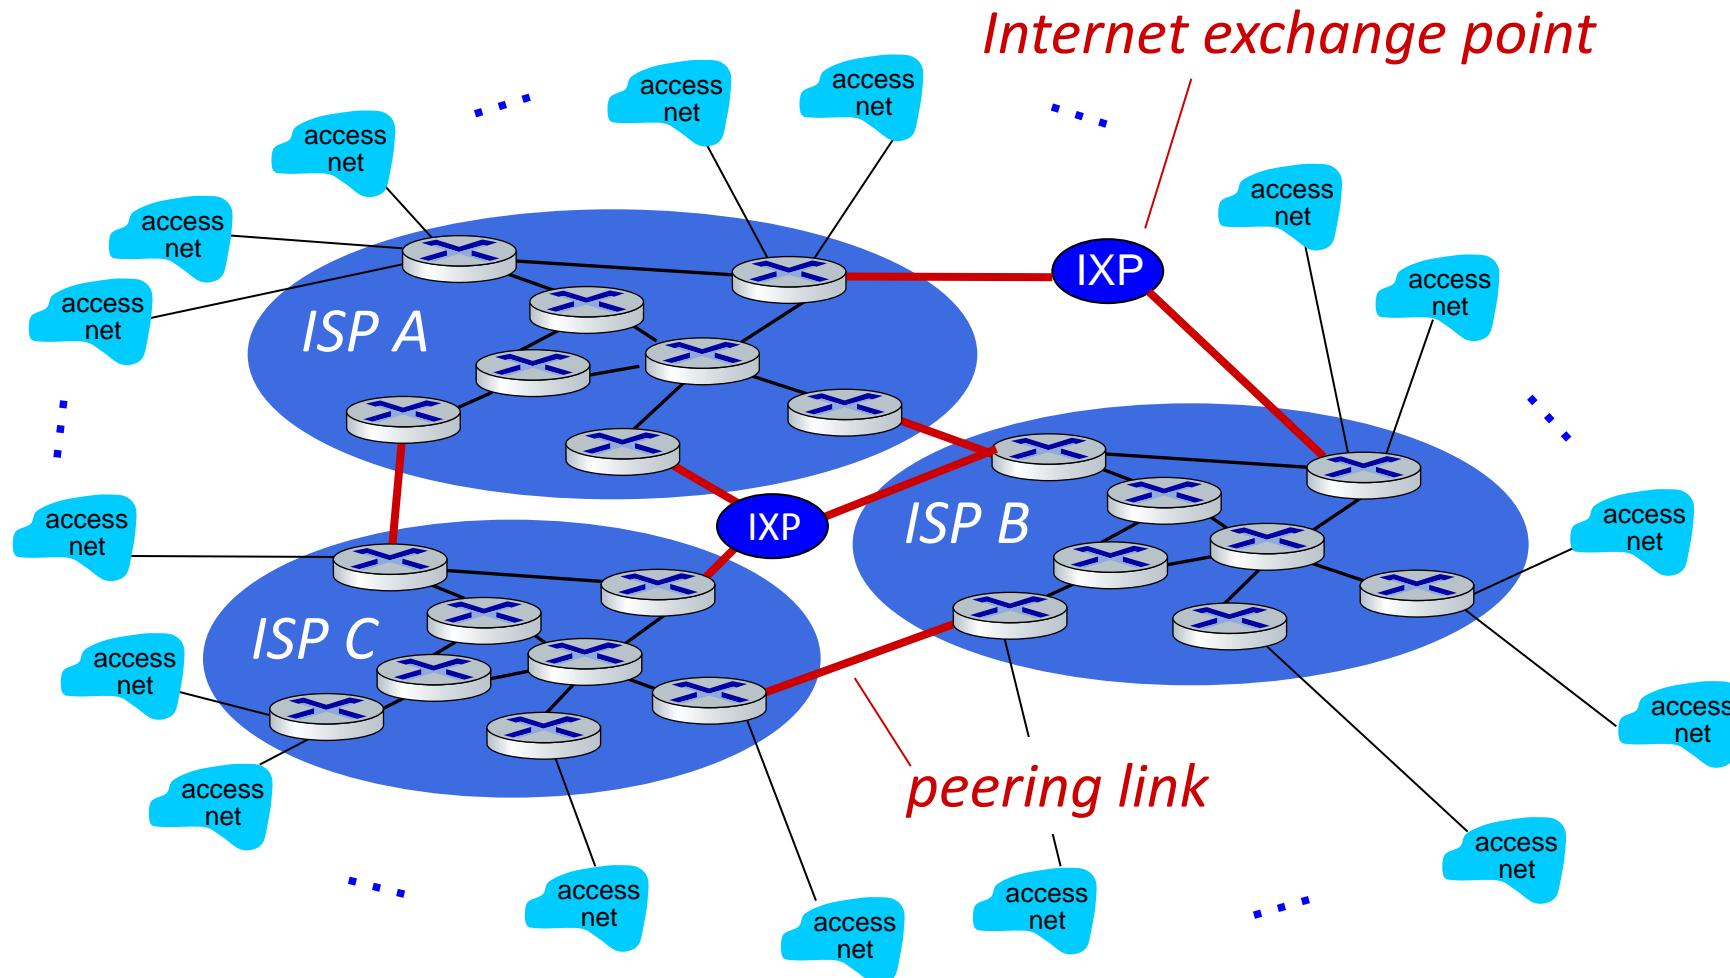
Internet structure: a “network of networks”

But if one global ISP is viable business, there will be competitors



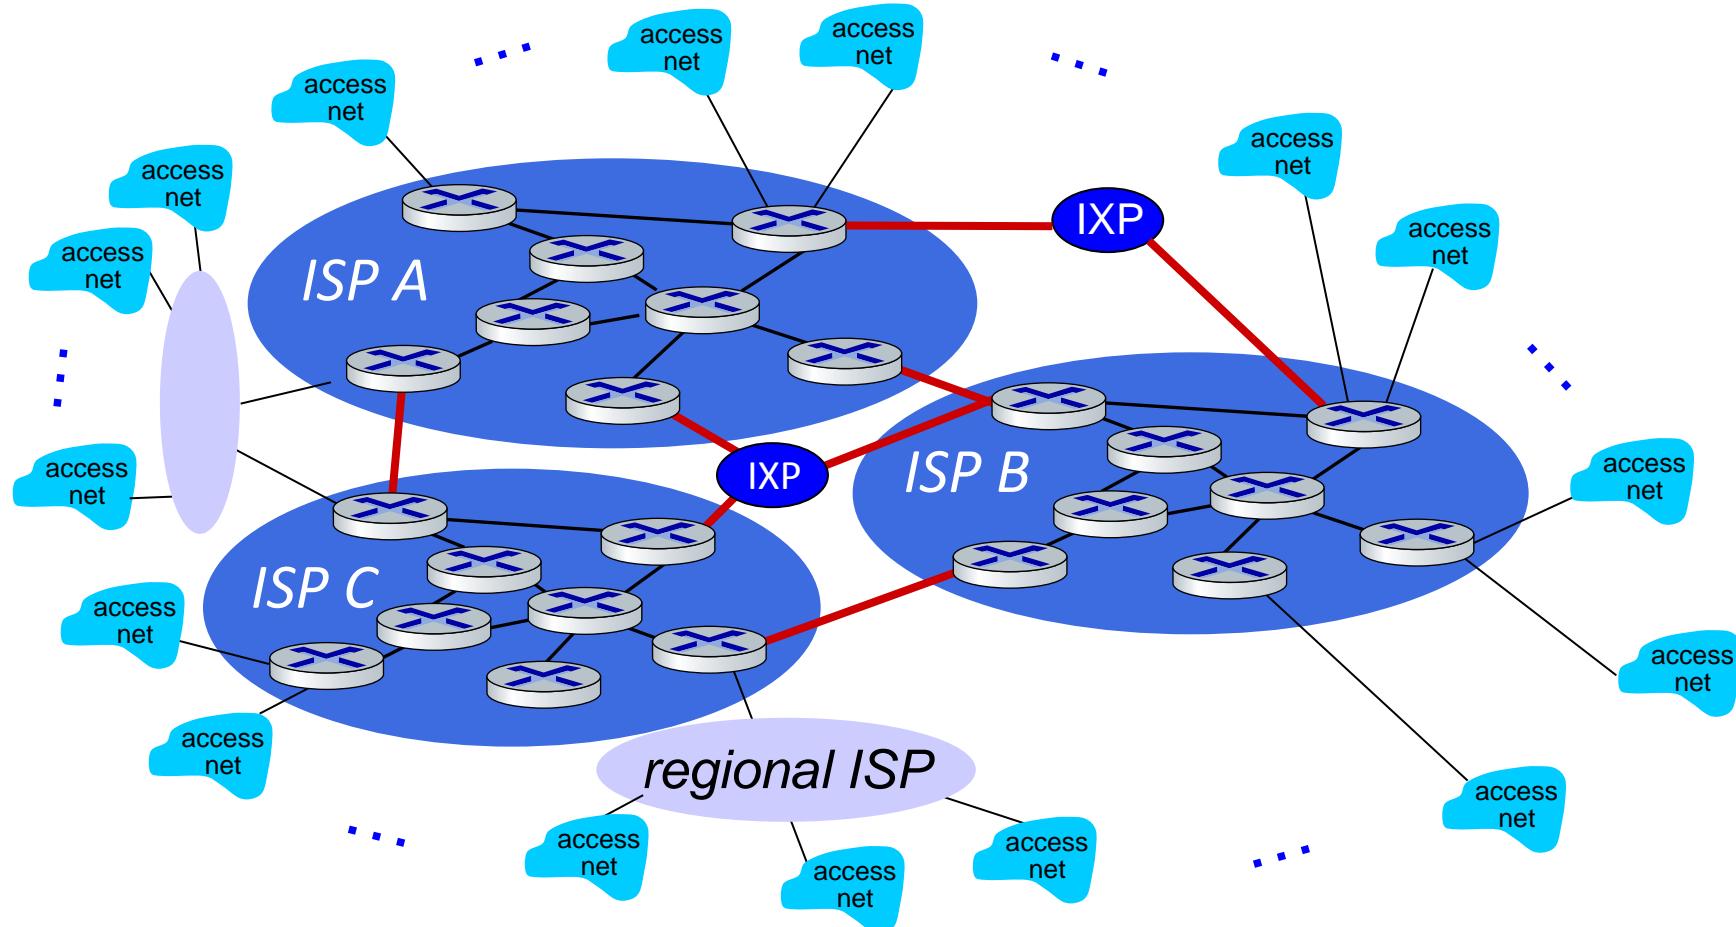
Internet structure: a “network of networks”

But if one global ISP is viable business, there will be competitors ... who will want to be connected



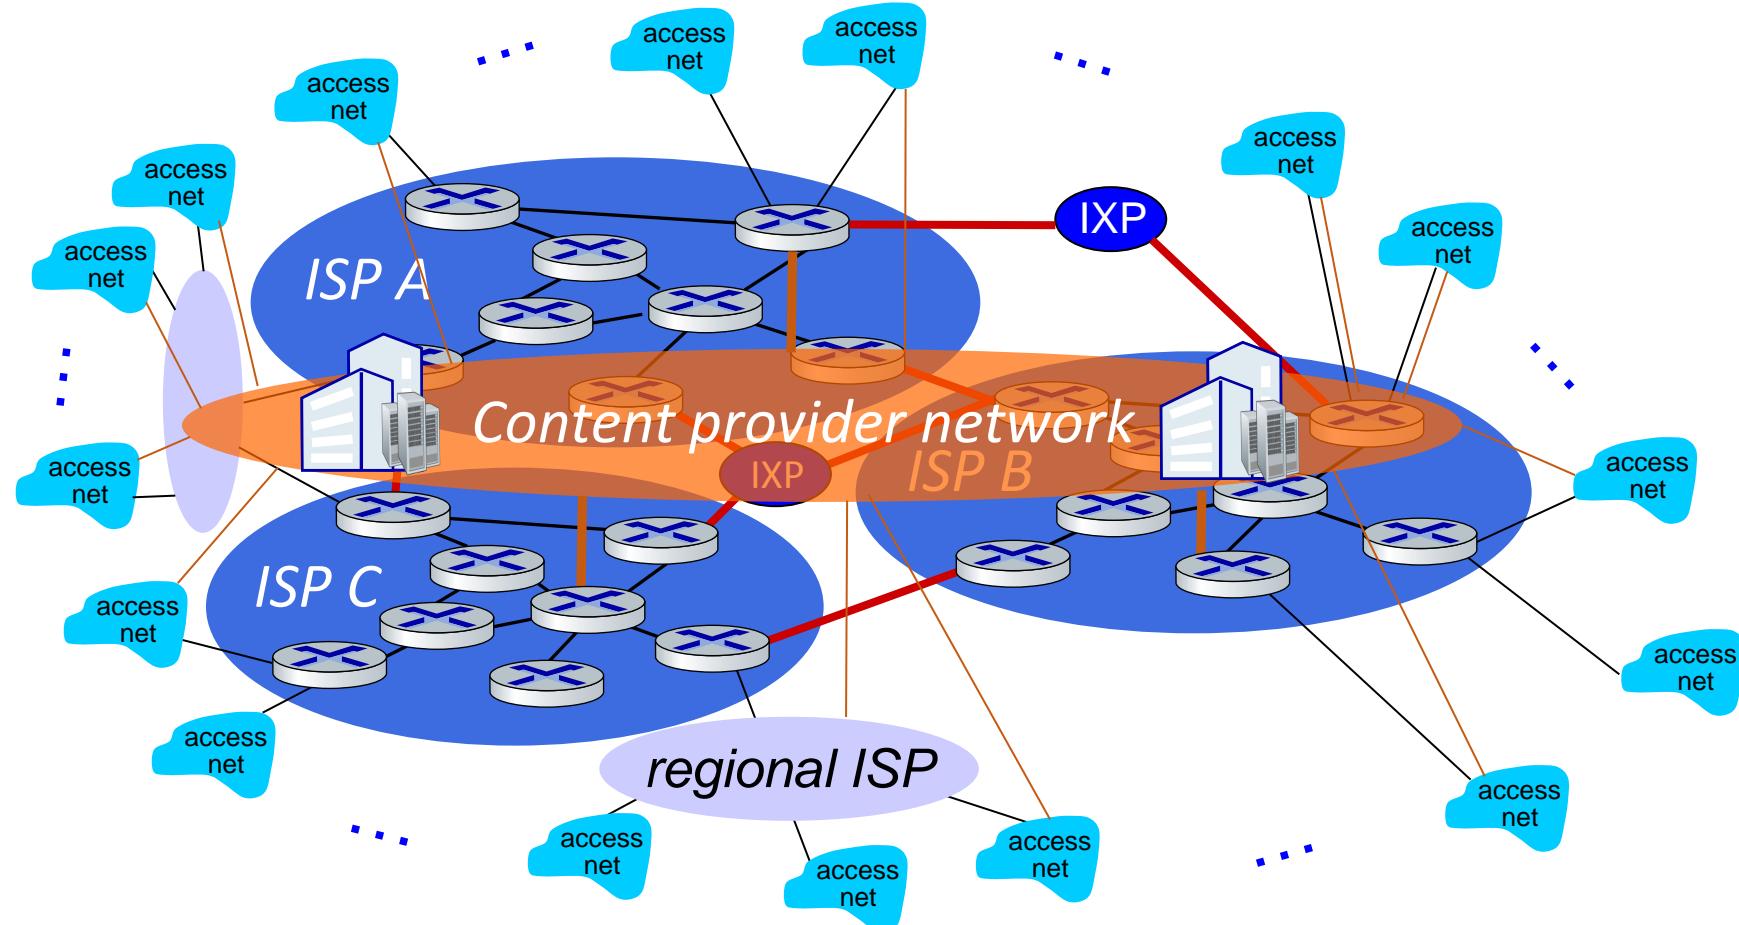
Internet structure: a “network of networks”

... and regional networks may arise to connect access nets to ISPs

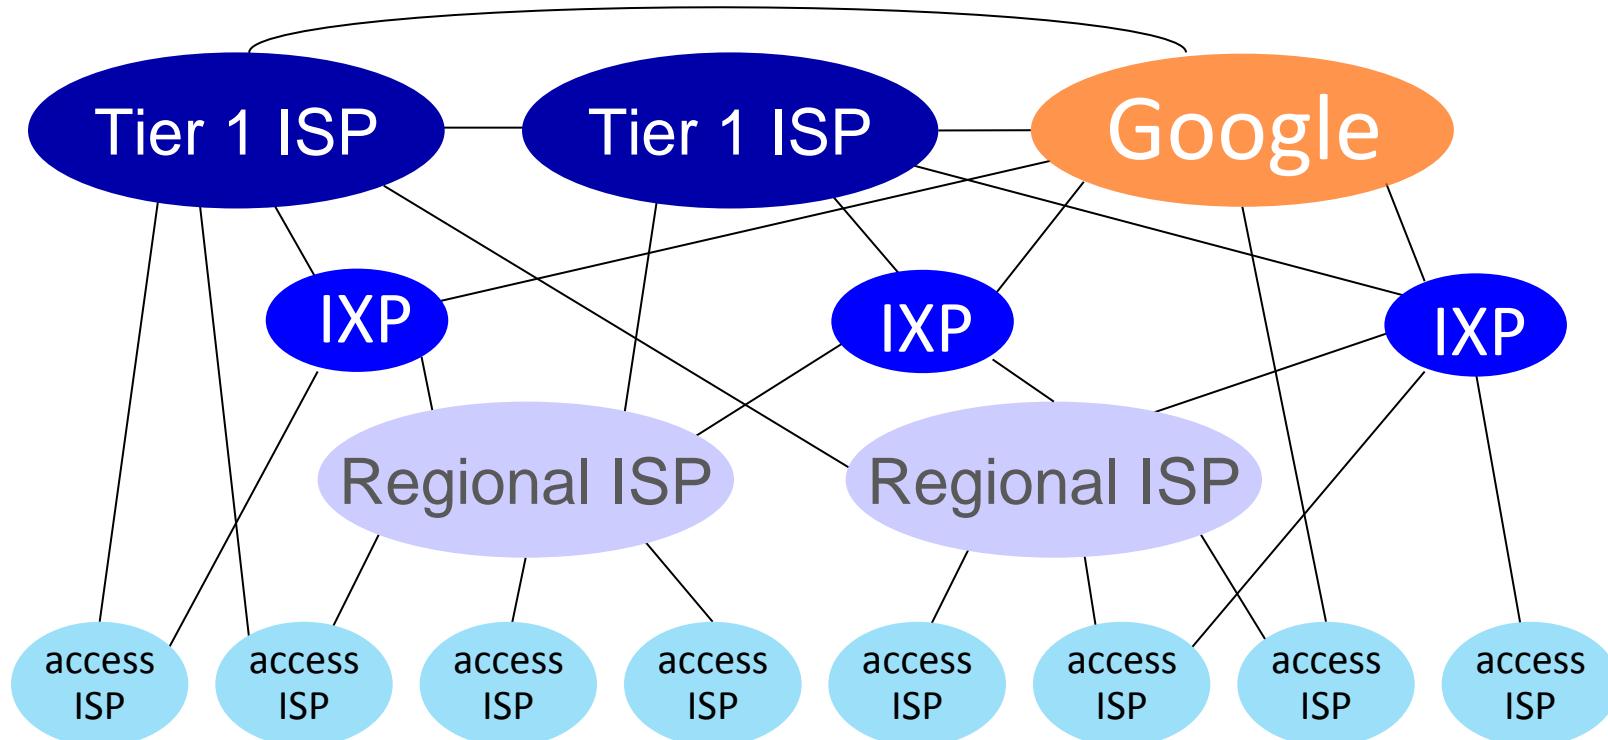


Internet structure: a “network of networks”

... and content provider networks (e.g., Google, Microsoft, Akamai) may run their own network, to bring services, content close to end users



Internet structure: a “network of networks”



At “center”: small # of well-connected large networks

- **“tier-1” commercial ISPs** (e.g., Level 3, Sprint, AT&T, NTT), national & international coverage
- **content provider networks** (e.g., Google, Facebook): private network that connects its data centers to Internet, often bypassing tier-1, regional ISPs

Chapter 1: roadmap

- What *is* the Internet?
- What *is* a protocol?
- Network edge: hosts, access network, physical media
- Network core: packet/circuit switching, internet structure
- **Performance:** loss, delay, throughput
- Security
- Protocol layers, service models
- History



Performance

Delay

Throughput

Loss

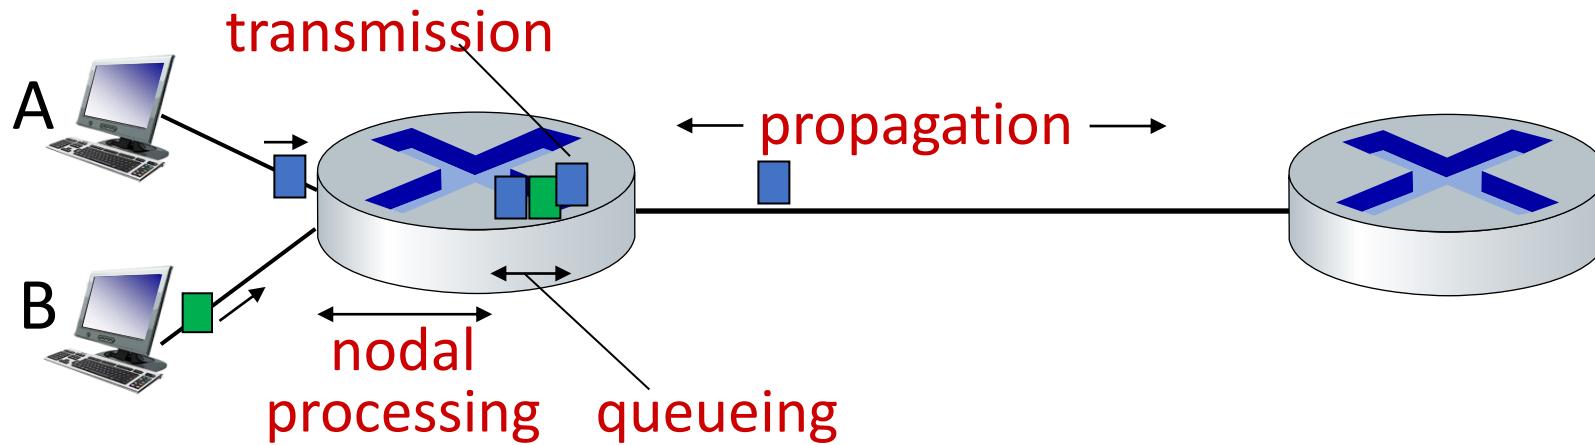
Performance

Delay

Throughput

Loss

Packet delay: four sources



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

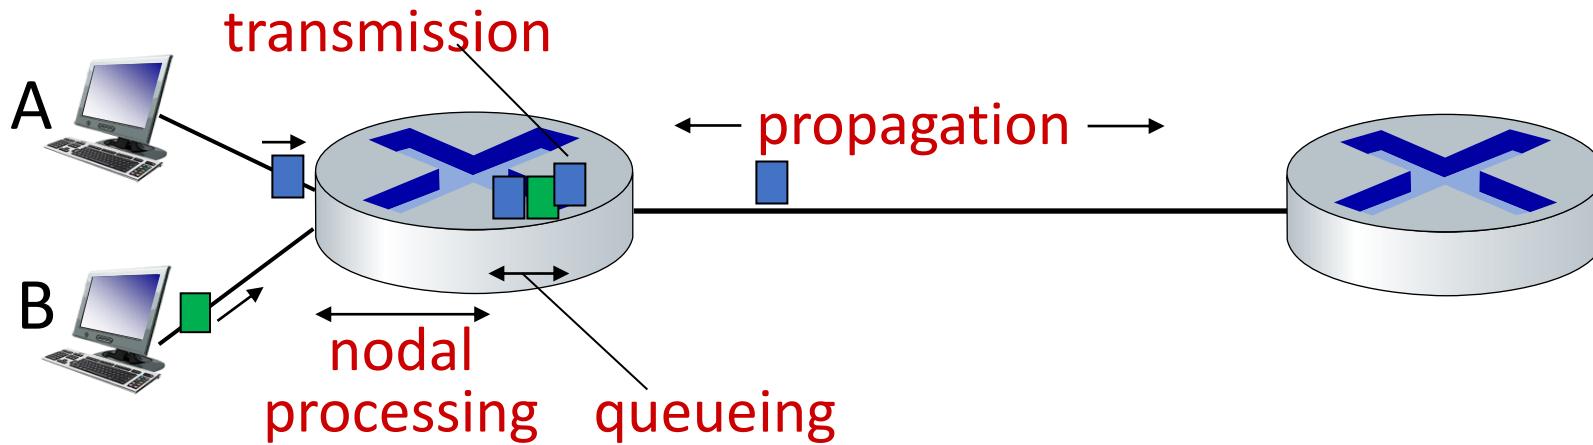
d_{proc} : nodal processing

- check bit errors
- determine output link
- typically < microsecs

d_{queue} : queueing delay

- time waiting at output link for transmission
- depends on congestion level of router

Packet delay: four sources



$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

d_{trans} : transmission delay:

- L : packet length (bits)
- R : link *transmission rate (bps)*

$$\boxed{d_{\text{trans}} = L/R}$$

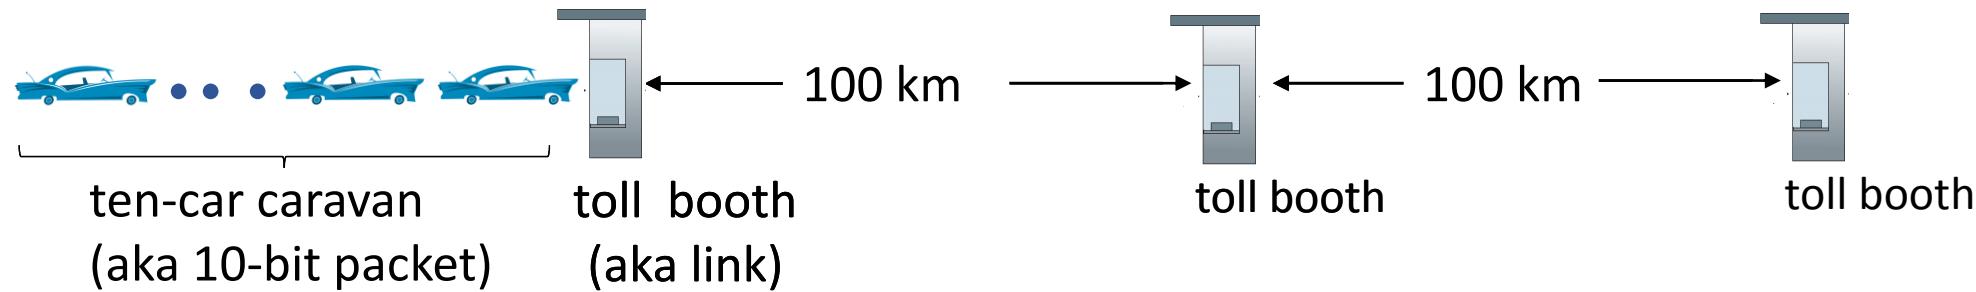
d_{prop} : propagation delay:

- d : length of physical link
- s : propagation speed ($\sim 2 \times 10^8$ m/sec)

$$\boxed{d_{\text{prop}} = d/s}$$

d_{trans} and d_{prop}
very different

Caravan analogy



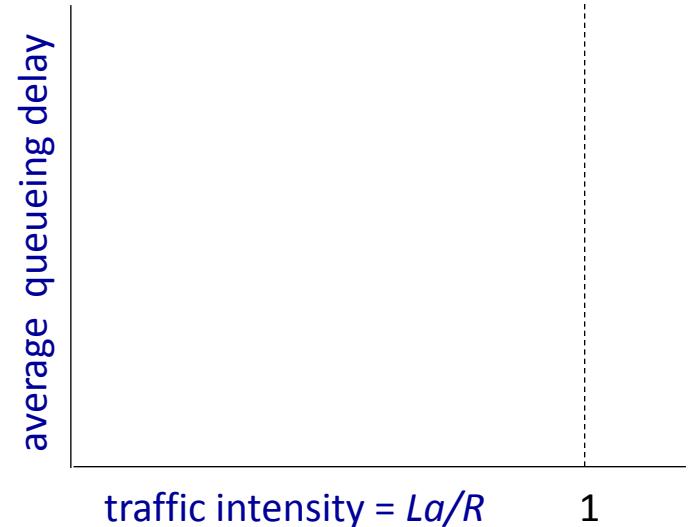
- car ~ bit; caravan ~ packet; toll service ~ link transmission
- toll booth takes 12 sec to service car (bit transmission time)
- “propagate” at 100 km/hr
- **Q: How long until caravan is lined up before 2nd toll booth?**
- time to “push” entire caravan through toll booth onto highway = $12 * 10 = 120$ sec
- time for last car to propagate from 1st to 2nd toll both: $100\text{km}/(100\text{km/hr}) = 1$ hr
- **A: 62 minutes**

Packet queueing delay (revisited)

- a : average packet arrival rate
- L : packet length (bits)
- R : link bandwidth (bit transmission rate)

$$\frac{L \cdot a}{R} : \frac{\text{arrival rate of bits}}{\text{service rate of bits}}$$

“traffic intensity”



- $La/R \sim 0$: avg. queueing delay small
- $La/R > 1$: more “work” arriving is more than can be serviced - average delay infinite!



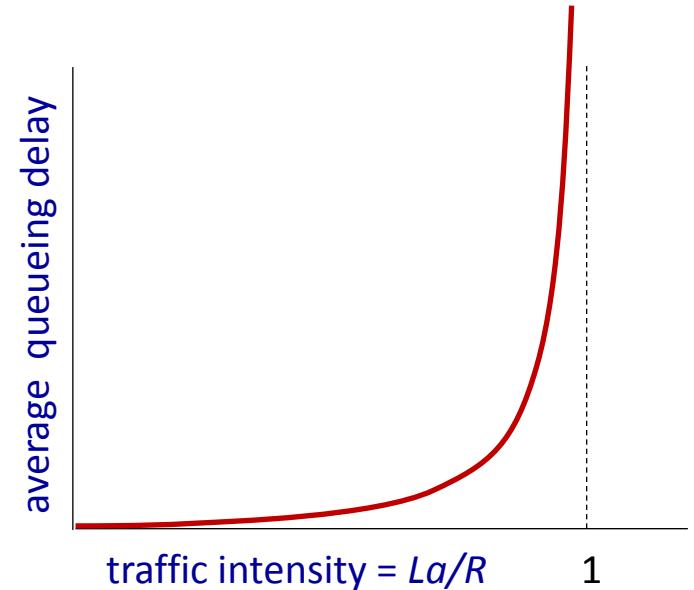
$La/R > 1$

Packet queueing delay (revisited)

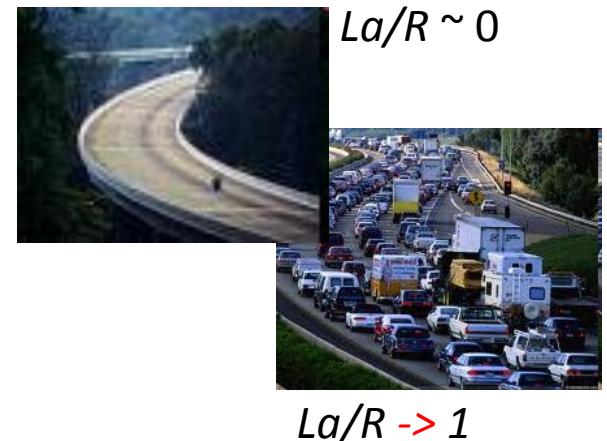
- a : average packet arrival rate
- L : packet length (bits)
- R : link bandwidth (bit transmission rate)

$$\frac{L \cdot a}{R} : \frac{\text{arrival rate of bits}}{\text{service rate of bits}}$$

“traffic intensity”

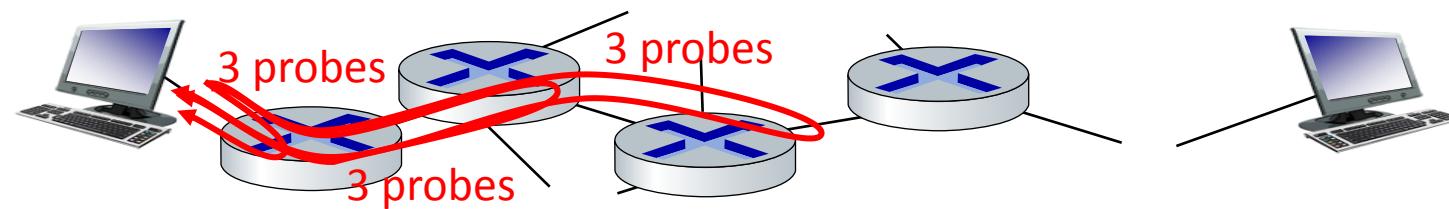


- $La/R \sim 0$: avg. queueing delay small
- $La/R \rightarrow 1$: avg. queueing delay large
- $La/R > 1$: more “work” arriving is more than can be serviced - average delay infinite!



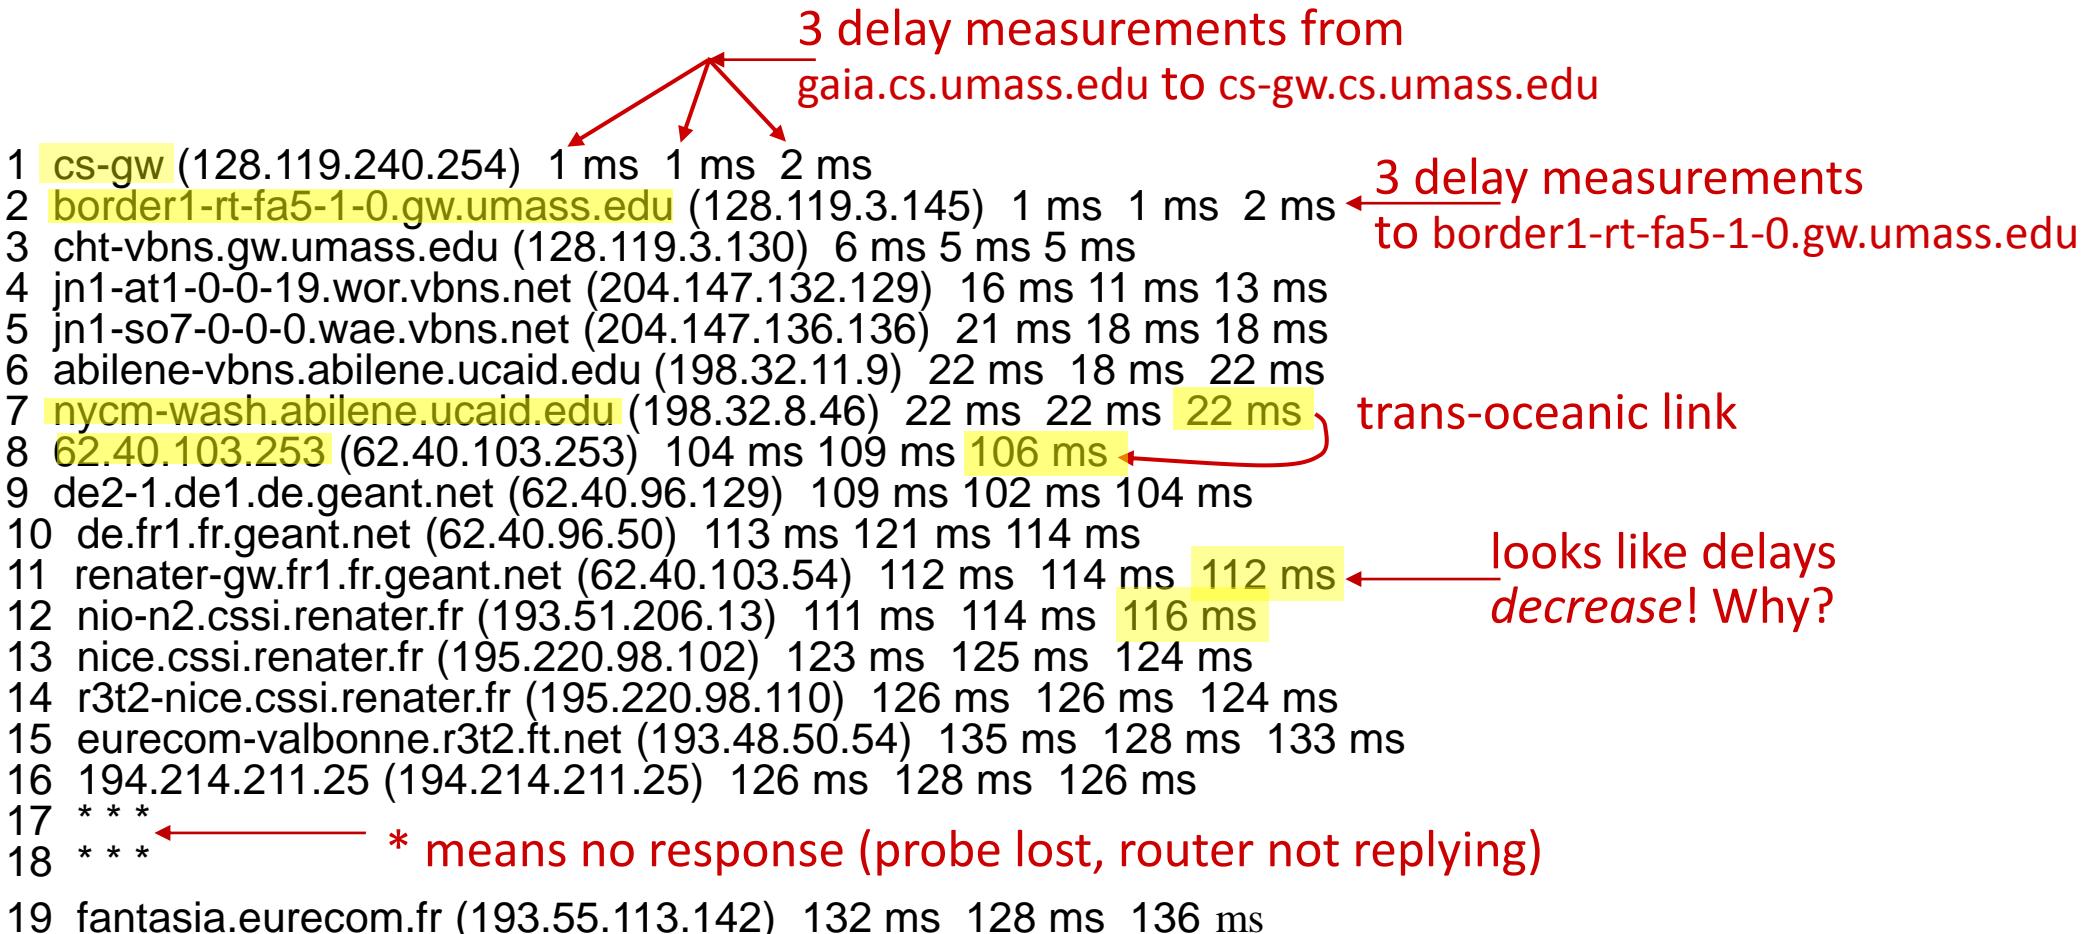
“Real” Internet delays and routes

- what do “real” Internet delay & loss look like?
- **traceroute** program: provides delay measurement from source to router along end-end Internet path towards destination. For all i :
 - sends three packets that will reach router i on path towards destination (with time-to-live field value of i)
 - router i will return packets to sender
 - sender measures time interval between transmission and reply



Real Internet delays and routes

traceroute: gaia.cs.umass.edu to www.eurecom.fr



* Do some traceroutes from exotic countries at www.traceroute.org

Performance

Delay

Throughput

Loss

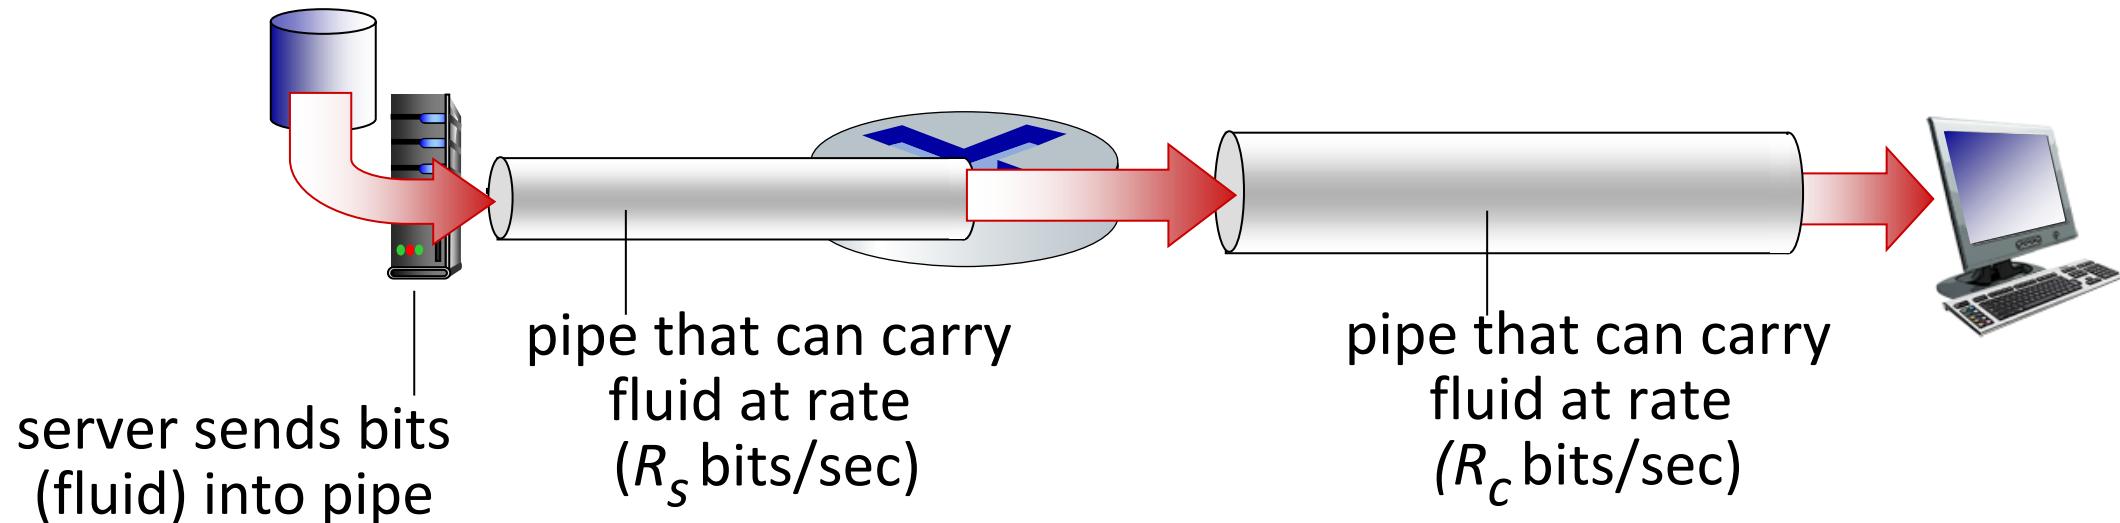
Throughput

- *throughput*: rate (bits/time unit) at which bits are being sent from sender to receiver
 - *instantaneous*: rate at given point in time
 - *average*: rate over longer period of time

$$\text{Throughput} = \frac{\text{Packet size (bit)}}{\text{delay (s)}}$$

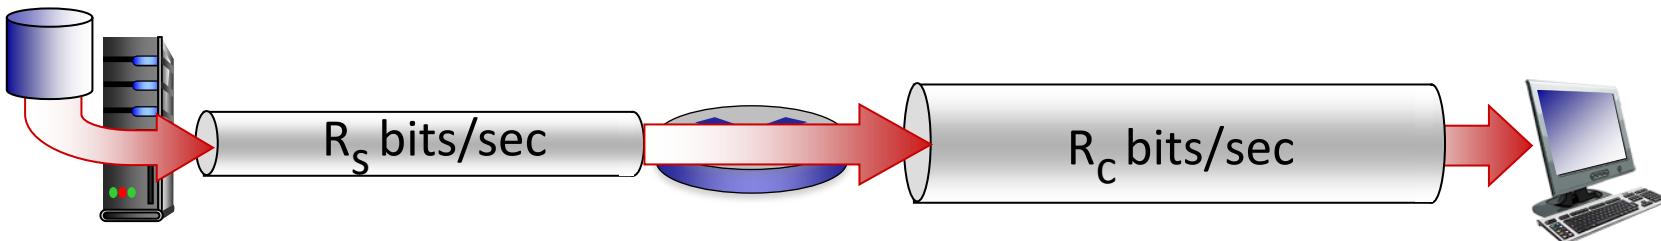
Throughput calculation: Approximate, accurate

Links as pipes

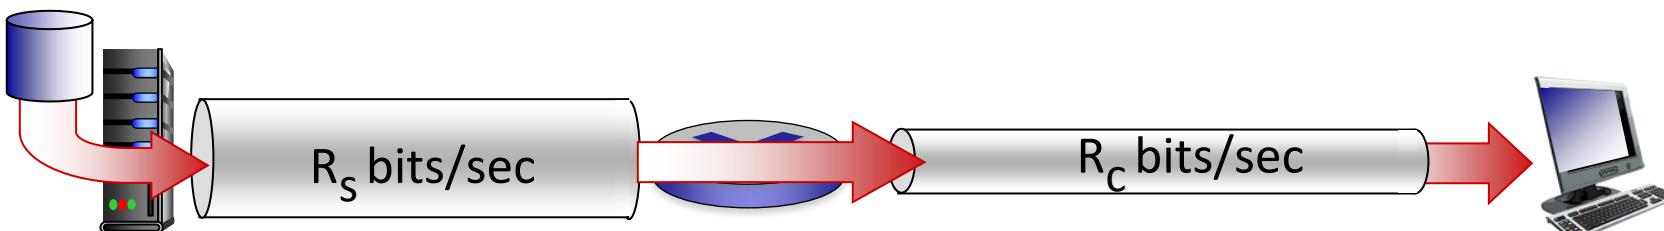


Throughput calculation: Approximate, accurate

$R_s < R_c$ What is average end-end throughput?



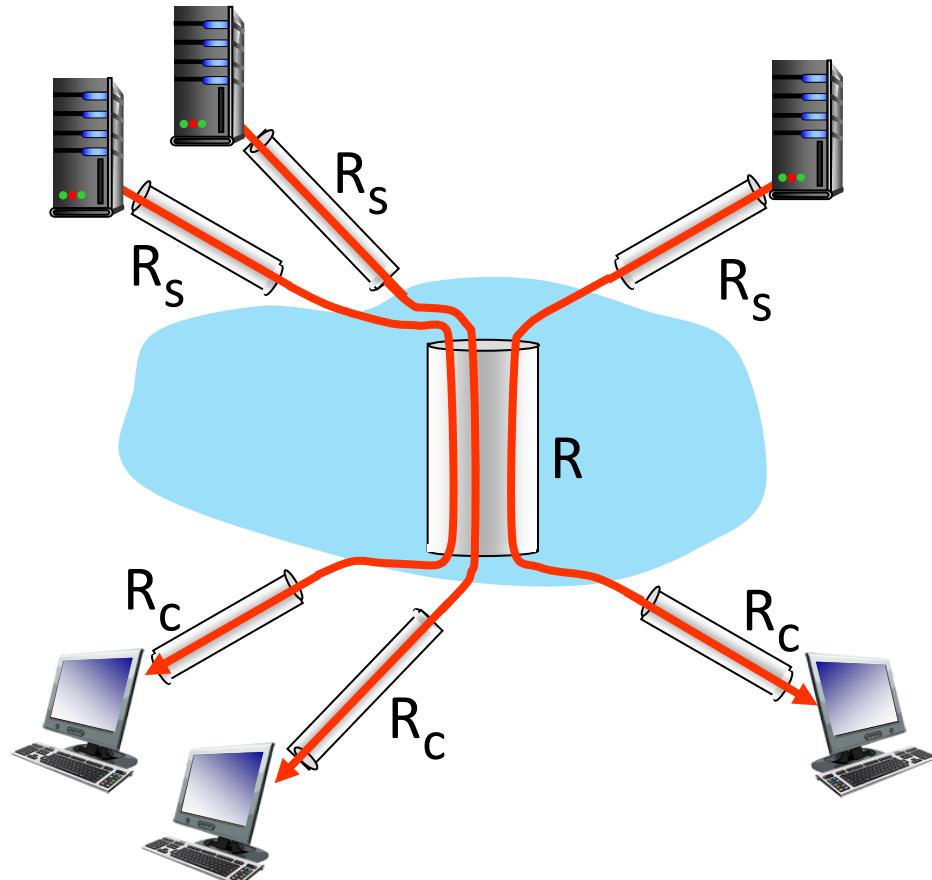
$R_s > R_c$ What is average end-end throughput?



bottleneck link

link on end-end path that constrains end-end throughput

Throughput calculation: Approximate, accurate

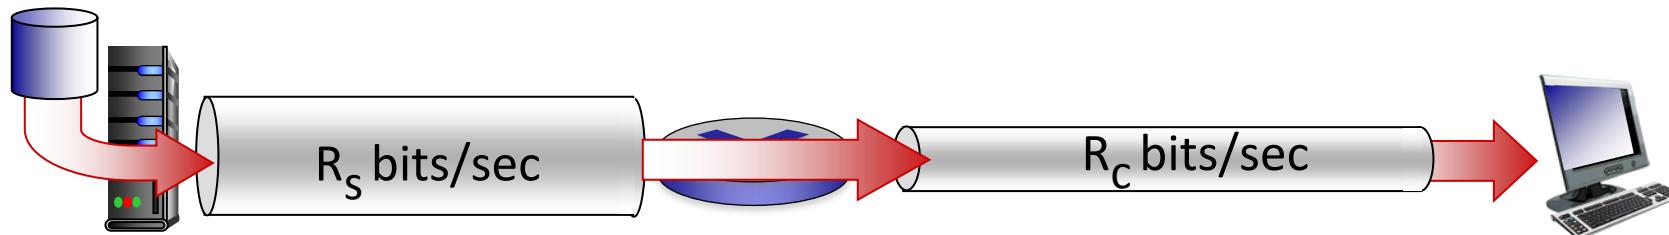


10 connections (fairly) share
backbone bottleneck link R bits/sec

- per-connection end-end throughput: $\min(R_c, R_s, R/10)$
- in practice: R_c or R_s is often bottleneck

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/

Throughput calculation: Approximate, accurate



$$delay = \frac{L}{R_S} + \frac{d_1}{S_1} + d_{proc} + d_{queue} + \frac{L}{R_C} + \frac{d_2}{S_2} \rightarrow Throughput = \frac{L}{\frac{L}{R_S} + \frac{d_1}{S_1} + d_{proc} + d_{queue} + \frac{L}{R_C} + \frac{d_2}{S_2}}$$

Just considering transmission delay $\Rightarrow Throughput \simeq \frac{L}{\frac{L}{R_S} + \frac{L}{R_C}} = \frac{R_S R_C}{R_S + R_C} \simeq \min(R_S, R_C)$

Performance

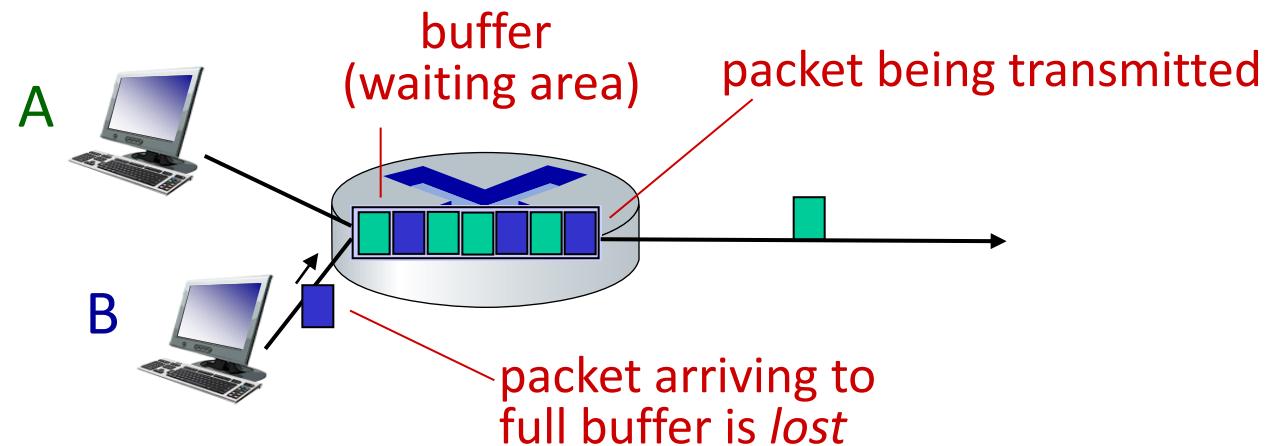
Delay

Throughput

Loss

Packet loss

- queue (aka buffer) preceding link in buffer has finite capacity
- packet arriving to full queue dropped (aka lost)
- lost packet may be retransmitted by previous node, by source end system, or not at all



* Check out the Java applet for an interactive animation (on publisher's website) of queuing and loss

Lines of defense:

- **authentication**: proving you are who you say you are
 - cellular networks provides hardware identity via SIM card; no such hardware assist in traditional Internet
- **confidentiality**: via encryption
- **integrity checks**: digital signatures prevent/detect tampering
- **access restrictions**: password-protected VPNs
- **firewalls**: specialized “middleboxes” in access and core networks:
 - off-by-default: filter incoming packets to restrict senders, receivers, applications
 - detecting/reacting to DOS attacks

... lots more on security (throughout, Chapter 8)

Chapter 1: roadmap

- What *is* the Internet?
- What *is* a protocol?
- Network edge: hosts, access network, physical media
- Network core: packet/circuit switching, internet structure
- Performance: loss, delay, throughput
- Security
- **Protocol layers, service models**
- History



Protocol “layers” and reference models

Networks are complex,
with many “pieces”:

- hosts
- routers
- links of various media
- applications
- protocols
- hardware, software

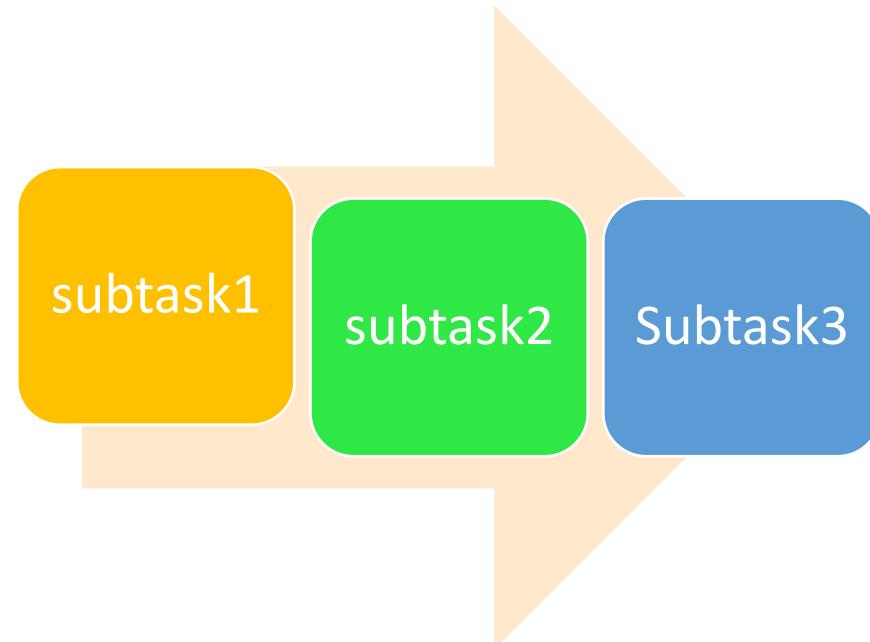
Question: is there any
hope of *organizing*
structure of network?

- and/or our *discussion*
of networks?

Layering is every where

Protocol layering (stack)

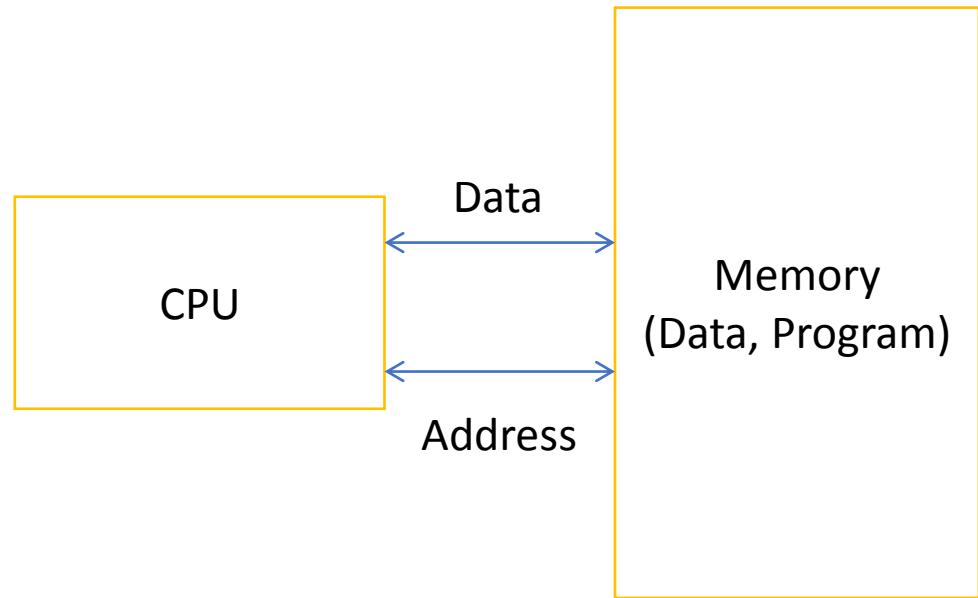
We usually organize our tasks into sub-tasks and define the interfaces



Example 1

We usually organize our tasks into sub-tasks and define the interfaces

- Task: Computation
- Sub-tasks: Hardware design & software design
- Interfaces: ISA



Example 2

We usually organize our tasks into sub-tasks and define the interfaces

- Task: Programming
- Sub-tasks: Procedures
- Interfaces:
Arguments/return values

```
#include <stdio.h>
int main(){
    printf("hello world");
}
```

Example 3

We usually organize our tasks into sub-tasks and define the interfaces

- Task: travel
- Sub-tasks: agency, check-in, gates, ...
- Interfaces: Passport, ticket, serial number, ...

ticket (purchase)
baggage (check)
gates (load)
runway takeoff
airplane routing

Example 3

We usually organize our tasks into sub-tasks and define the interfaces

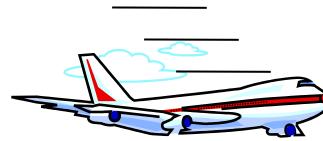
- Task: travel
- Sub-tasks: agency, check-in, gates, ...
- Interfaces: Passport, ticket, serial number, ...

ticket (purchase)
baggage (check)
gates (load)
runway takeoff
airplane routing

The Computer network's goal is similar to human's travel:

- >> Transporting packets from source application to the destination application
- >> Transporting humans from home airport to destination airport

So let's look deeper



— *end-to-end transfer of person plus baggage* —→

ticket (purchase)
baggage (check)
gates (load)
runway takeoff
airplane routing

ticket (complain)
baggage (claim)
gates (unload)
runway landing
airplane routing

airplane routing

How would you *define/discuss* the *system* of airline travel?

- a series of steps, involving many services

Example: organization of air travel



layers: each layer implements a service

- via its own internal-layer actions
- relying on services provided by layer below

Why layering?

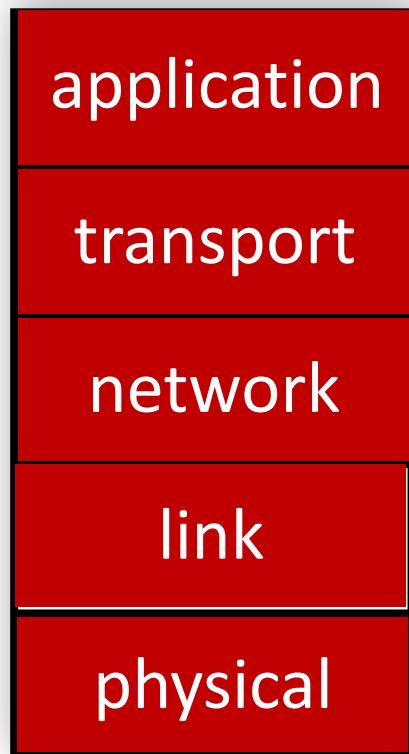
Approach to designing/discussing complex systems:

- explicit structure allows identification, relationship of system's pieces
 - layered *reference model* for discussion
- modularization eases maintenance, updating of system
 - change in layer's service *implementation*: transparent to rest of system
 - e.g., change in gate procedure doesn't affect rest of system

Layering is every where

Protocol layering (stack)

Layered Internet protocol stack

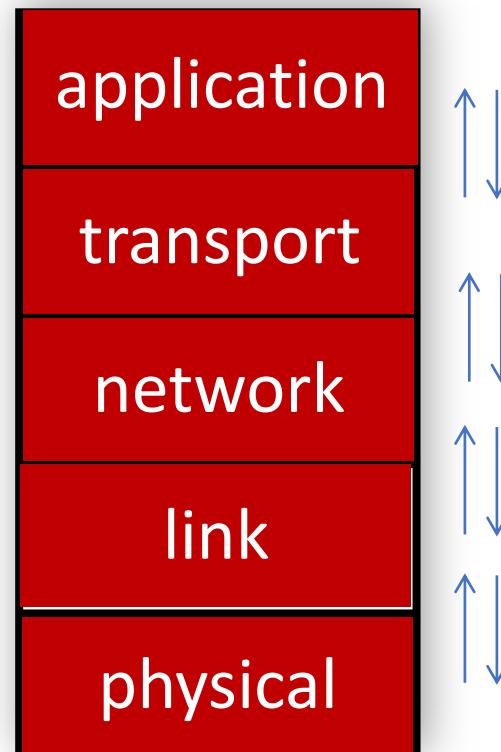


Layered Internet protocol stack

Stack:

Each layer connects to Just above and below layers

*Gets service from the below
Gives service to the above*

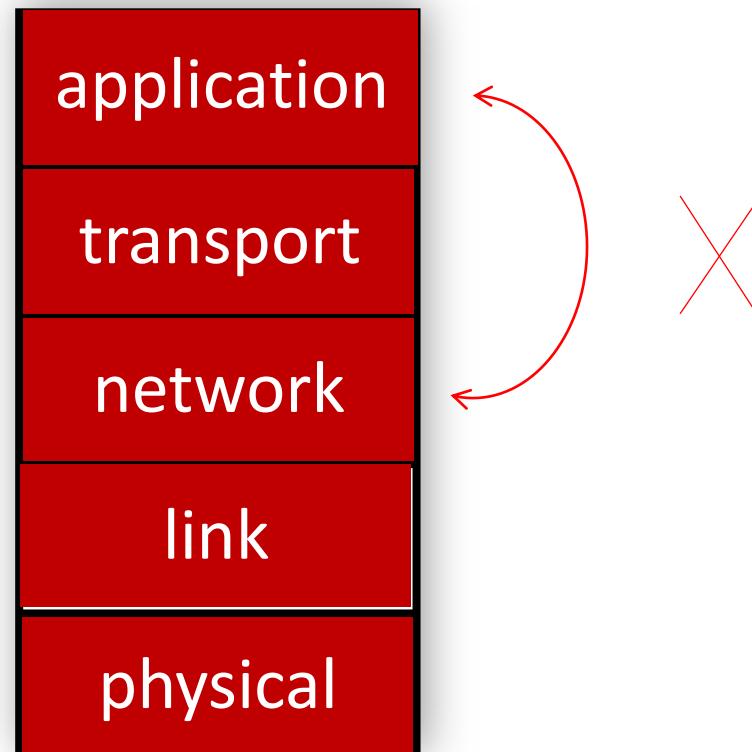


Layered Internet protocol stack

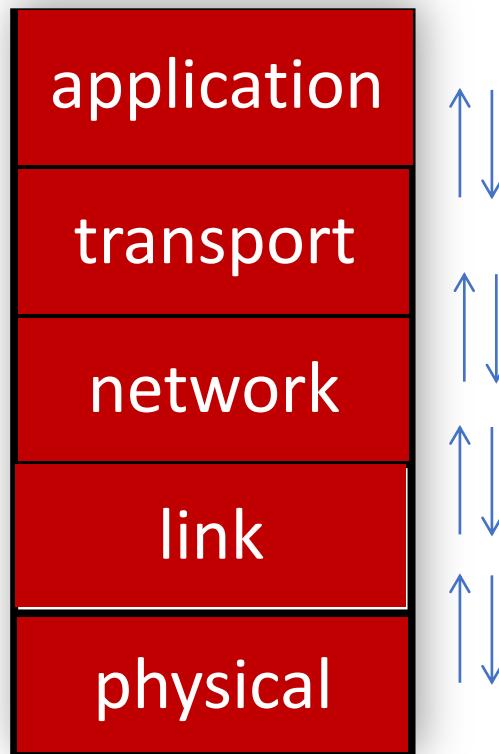
Stack:

Each layer connects to Just above and below layers

*Gets service from the below
Gives service to the above*



Layered Internet protocol stack

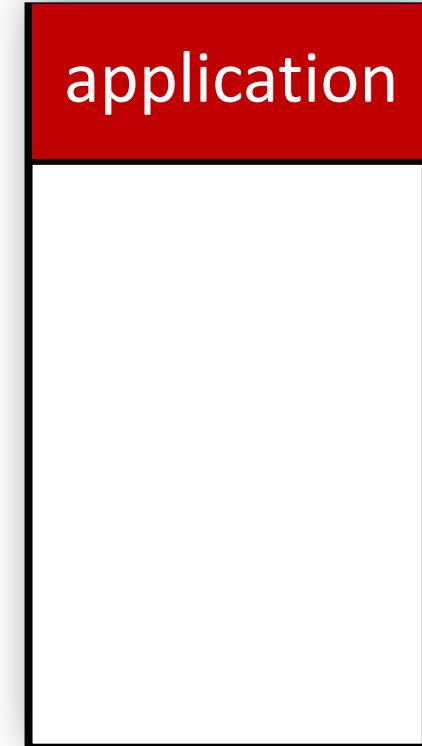


Each layer does his jobs
by adding its header to
the packet
(Encapsulation)

Layered Internet protocol stack

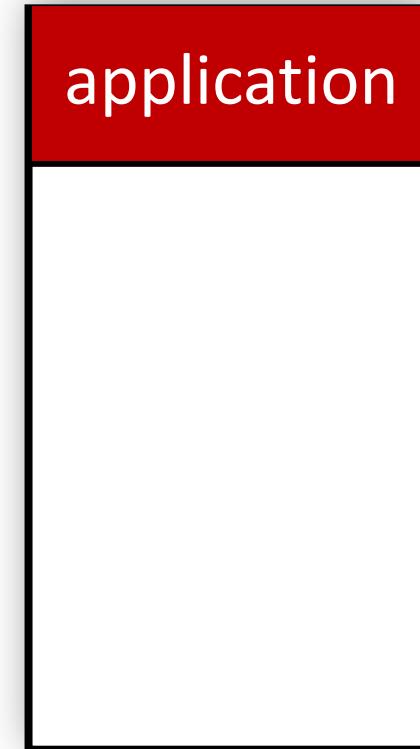
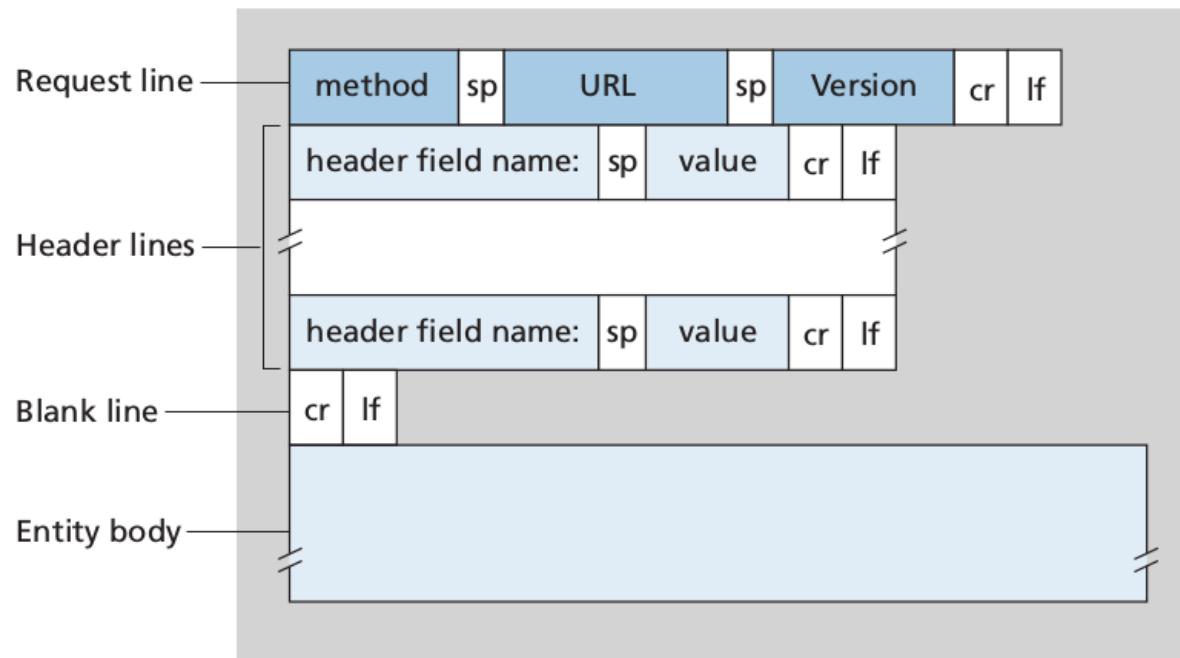
- *application*: supporting network applications
 - **HTTP**, IMAP, SMTP, DNS

HTTP Demo using telnet



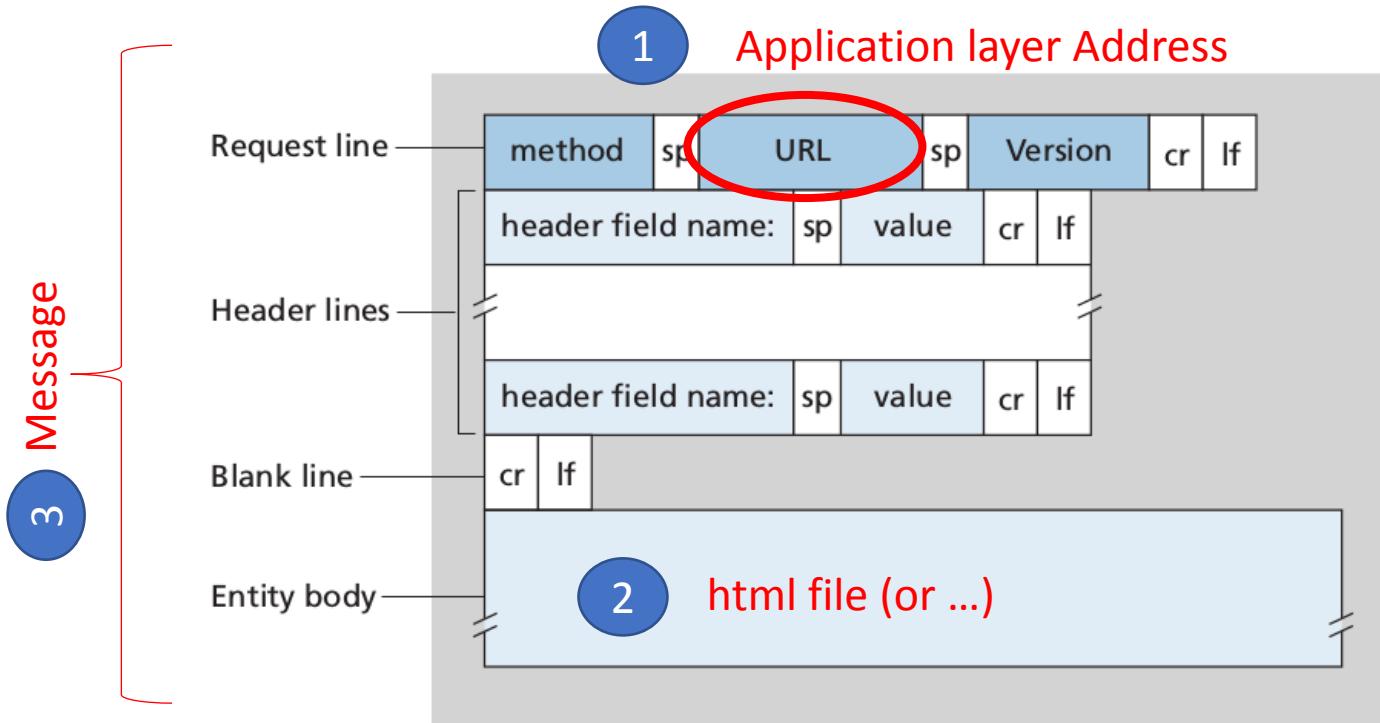
Layered Internet protocol stack

- *application*: supporting network applications
 - **HTTP, IMAP, SMTP, DNS**



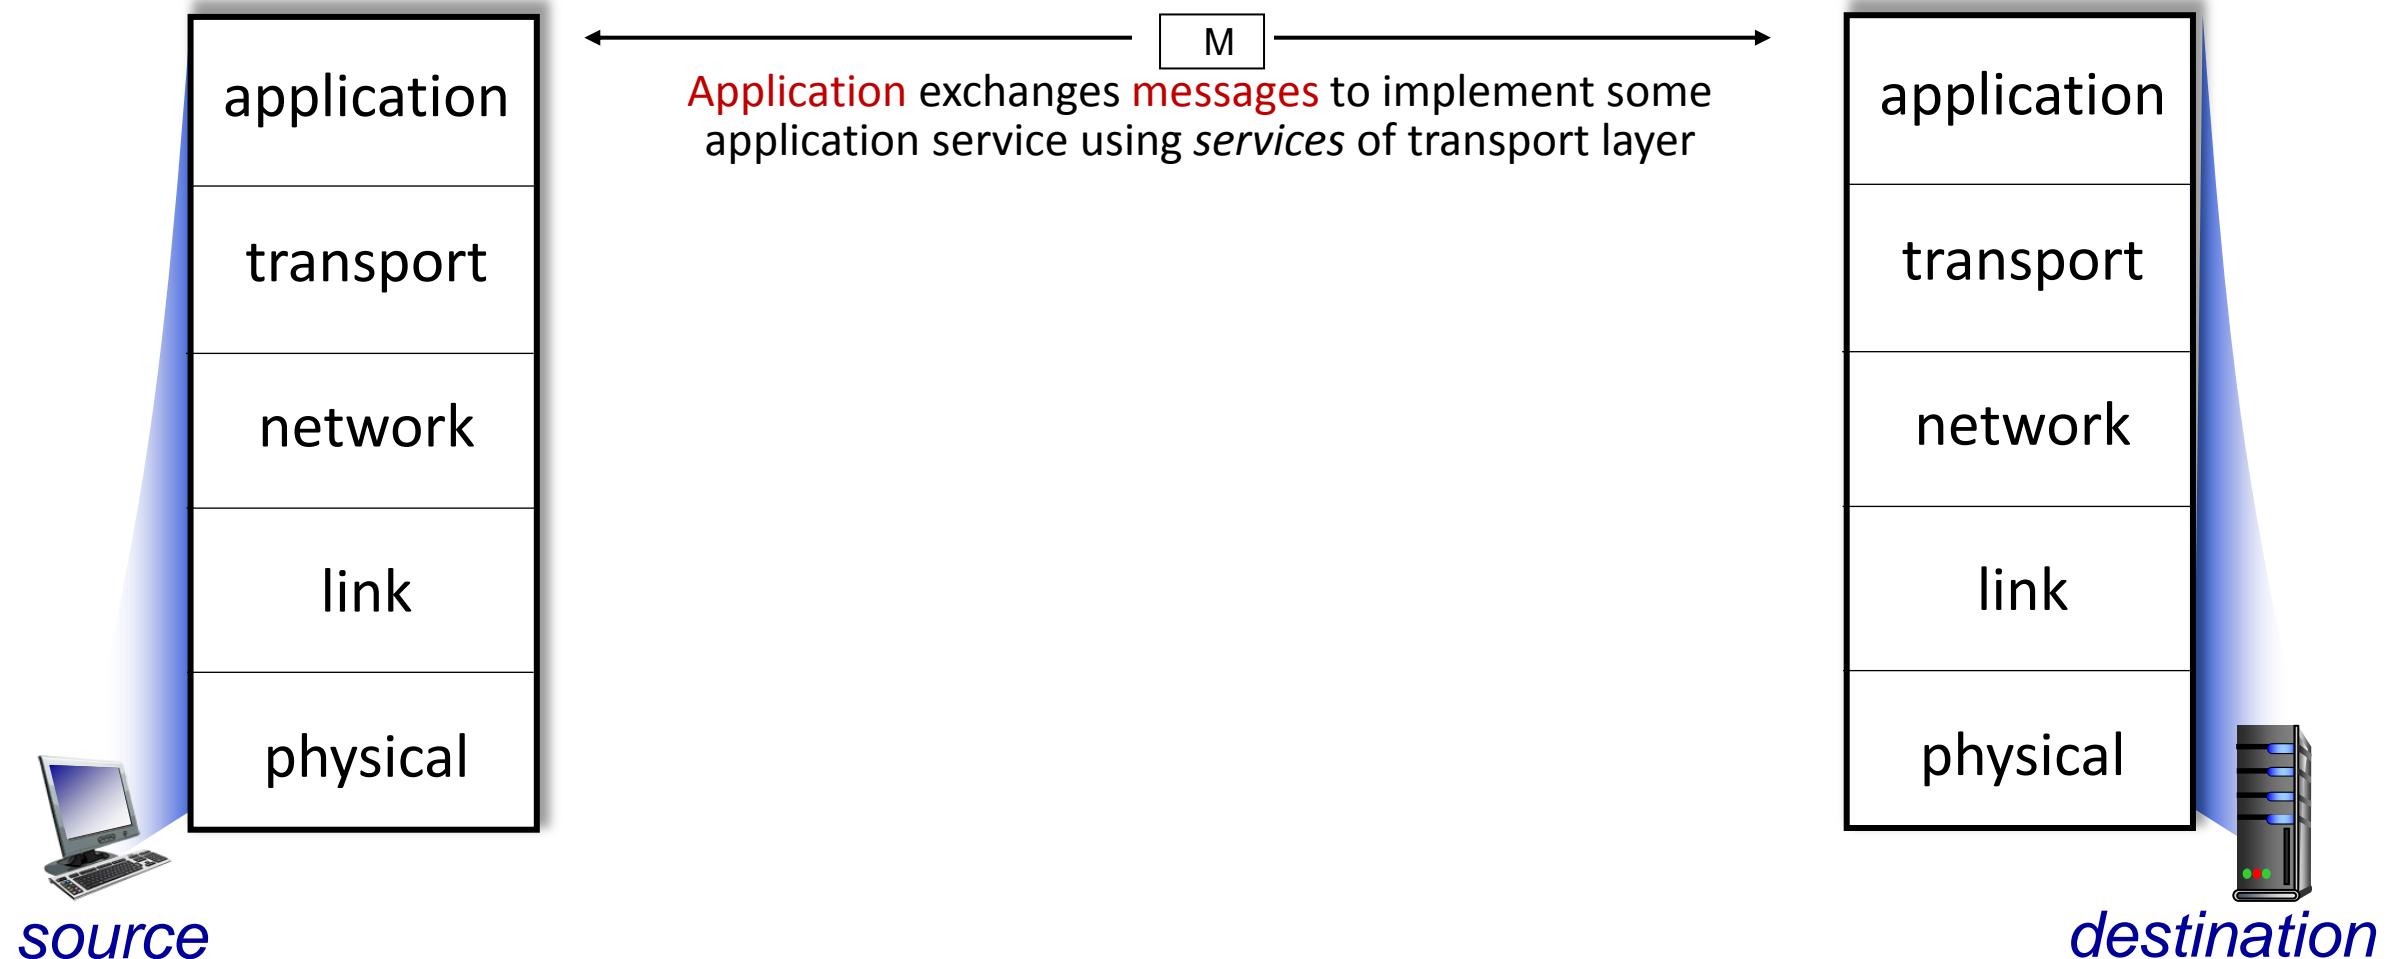
Layered Internet protocol stack

- *application*: supporting network applications
 - **HTTP, IMAP, SMTP, DNS**



4 Only in end-systems

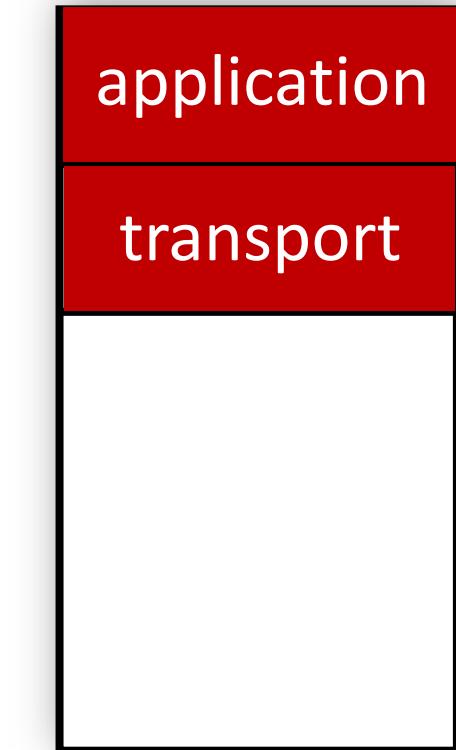
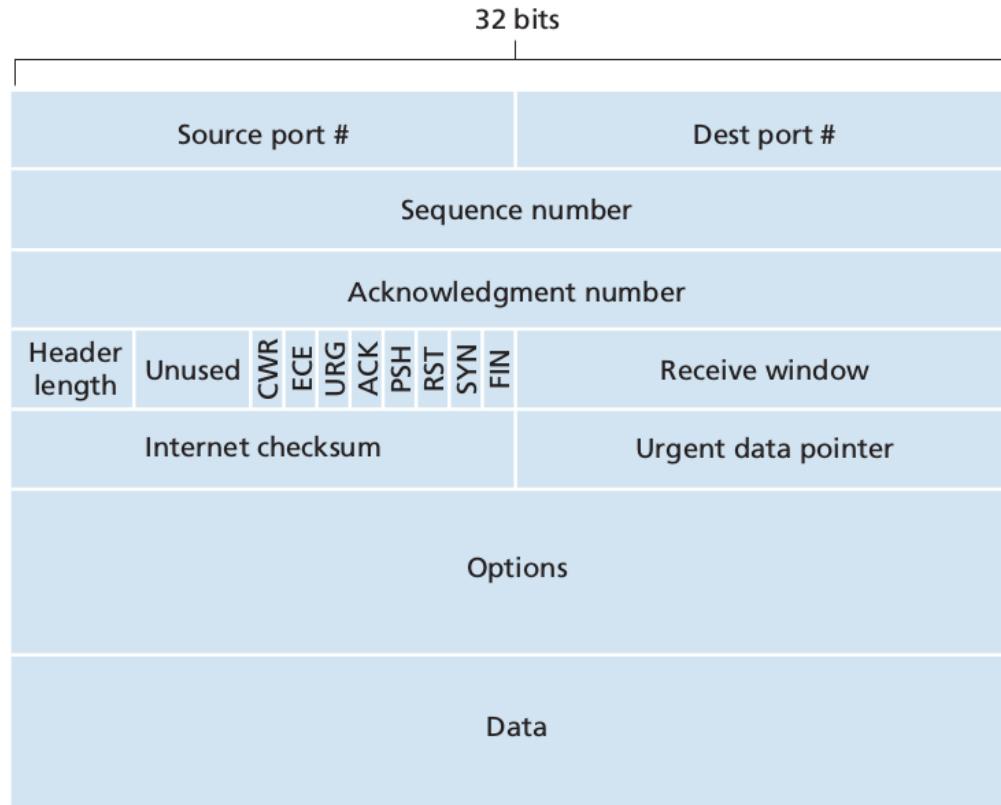
Encapsulation



Layered Internet protocol stack

- *transport*: process-process data transfer

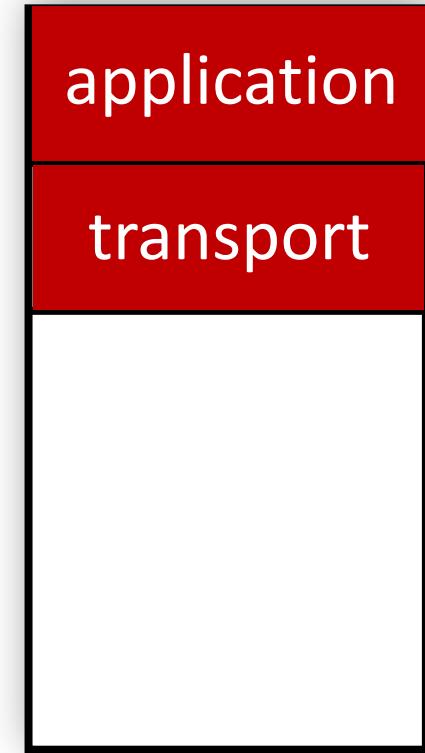
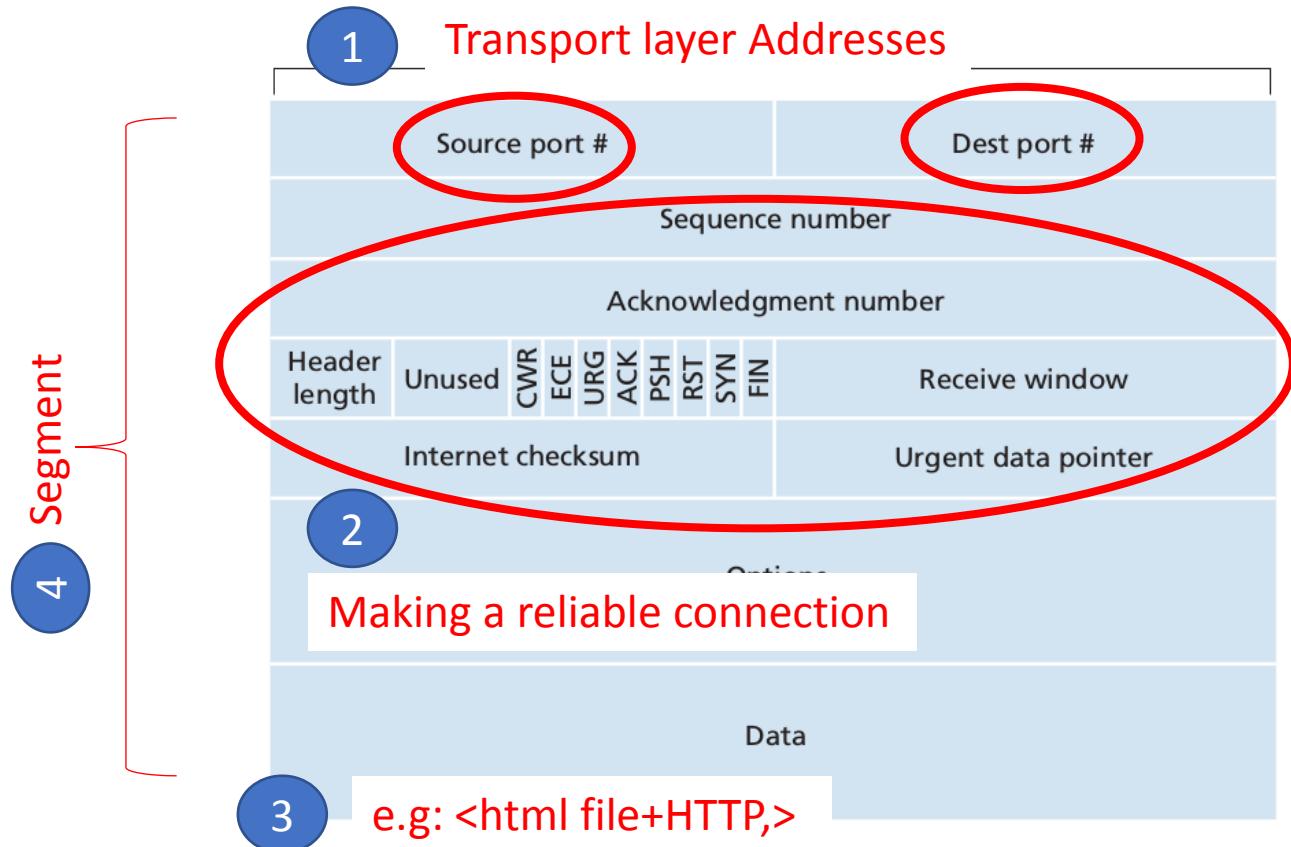
- TCP, UDP



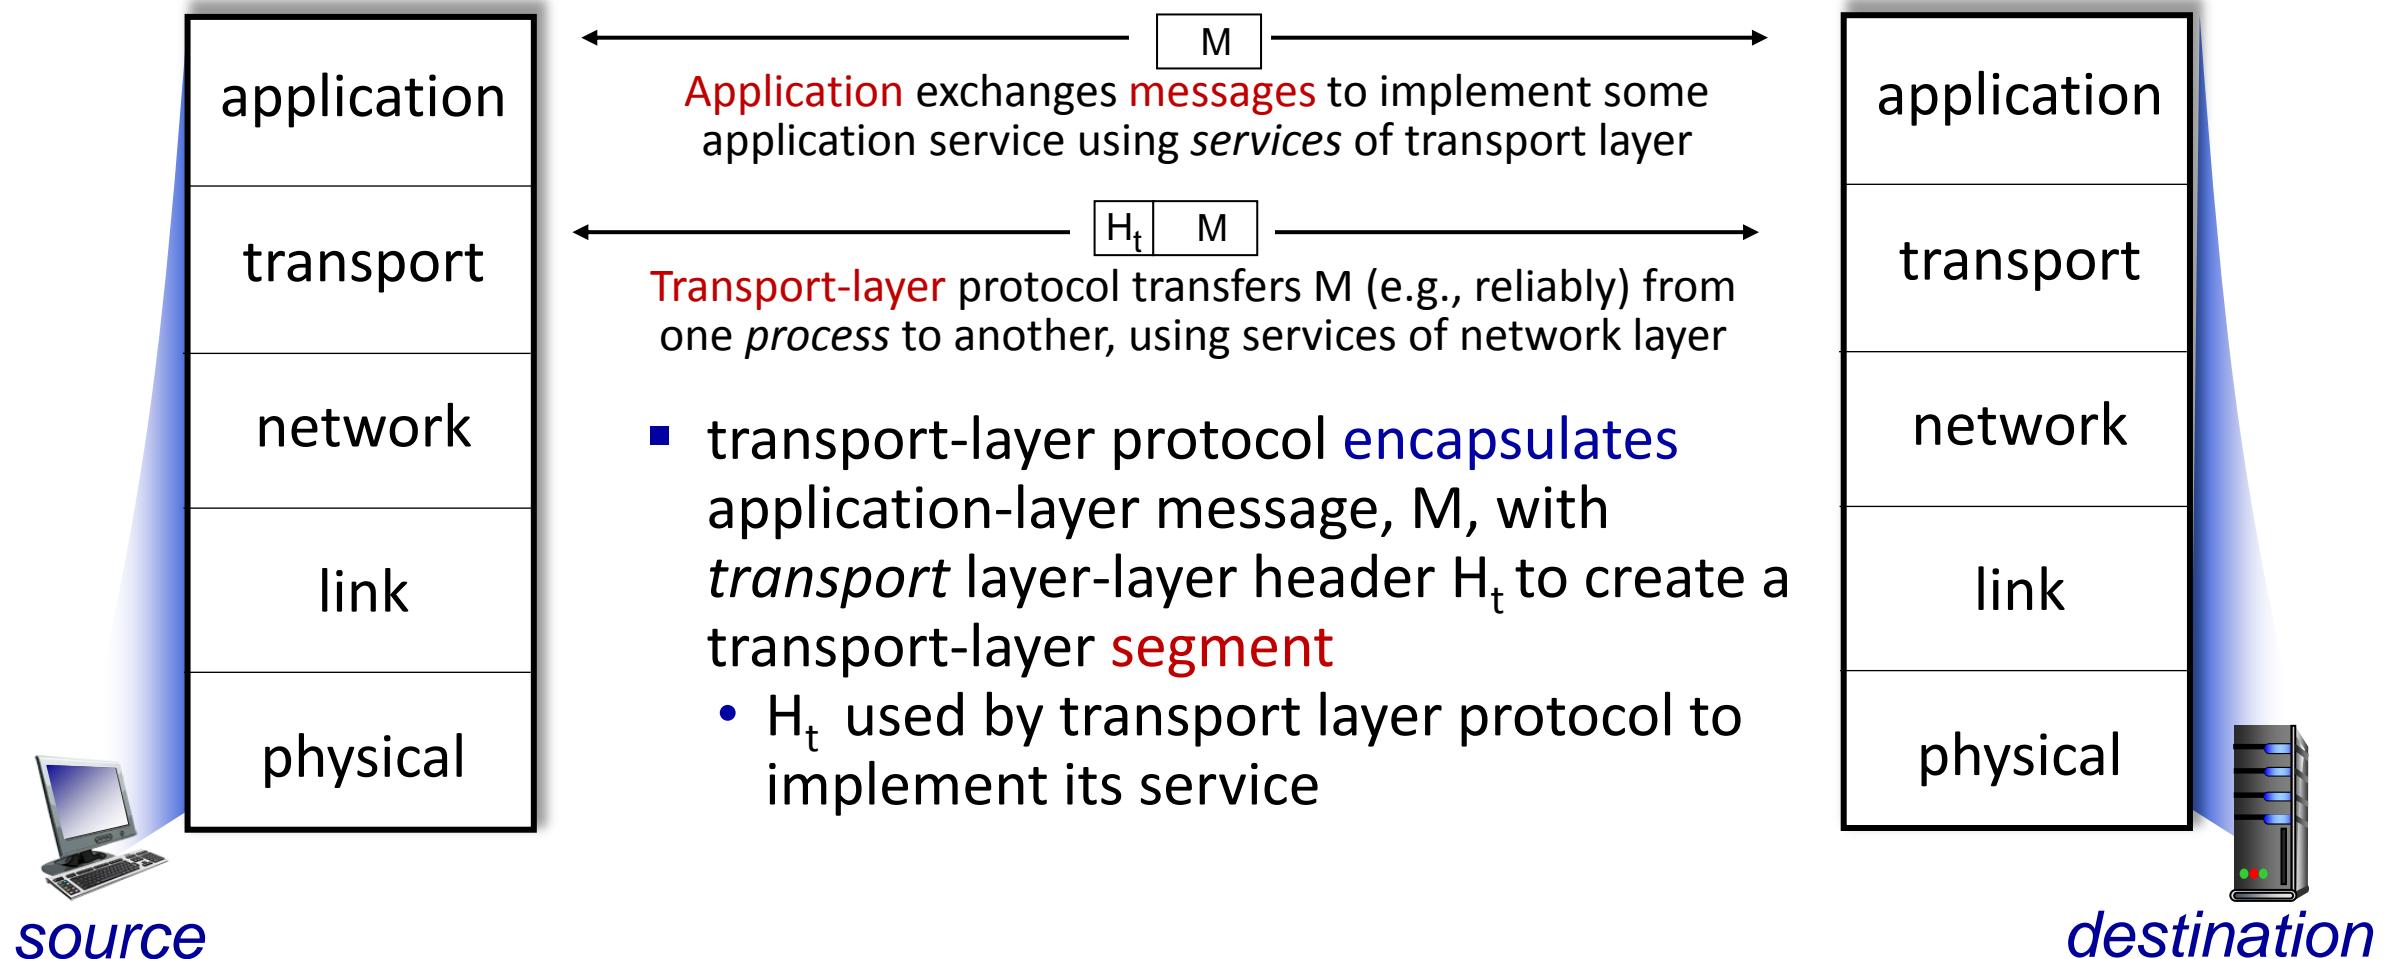
Layered Internet protocol stack

- *transport*: process-process data transfer

- TCP, UDP

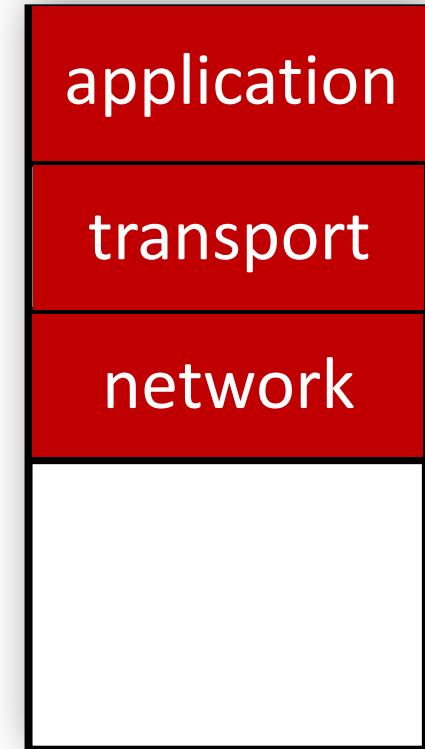
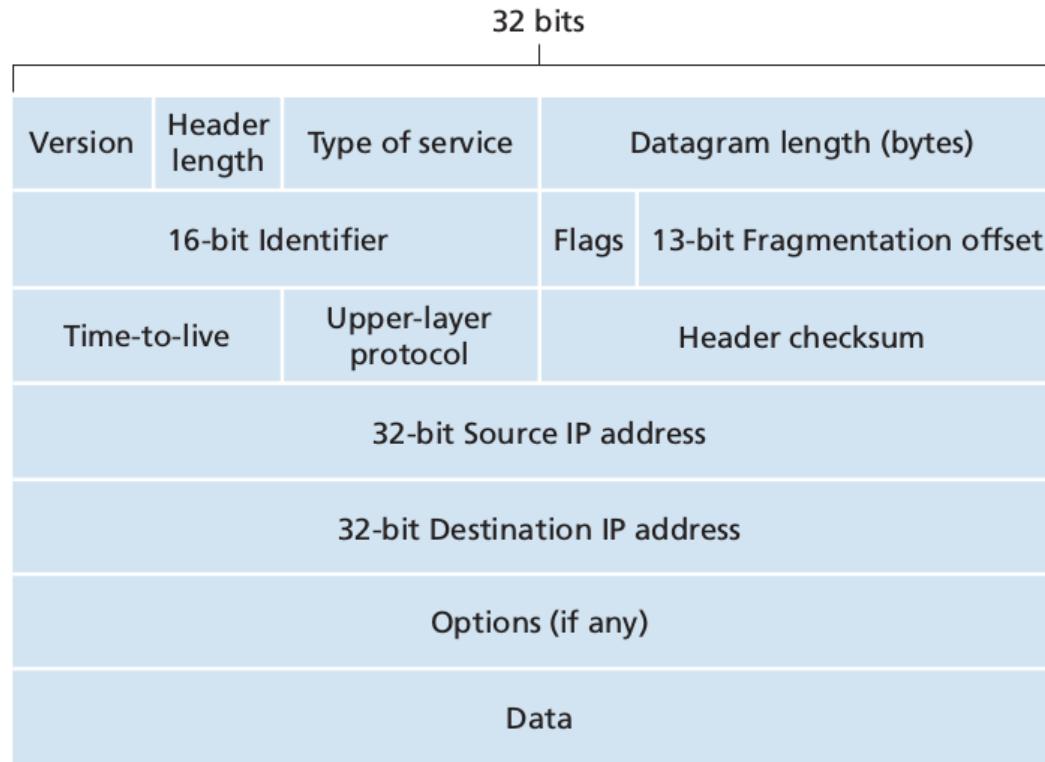


Encapsulation



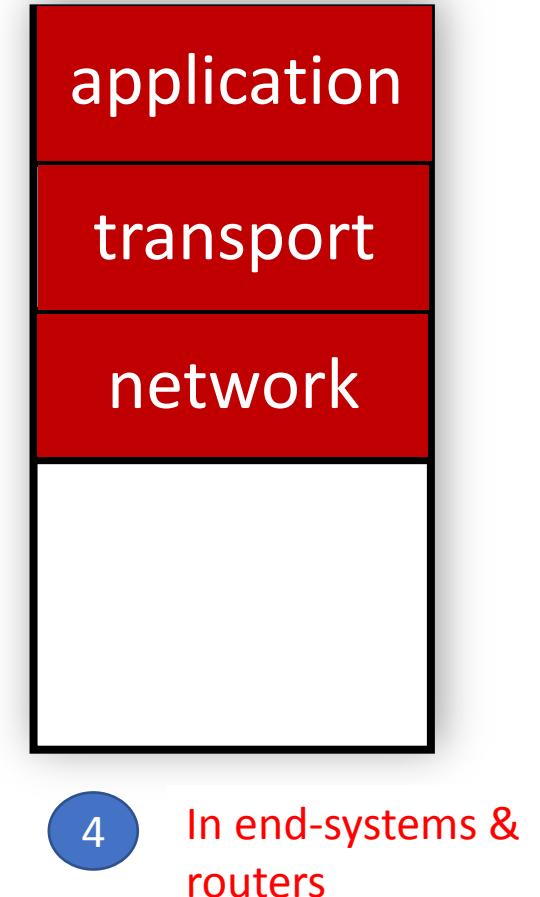
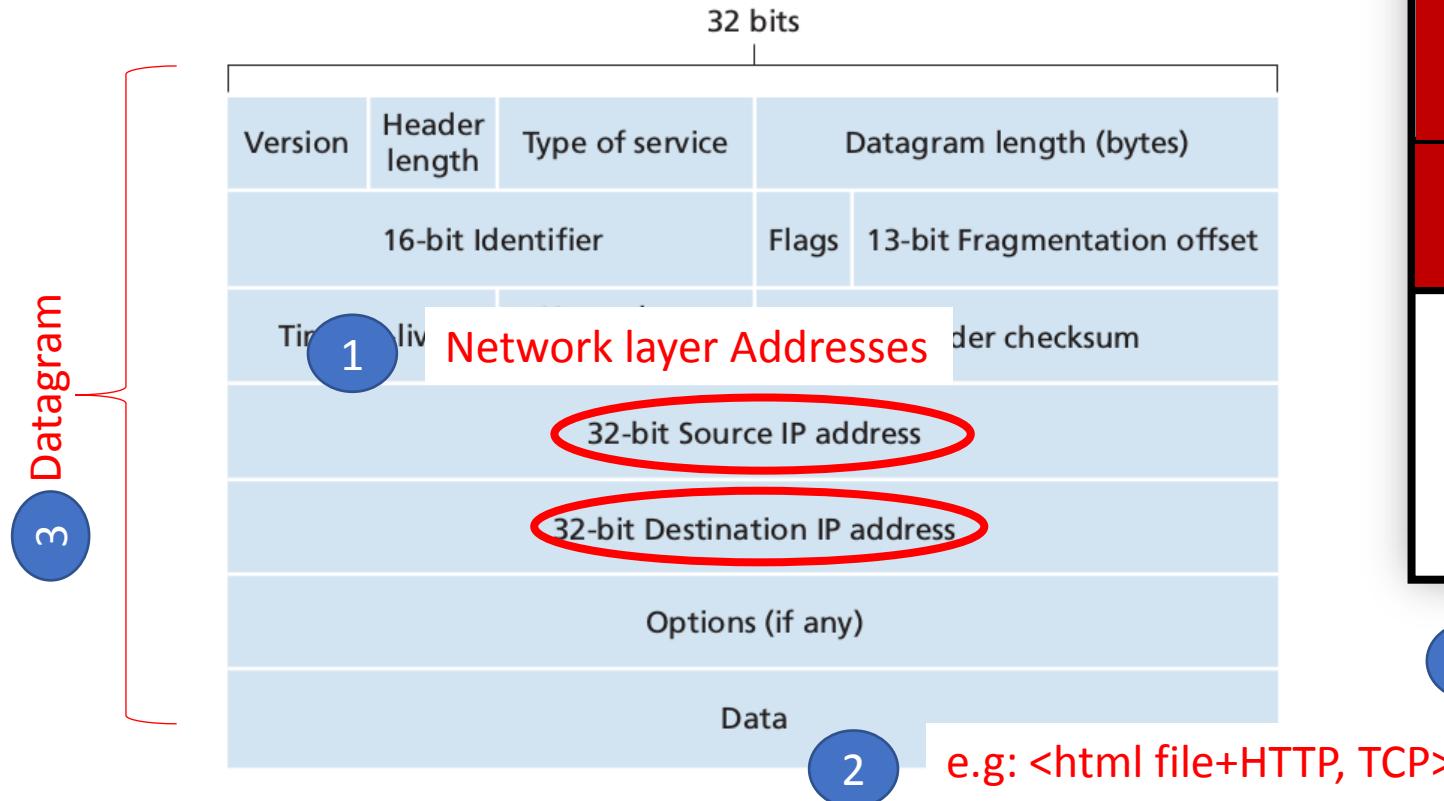
Layered Internet protocol stack

- *network*: routing of datagrams from source to destination
 - IP, routing protocols

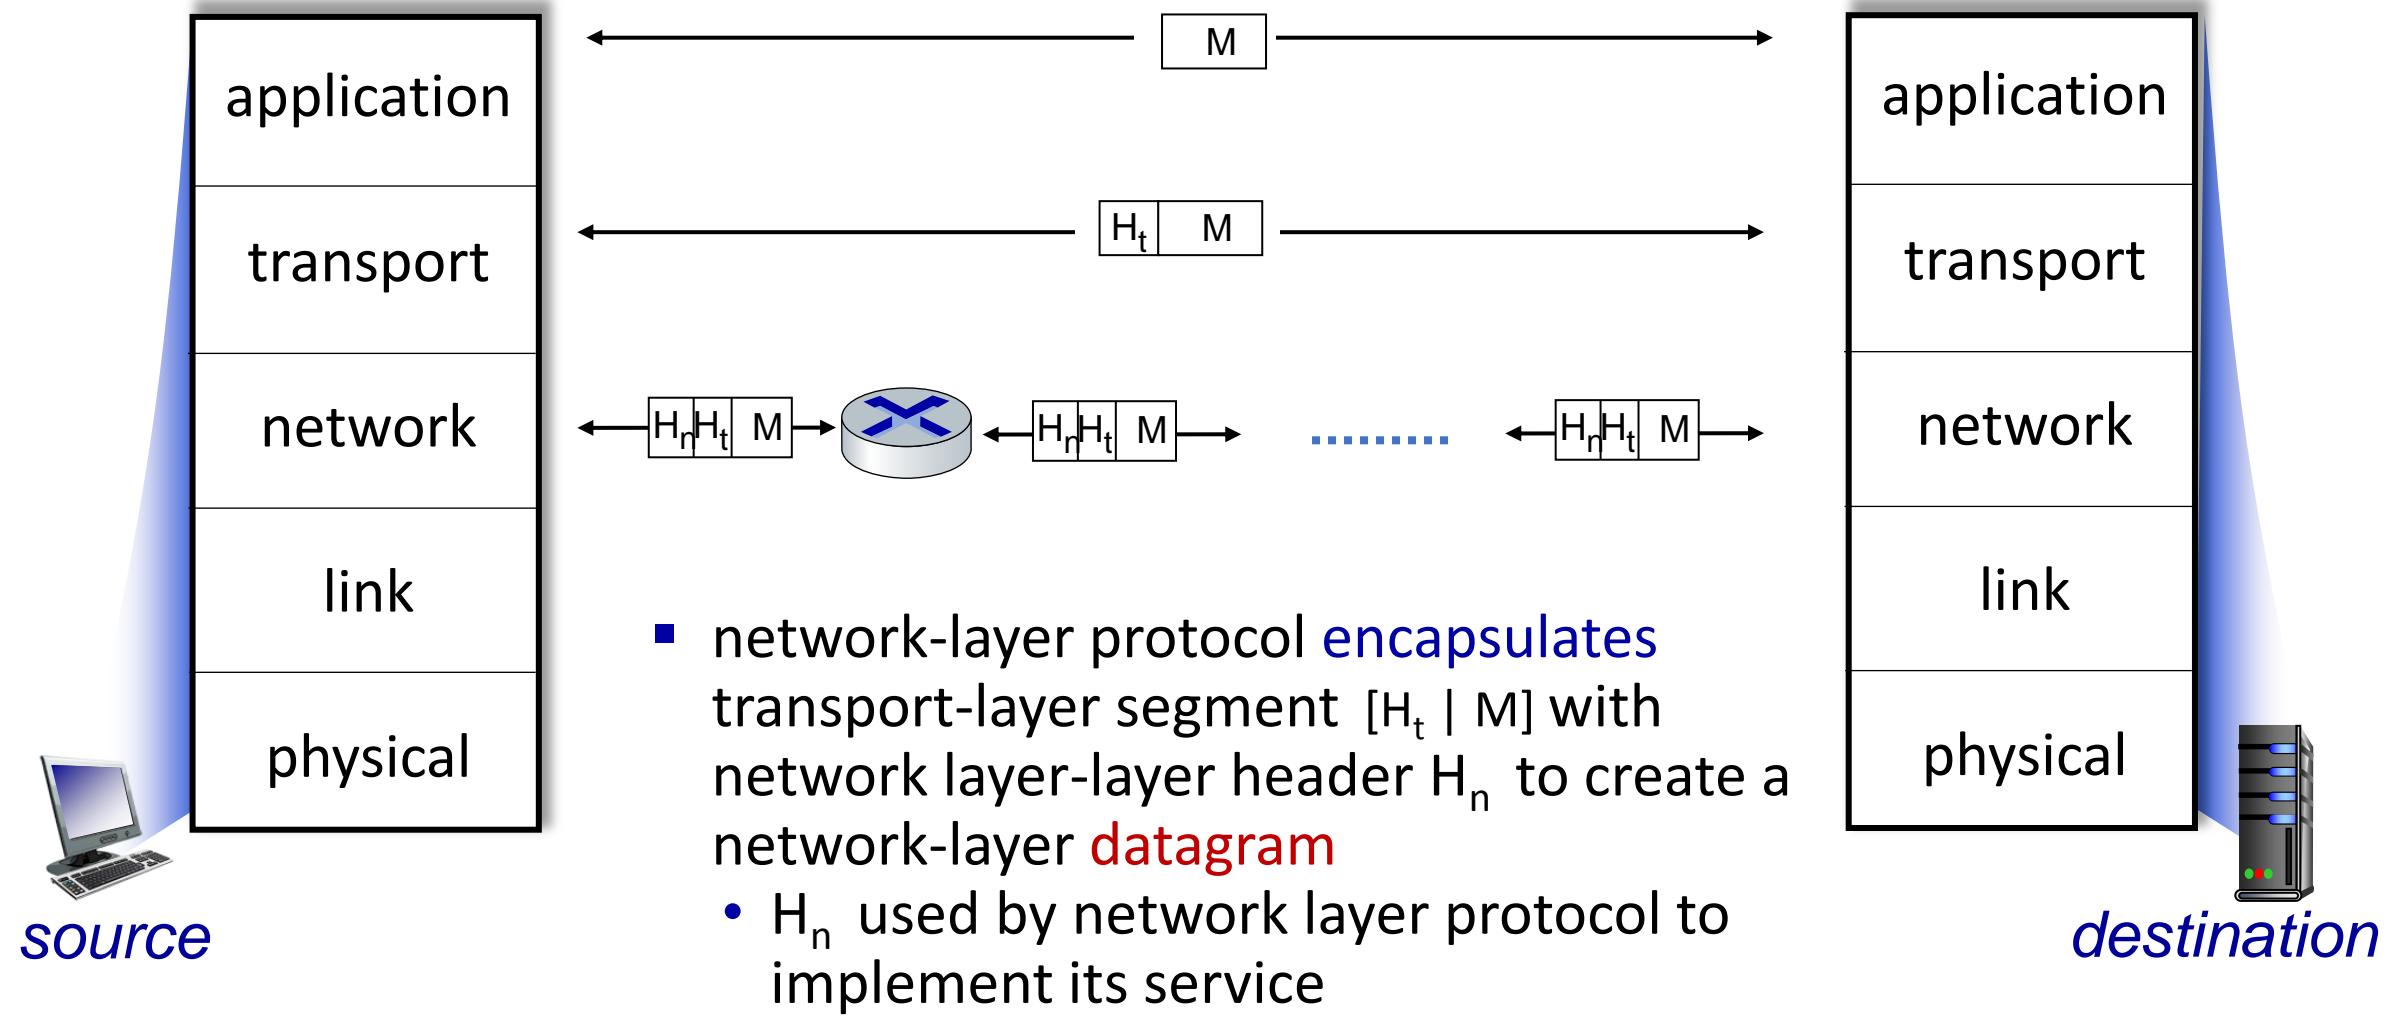


Layered Internet protocol stack

- *network*: routing of datagrams from source to destination
 - IP, routing protocols

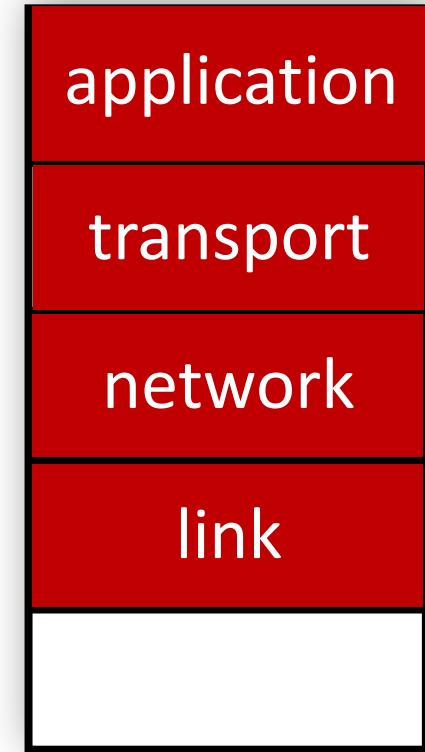
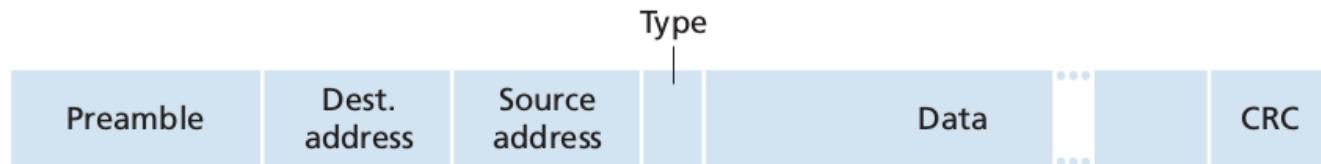


Encapsulation



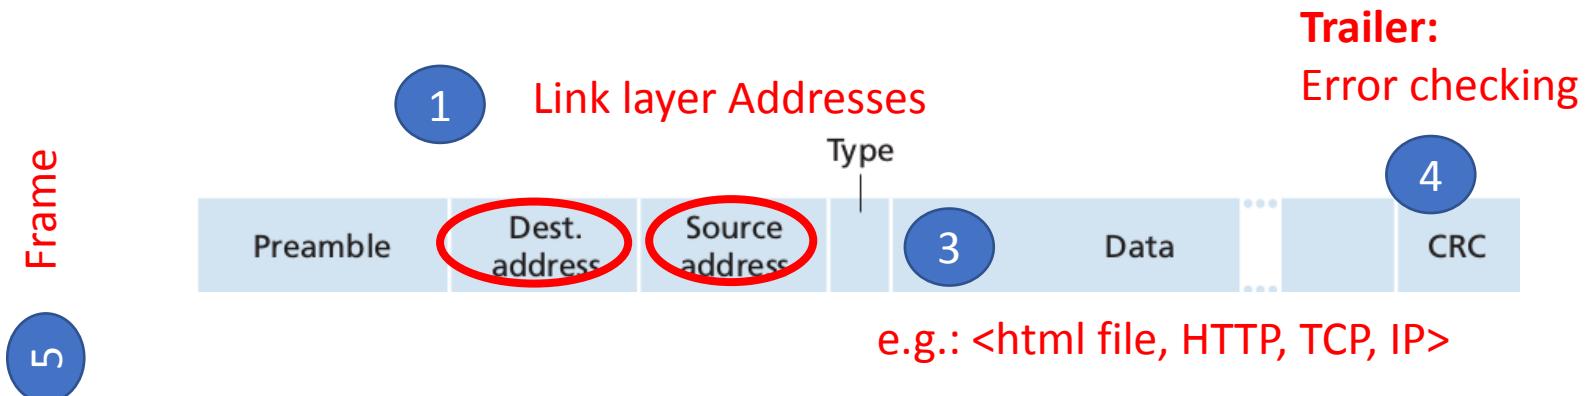
Layered Internet protocol stack

- *link*: data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP

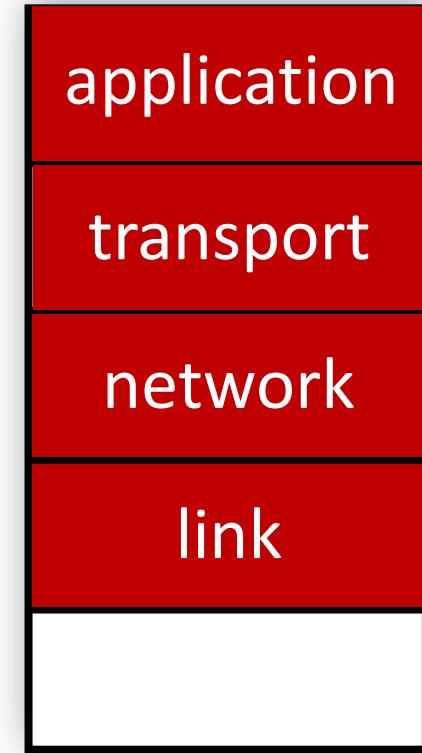


Layered Internet protocol stack

- **link:** data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP

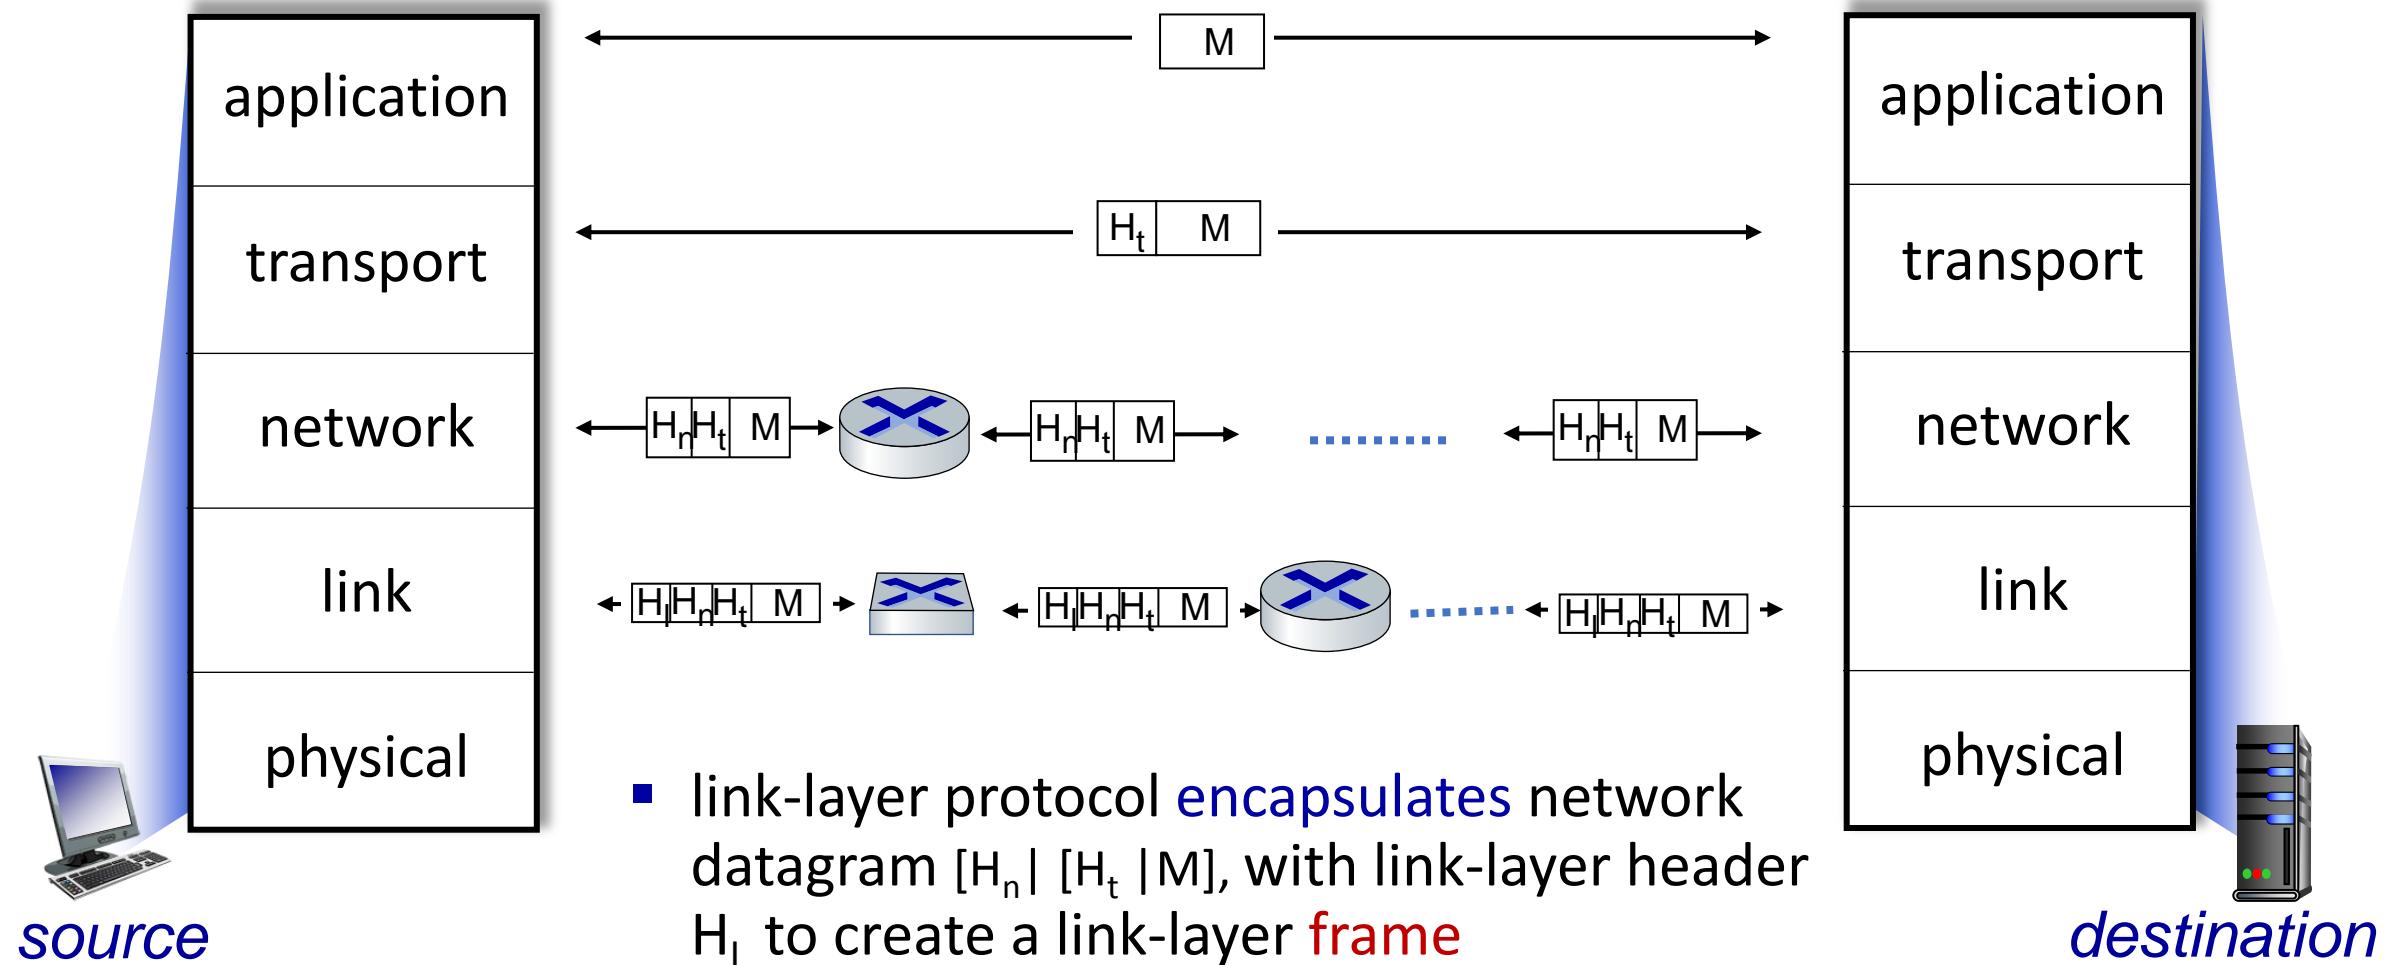


- 2 Link addresses (e.g. MAC) are different from IP address
 - >> IP addresses are fixed from source to destination
 - >>Link addresses change node-to-node (exception: switch)



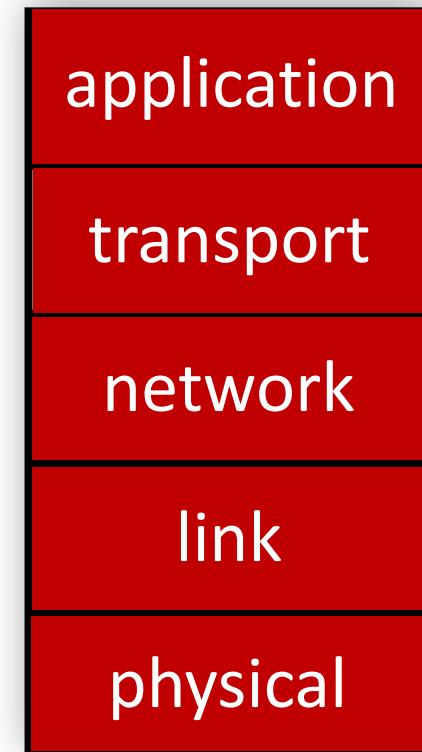
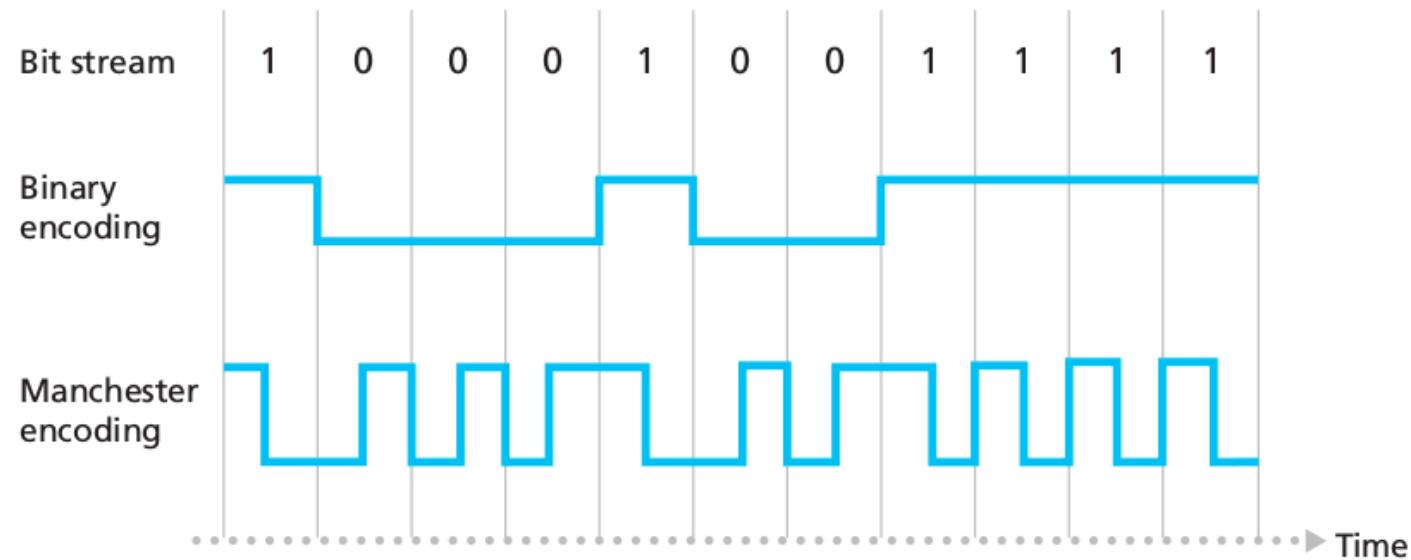
- 6 In all nodes

Encapsulation

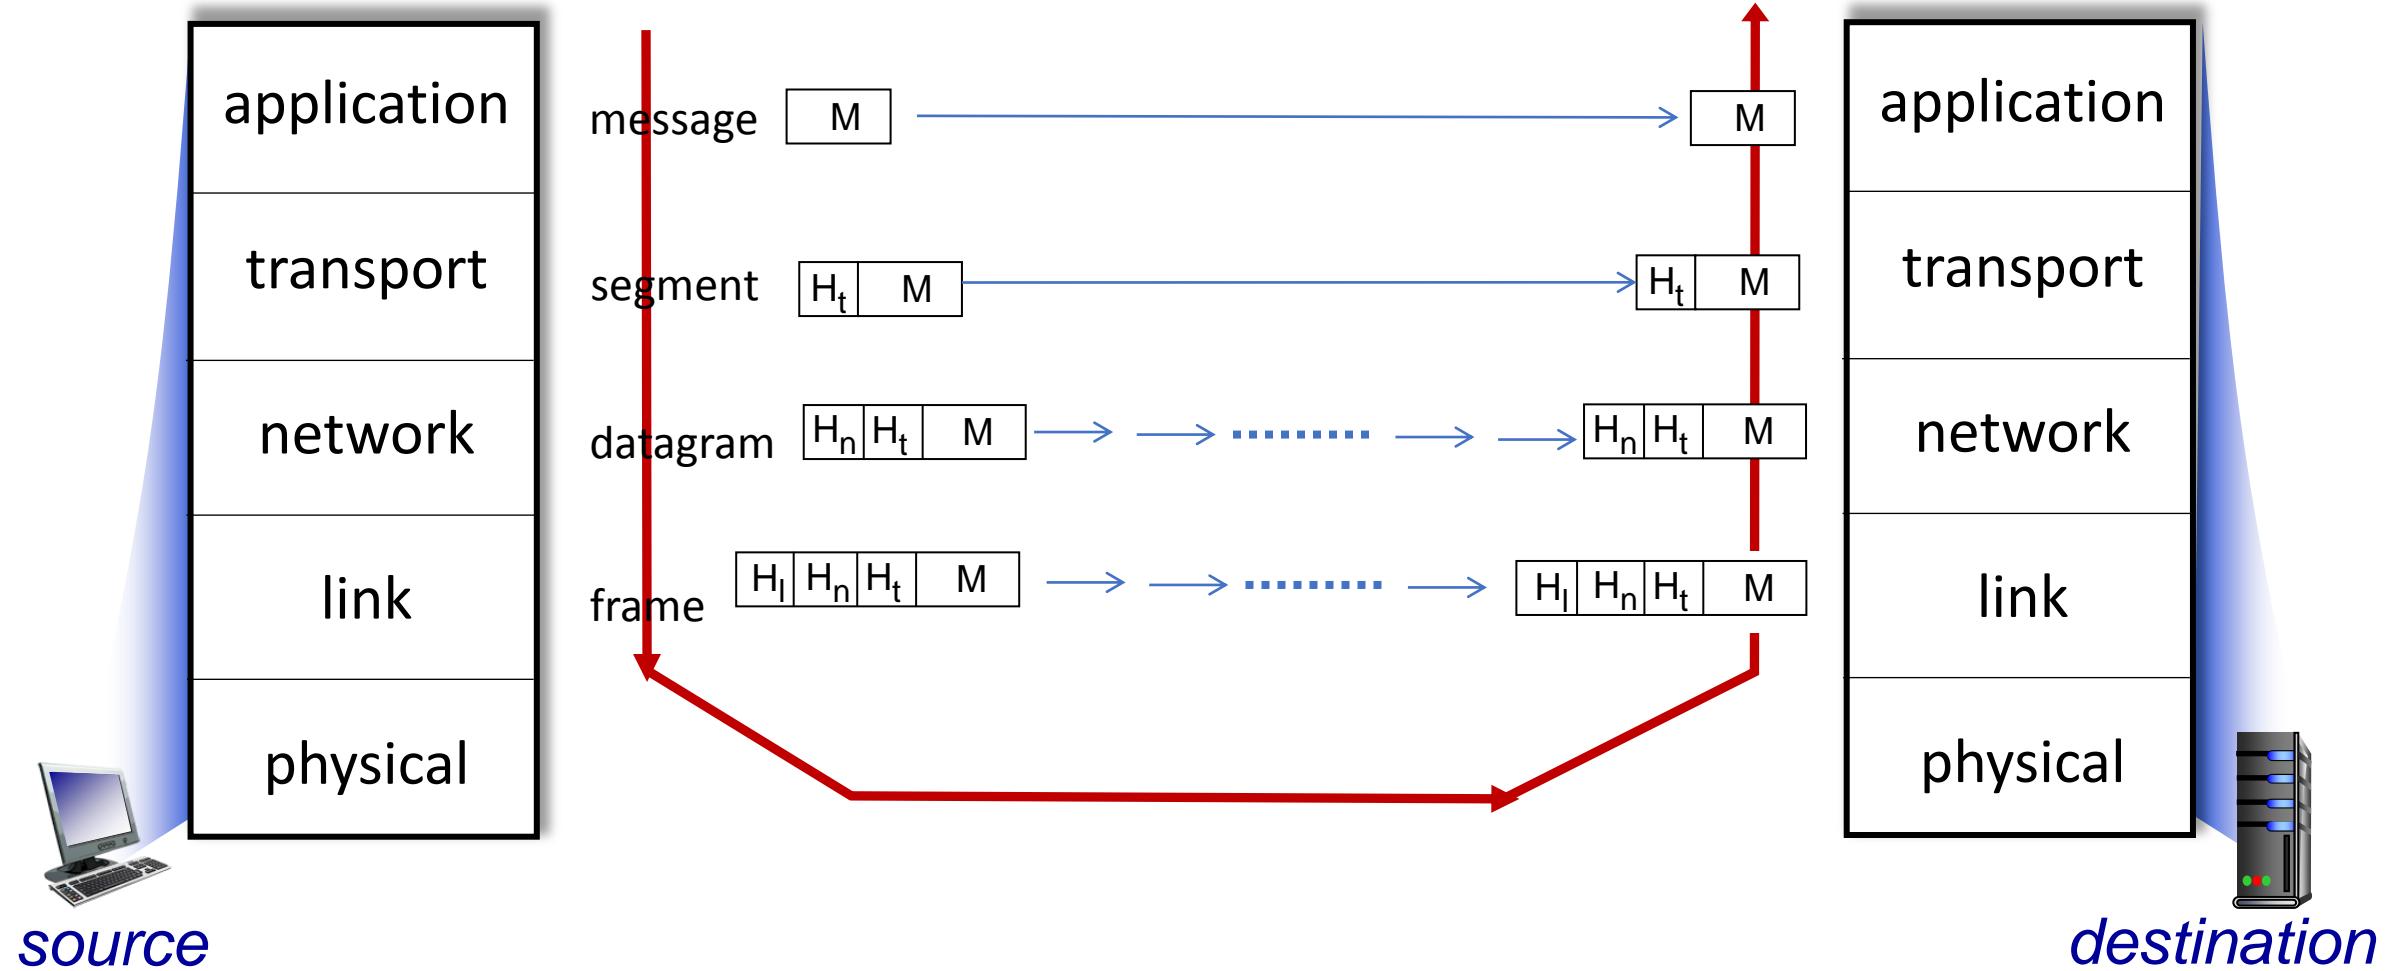


Layered Internet protocol stack

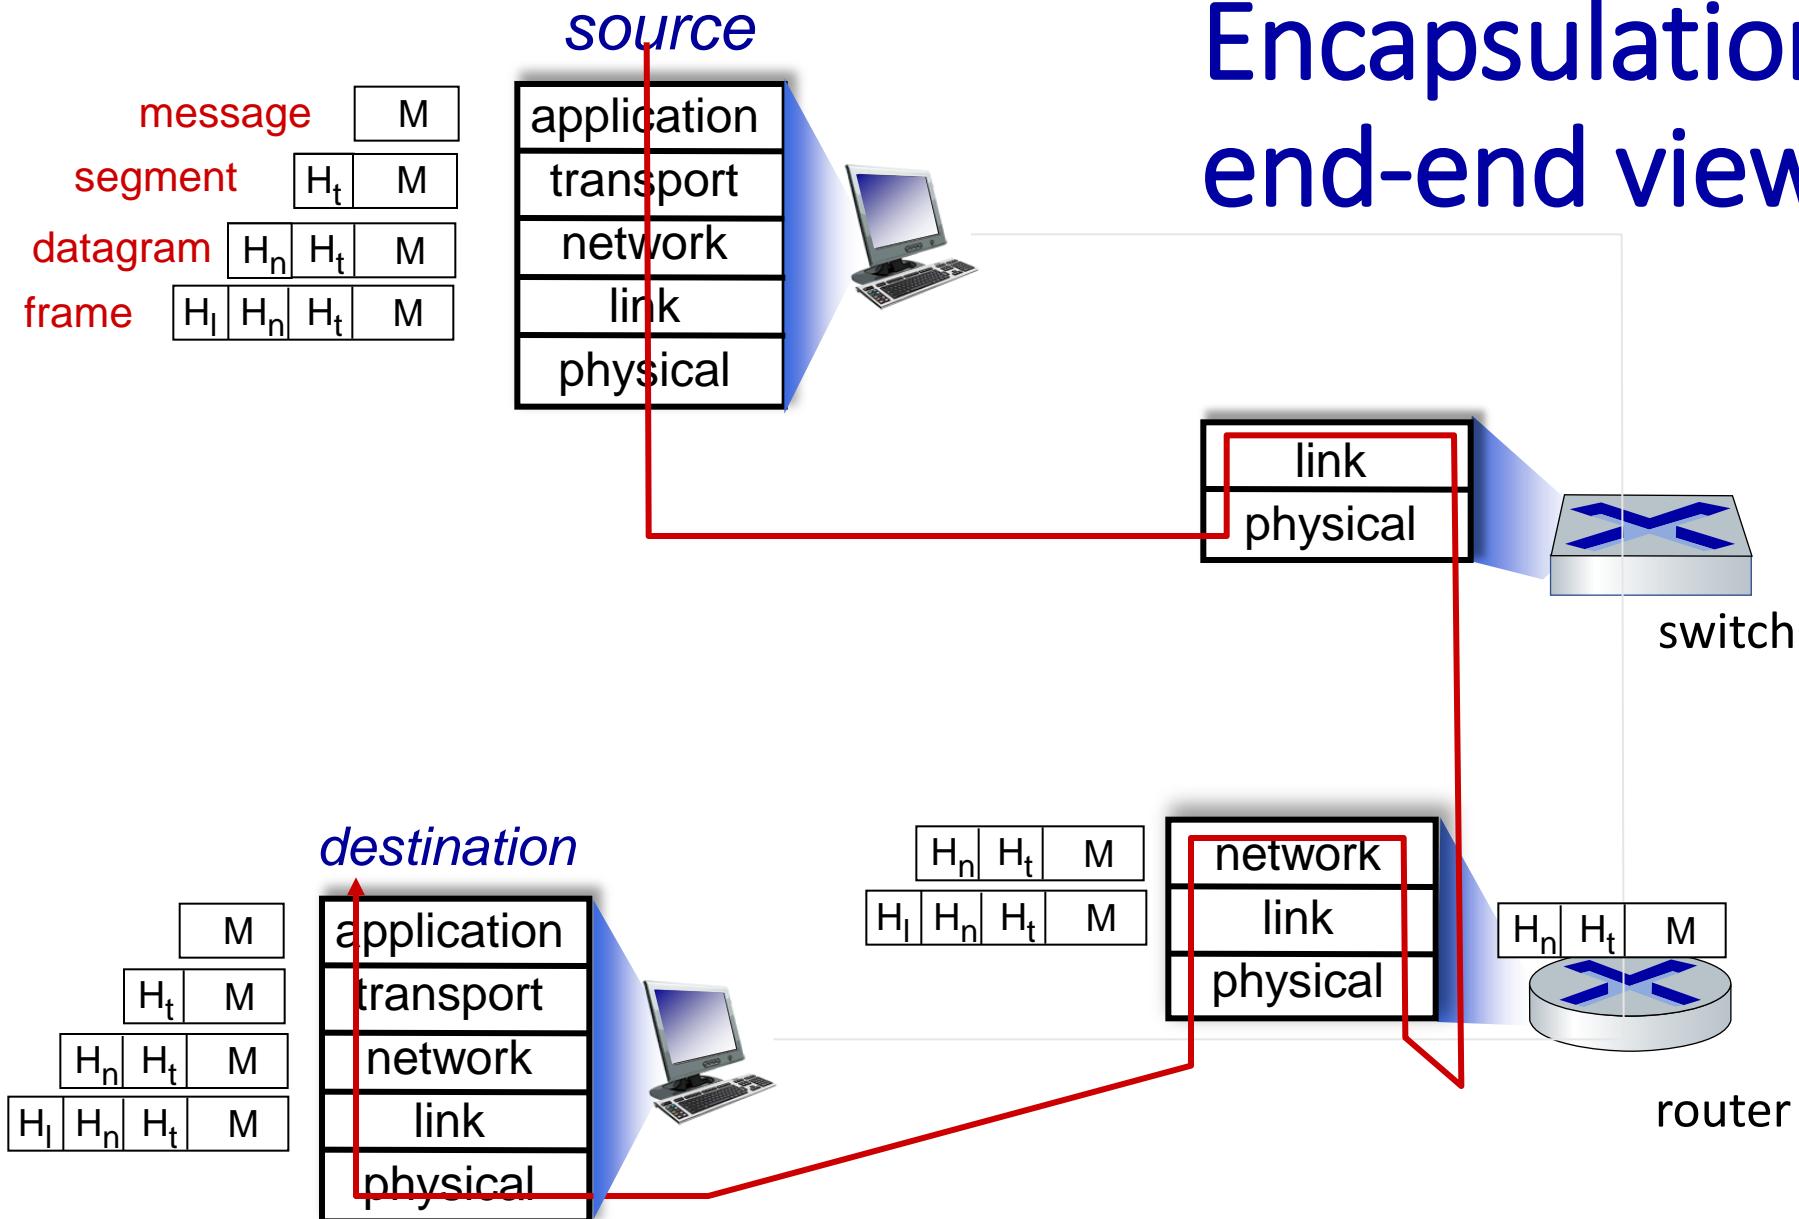
- *physical*: bits “on the wire”



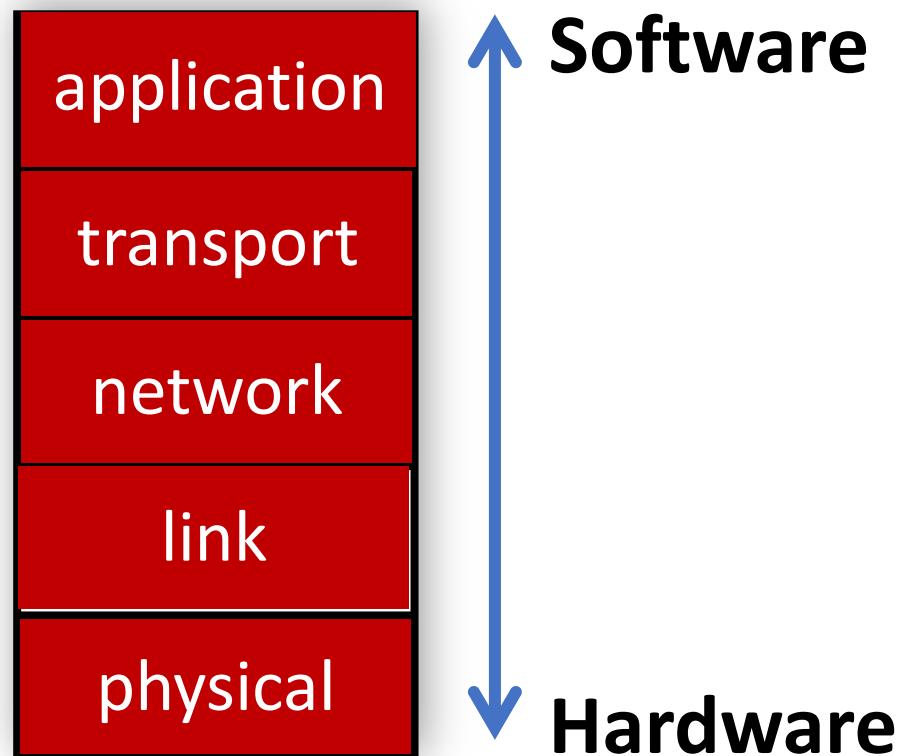
Encapsulation/De-encapsulation



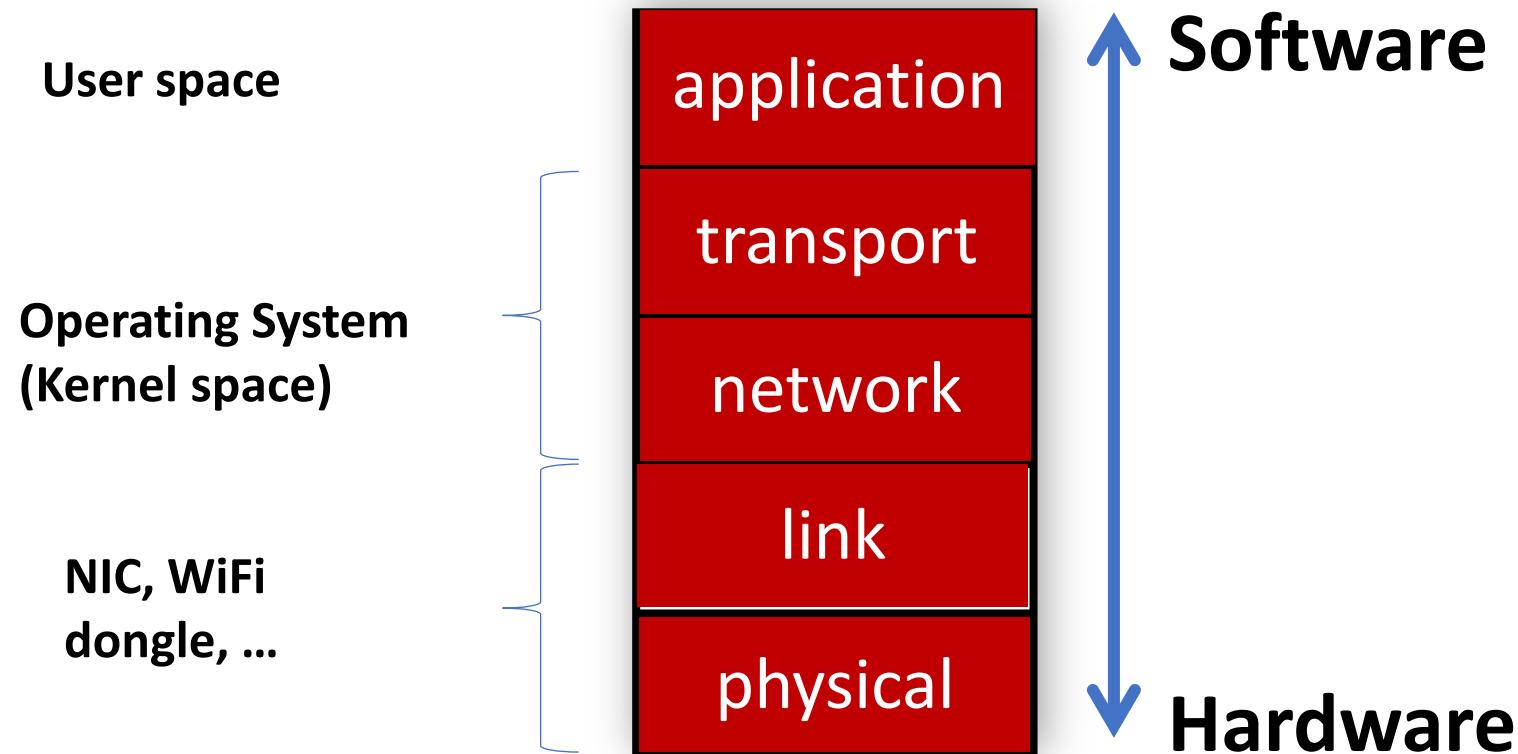
Encapsulation: an end-end view



Layered Internet protocol stack



Layered Internet protocol stack



Wireshark



Apply a display filter ... <Ctrl-/>

No.	Time	Source	Destination	Protocol	Length	Info
21	11:41:26.100360786	194.146.151.20	172.21.48.64	DNS	547	Standar
22	11:41:26.100513169	172.21.48.64	93.184.216.34	TCP	74	38182
23	11:41:26.305545043	93.184.216.34	172.21.48.64	TCP	74	80 → 3
24	11:41:26.305592198	172.21.48.64	93.184.216.34	TCP	66	38182
25	11:41:26.305752025	172.21.48.64	93.184.216.34	TCP	516	38182
26	11:41:26.434687441	172.21.48.64	52.35.31.120	TCP	66	44502
27	11:41:26.512493691	93.184.216.34	172.21.48.64	TCP	66	80 → 3

► Frame 25: 516 bytes on wire (4128 bits), 516 bytes captured (4128 bits) on interface 0
► Ethernet II, Src: LcfcHefe_d5:88:c1 (28:d2:44:d5:88:c1), Dst: Cisco_03:ea:49 (00:1d:71:03:ea:49)
► Internet Protocol Version 4, Src: 172.21.48.64, Dst: 93.184.216.34
► Transmission Control Protocol, Src Port: 38182, Dst Port: 80, Seq: 2336853736, Ack: 4278991565
► Data (450 bytes)

0000	00 1d 71 03 ea 49 28 d2 44 d5 88 c1 08 00 45 00	.q..I(. D.....E.
0010	01 f6 f9 45 40 00 40 06 2d 8c ac 15 30 40 5d b8	...E@. - ...@].
0020	d8 22 95 26 00 50 8b 49 8e e8 ff 0c 3a cd 80 18	.".&.P.I:....
0030	00 1d a5 da 00 00 01 01 08 0a 00 2c 77 d7 fa 48,w..H
0040	8d 80 47 45 54 20 2f 69 6e 64 65 78 2e 68 74 6d	.GET /i ndex.htm
0050	6c 20 48 54 54 50 2f 31 2e 31 0d 0a 48 6f 73 74	l HTTP/1 .1..Host
0060	3a 20 65 78 61 6d 70 6c 65 2e 63 6f 6d 0d 0a 55	: exempl e.com..U
0070	73 65 72 2d 41 67 65 6e 74 3a 20 4d 6f 7a 69 6c	ser-Agen t: Mozil
0080	6c 61 2f 35 2e 30 20 28 58 31 31 3b 20 55 62 75	la/5.0 (X11; Ubu
0090	6e 74 75 3b 20 4c 69 6e 75 78 20 78 38 36 5f 36	ntu; Lin ux x86_6
00a0	34 3b 20 72 76 3a 37 35 2e 30 29 20 47 65 63 6b	4; rv:75 .0) Geck
00b0	6f 2f 32 30 31 30 30 31 30 31 20 46 69 72 65 66	o/201001 01 Firef
00c0	6f 78 2f 37 35 2e 30 0d 0a 41 63 63 65 70 74 3a	ox/75.0. .Accept:
00d0	20 74 65 78 74 2f 68 74 6d 6c 2c 61 70 70 6c 69	text/ht ml,appli
00e0	63 61 74 69 6f 6e 2f 78 68 74 6d 6c 2b 78 6d 6c	cation/x html+xml

Chapter 1: roadmap

- What *is* the Internet?
- What *is* a protocol?
- Network edge: hosts, access network, physical media
- Network core: packet/circuit switching, internet structure
- Performance: loss, delay, throughput
- **Security**
- Protocol layers, service models
- History



Network security

- Internet not originally designed with (much) security in mind
 - *original vision:* “a group of mutually trusting users attached to a transparent network” ☺
 - Internet protocol designers playing “catch-up”
 - security considerations in all layers!
- We now need to think about:
 - how bad guys can attack computer networks
 - how we can defend networks against attacks
 - how to design architectures that are immune to attacks

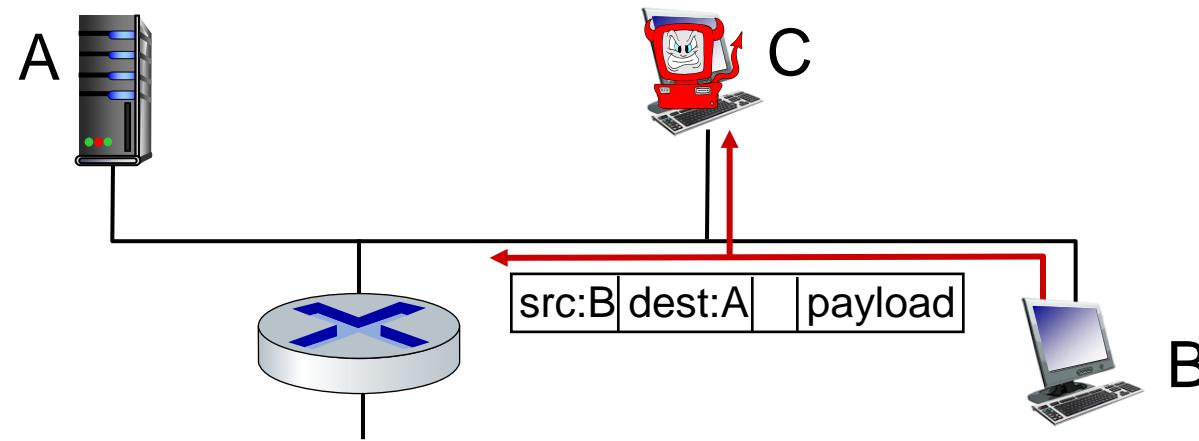
Network security

- Internet not originally designed with (much) security in mind
 - *original vision:* “a group of mutually trusting users attached to a transparent network” ☺
 - Internet protocol designers playing “catch-up”
 - security considerations in all layers!
- We now need to think about:
 - how bad guys can attack computer networks
 - how we can defend networks against attacks
 - how to design architectures that are immune to attacks

Bad guys: packet interception

packet “sniffing”:

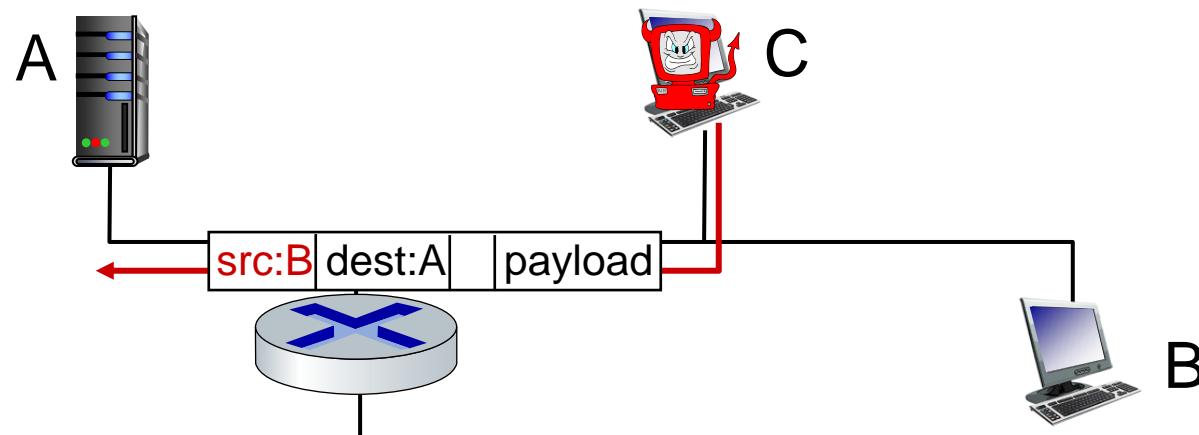
- broadcast media (shared Ethernet, wireless)
- promiscuous network interface reads/records all packets (e.g., including passwords!) passing by



Wireshark software used for our end-of-chapter labs is a (free) packet-sniffer

Bad guys: fake identity

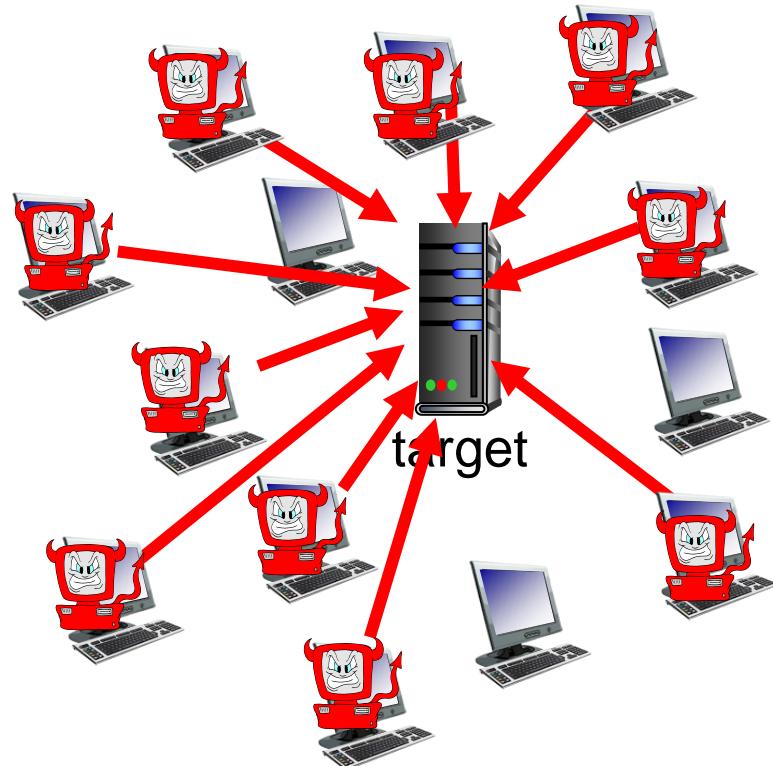
IP spoofing: injection of packet with false source address



Bad guys: denial of service

Denial of Service (DoS): attackers make resources (server, bandwidth) unavailable to legitimate traffic by overwhelming resource with bogus traffic

1. select target
2. break into hosts
around the network
(see botnet)
3. send packets to target
from compromised
hosts



Chapter 1: roadmap

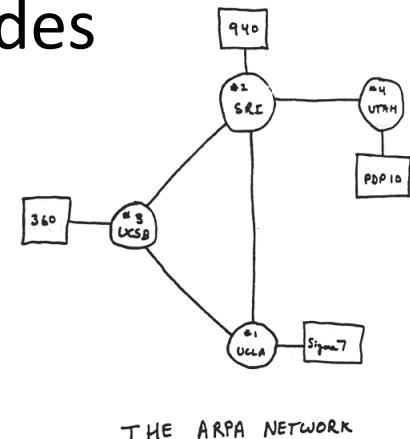
- What *is* the Internet?
- What *is* a protocol?
- Network edge: hosts, access network, physical media
- Network core: packet/circuit switching, internet structure
- Performance: loss, delay, throughput
- Security
- Protocol layers, service models
- History



Internet history

1961-1972: Early packet-switching principles

- 1961: Kleinrock - queueing theory shows effectiveness of packet-switching
- 1964: Baran - packet-switching in military nets
- 1967: ARPAnet conceived by Advanced Research Projects Agency
- 1969: first ARPAnet node operational
- 1972:
 - ARPAnet public demo
 - NCP (Network Control Protocol) first host-host protocol
 - first e-mail program
 - ARPAnet has 15 nodes



Internet history

1972-1980: Internetworking, new and proprietary networks

- 1970: ALOHAnet satellite network in Hawaii
- 1974: Cerf and Kahn - architecture for interconnecting networks
- 1976: Ethernet at Xerox PARC
- late70's: proprietary architectures: DECnet, SNA, XNA
- 1979: ARPAnet has 200 nodes

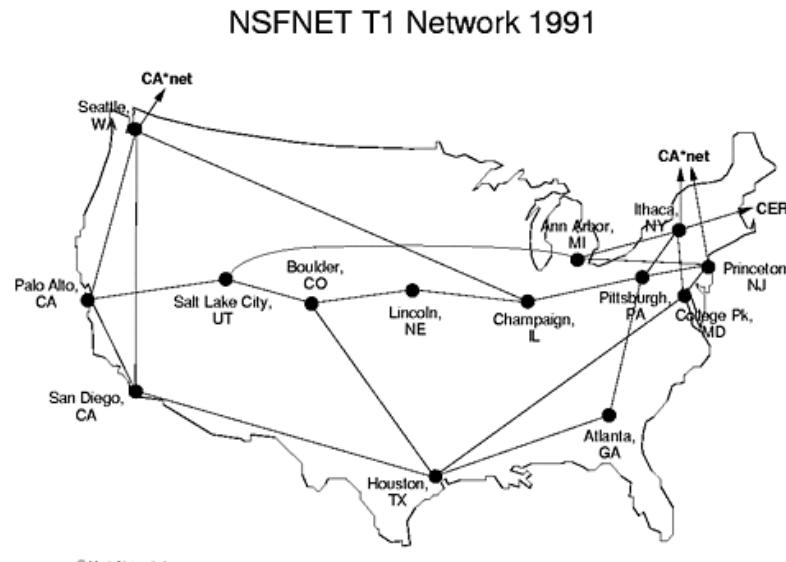
Cerf and Kahn's internetworking principles:

- minimalism, autonomy - no internal changes required to interconnect networks
 - best-effort service model
 - stateless routing
 - decentralized control
- define today's Internet architecture

Internet history

1980-1990: new protocols, a proliferation of networks

- 1983: deployment of TCP/IP
- 1982: smtp e-mail protocol defined
- 1983: DNS defined for name-to-IP-address translation
- 1985: ftp protocol defined
- 1988: TCP congestion control
- new national networks: CSnet, BITnet, NSFnet, Minitel
- 100,000 hosts connected to confederation of networks



Internet history

1990, 2000s: commercialization, the Web, new applications

- early 1990s: ARPAnet decommissioned
 - 1991: NSF lifts restrictions on commercial use of NSFnet (decommissioned, 1995)
 - early 1990s: Web
 - hypertext [Bush 1945, Nelson 1960's]
 - HTML, HTTP: Berners-Lee
 - 1994: Mosaic, later Netscape
 - late 1990s: commercialization of the Web
- late 1990s – 2000s:
- more killer apps: instant messaging, P2P file sharing
 - network security to forefront
 - est. 50 million host, 100 million+ users
 - backbone links running at Gbps

Internet history

2005-present: scale, SDN, mobility, cloud

- aggressive deployment of broadband home access (10-100's Mbps)
- 2008: software-defined networking (SDN)
- increasing ubiquity of high-speed wireless access: 4G/5G, WiFi
- service providers (Google, FB, Microsoft) create their own networks
 - bypass commercial Internet to connect “close” to end user, providing “instantaneous” access to social media, search, video content, ...
- enterprises run their services in “cloud” (e.g., Amazon Web Services, Microsoft Azure)
- rise of smartphones: more mobile than fixed devices on Internet (2017)
- ~18B devices attached to Internet (2017)

Chapter 1: summary

We've covered a "ton" of material!

- Internet overview
- what's a protocol?
- network edge, access network, core
 - packet-switching versus circuit-switching
 - Internet structure
- performance: loss, delay, throughput
- layering, service models
- security
- history

You now have:

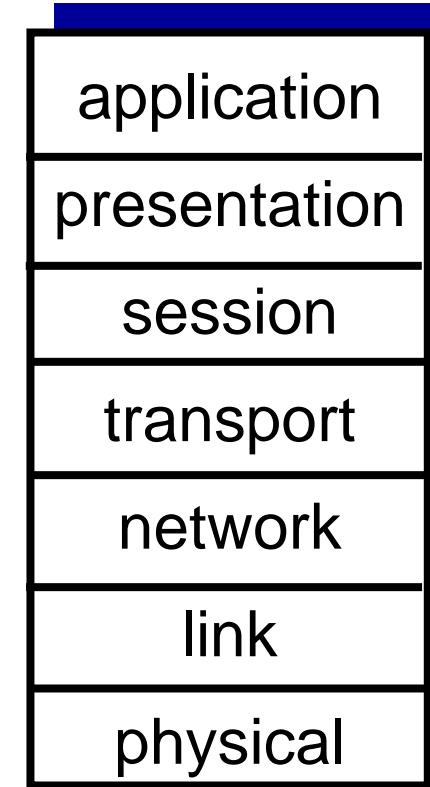
- context, overview, vocabulary, "feel" of networking
- more depth, detail, *and fun* to follow!

Additional Chapter 1 slides

ISO/OSI reference model

Two layers not found in Internet protocol stack!

- *presentation*: allow applications to interpret meaning of data, e.g., encryption, compression, machine-specific conventions
- *session*: synchronization, checkpointing, recovery of data exchange
- Internet stack “missing” these layers!
 - these services, *if needed*, must be implemented in application
 - needed?



The seven layer OSI/ISO reference model

Wireshark

