

شبکه های کامپیوتری ۲

جلسه ۱۰ فصل ۳

Congestion Control

دانشگاه صنعتی اصفهان
دانشکده مهندسی برق و کامپیوتر

Chapter 3

Transport Layer

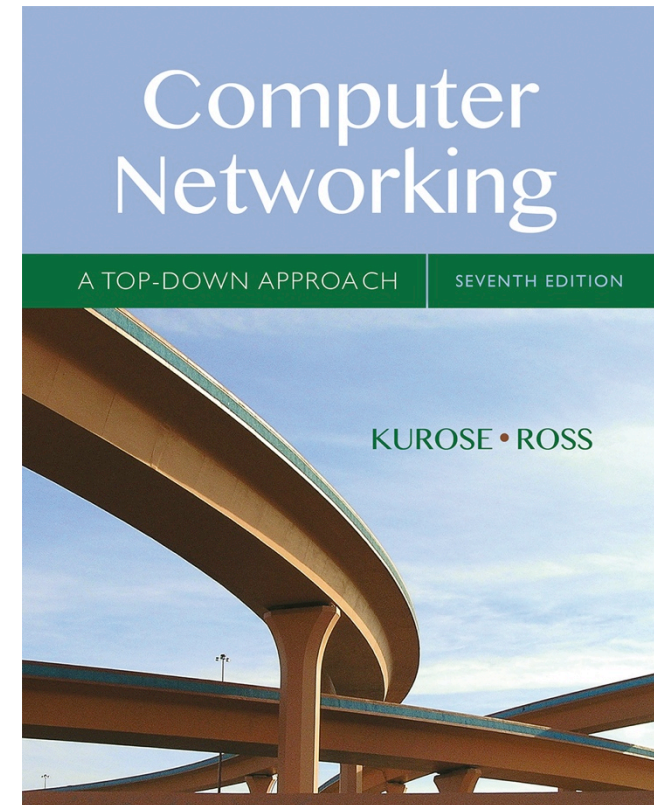
A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016
J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

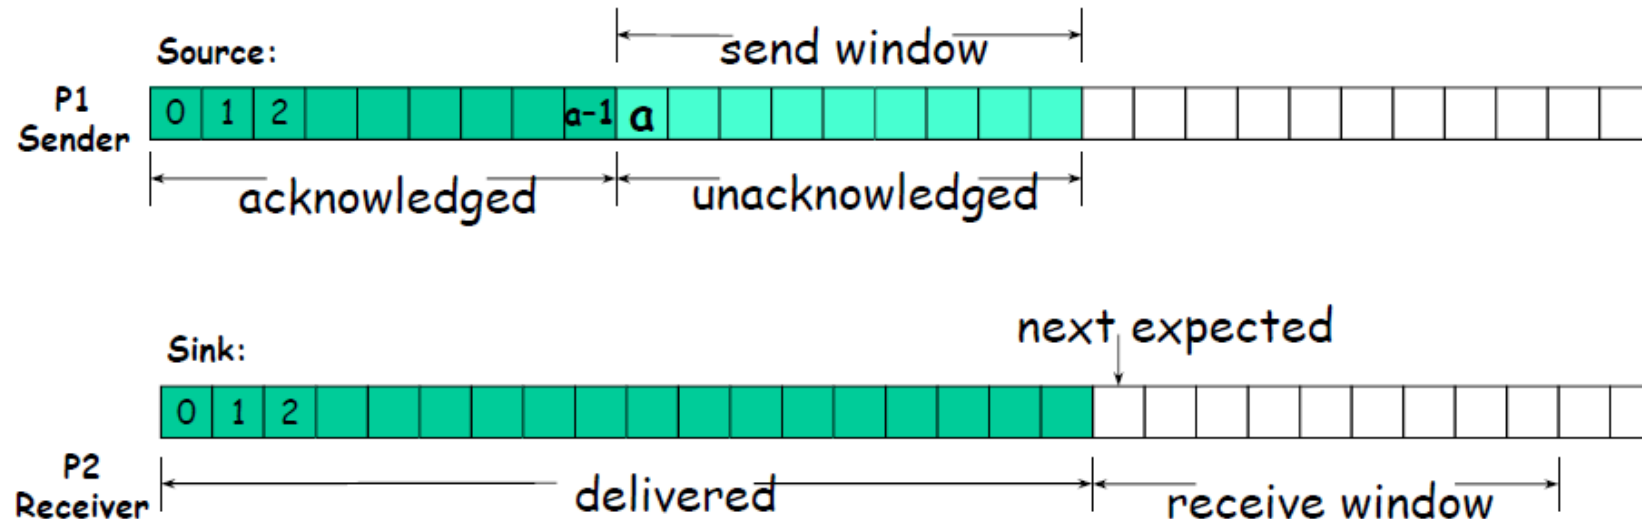
- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

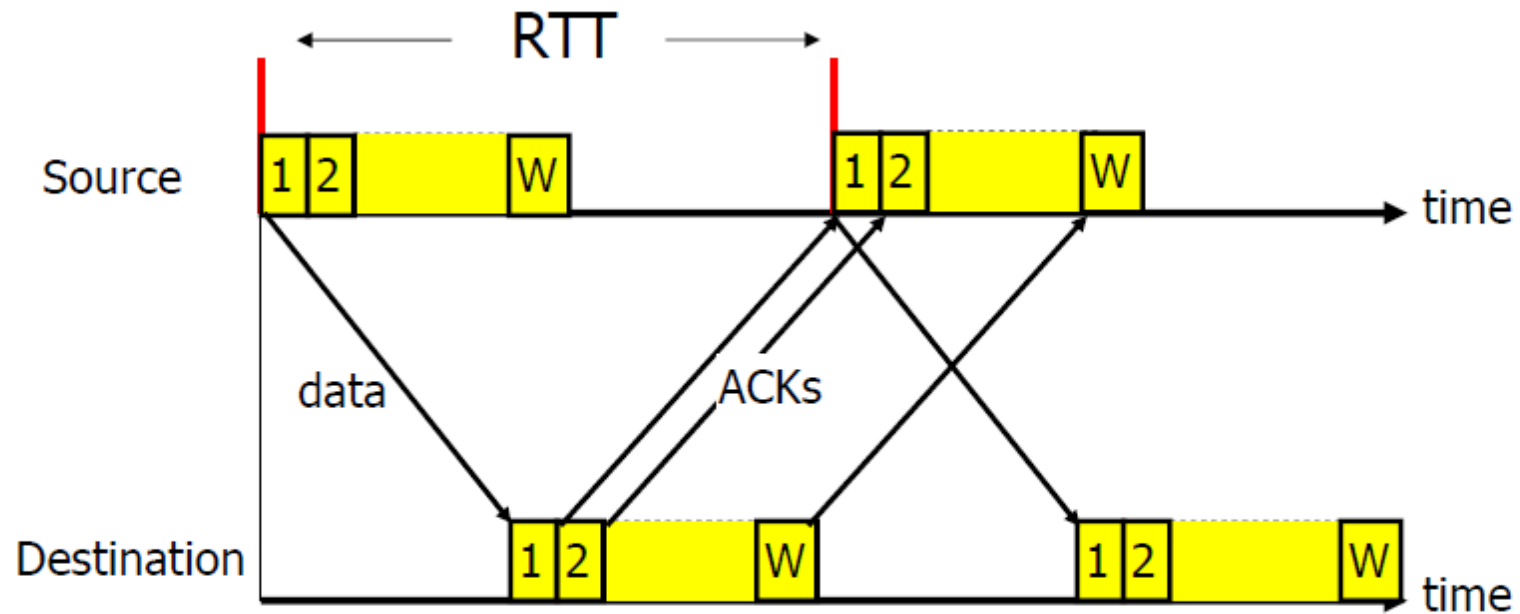
3.7 TCP congestion control

Sliding Window protocol

- Functions provided
 - reliable delivery (error and loss control)
 - in-order delivery
 - flow and congestion control
 - by varying send window size



Window Size Controls Sending Rate



□ ~ W packets per RTT when **no loss**

Throughput

Max. throughput = W / RTT bytes/sec

- This is an upper bound
- Actual throughput is smaller
 - Average number in the send buffer is less than W
 - Retransmissions
- The throughput of a host's TCP send buffer is the host's send rate into the network (including original transmissions and retransmissions)

TCP Send Window Size

- TCP flow control
 - Avoid overloading receiver
 - Receiver calculates flow control window size (**rwnd**) based on the available receiver buffer space
 - Receiver sends flow control window size to sender in TCP segment header
 - Sender keeps Send Window size less than most recently received **rwnd** value
- TCP Congestion Control
 - Avoid overloading network
 - Sender estimates network congestion from “loss indications”
 - Sender calculates congestion window size (**cwnd**)
 - Sender keeps Send Window size less than a maximum **cwnd** value
- **Sender sets $W = \min(cwnd, rwnd)$**

TCP Congestion Control

- end-to-end control (no network assistance)
- Sender limits transmission

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

$$\text{Throughput} \leq \text{cwnd} / \text{RTT} \quad \text{bytes/sec}$$

Note: For now consider **rwnd** to be very large such that the send window size is always set equal to **cwnd**

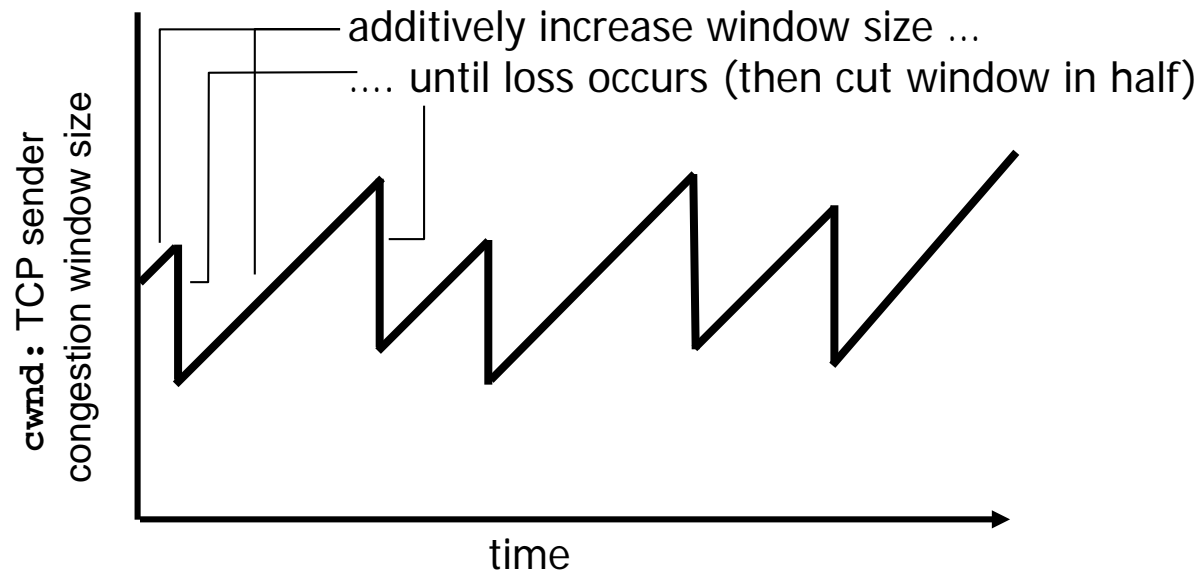
TCP Congestion Control

- How does sender estimate network congestion?
 - Packet loss is considered as an indication of network congestion
 - Time Out
 - Duplicate Acks
 - TCP sender reduces cwnd after a loss event
- How does sender determine **cwnd** size?
 - Sender adjusts existing cwnd according to the loss events
 - AIMD (Additive Increase Multiplicative Decrease)

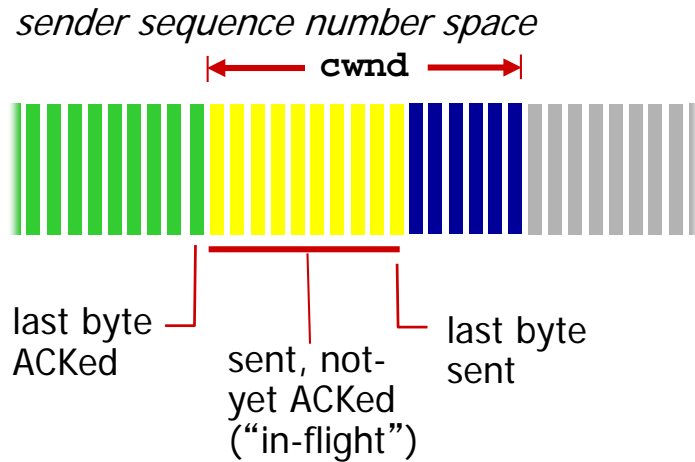
Additive Increase Multiplicative Decrease

- *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase `cwnd` by 1 MSS every RTT until loss detected
 - *multiplicative decrease*: cut `cwnd` in half after loss

AIMD saw tooth behavior: probing for bandwidth



TCP Congestion Control: details



- sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- **cwnd** is dynamic, function of perceived network congestion

TCP sending rate:

- *roughly*: send cwnd bytes, wait RTT for ACKS, then send more bytes

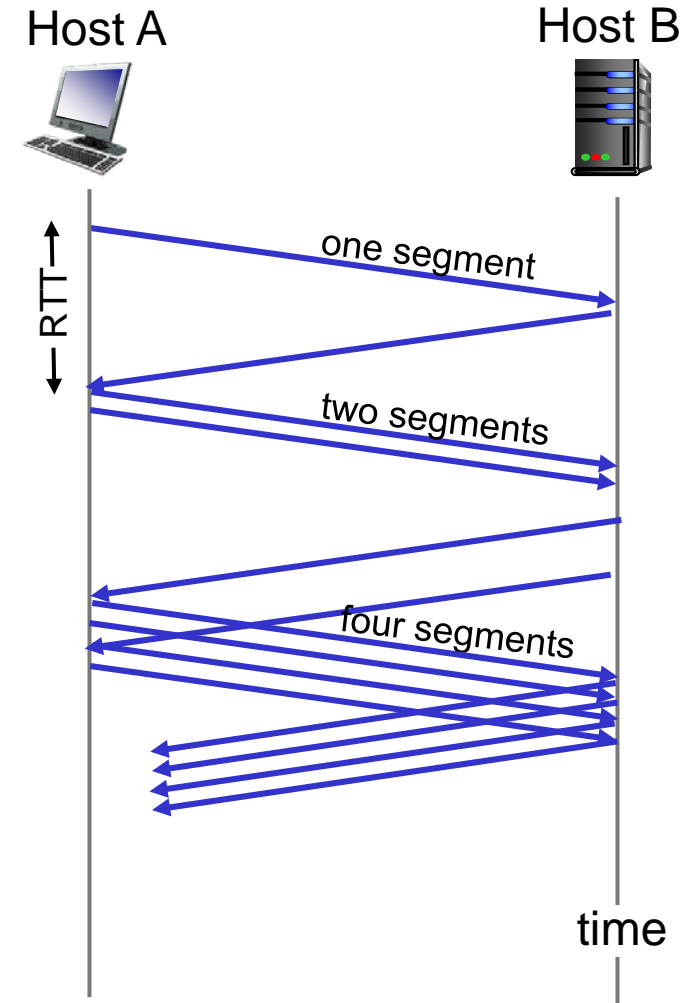
$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

At the beginning?

- TCP Slow Start:
 - Probing for usable bandwidth
 - When connection begins, **cwnd** = 1 MSS
 - Example: MSS = 500 bytes & RTT = 200 msec
 - Initial rate = 2500 bytes/sec = 20 kbps
 - Available bandwidth may be \gg MSS/RTT
 - Desirable to quickly ramp up to a higher rate

TCP Slow Start

- when connection begins, increase rate exponentially until first loss event:
 - initially `cwnd` = 1 MSS
 - double `cwnd` every RTT
 - done by incrementing `cwnd` for every ACK received
- summary: initial rate is slow but ramps up exponentially fast



TCP: detecting, reacting to loss

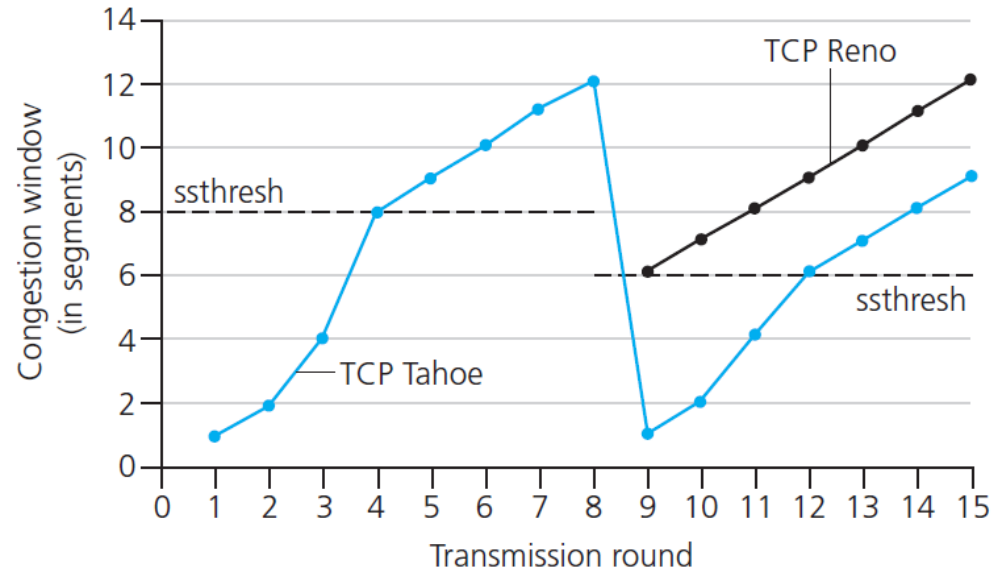
- loss indicated by timeout:
 - `cwnd` set to 1 MSS;
 - window then grows exponentially (as in slow start) to threshold, then grows linearly
- loss indicated by 3 duplicate ACKs: TCP RENO
 - dup ACKs indicate network capable of delivering some segments
 - `cwnd` is cut in half window then grows linearly
- TCP Tahoe always sets `cwnd` to 1 (timeout or 3 duplicate acks)

TCP: switching from slow start to CA

Q: when should the exponential increase switch to linear?

A: when `cwnd` gets to 1/2 of its value before timeout.

This is called slow start threshold (ssthreshold)



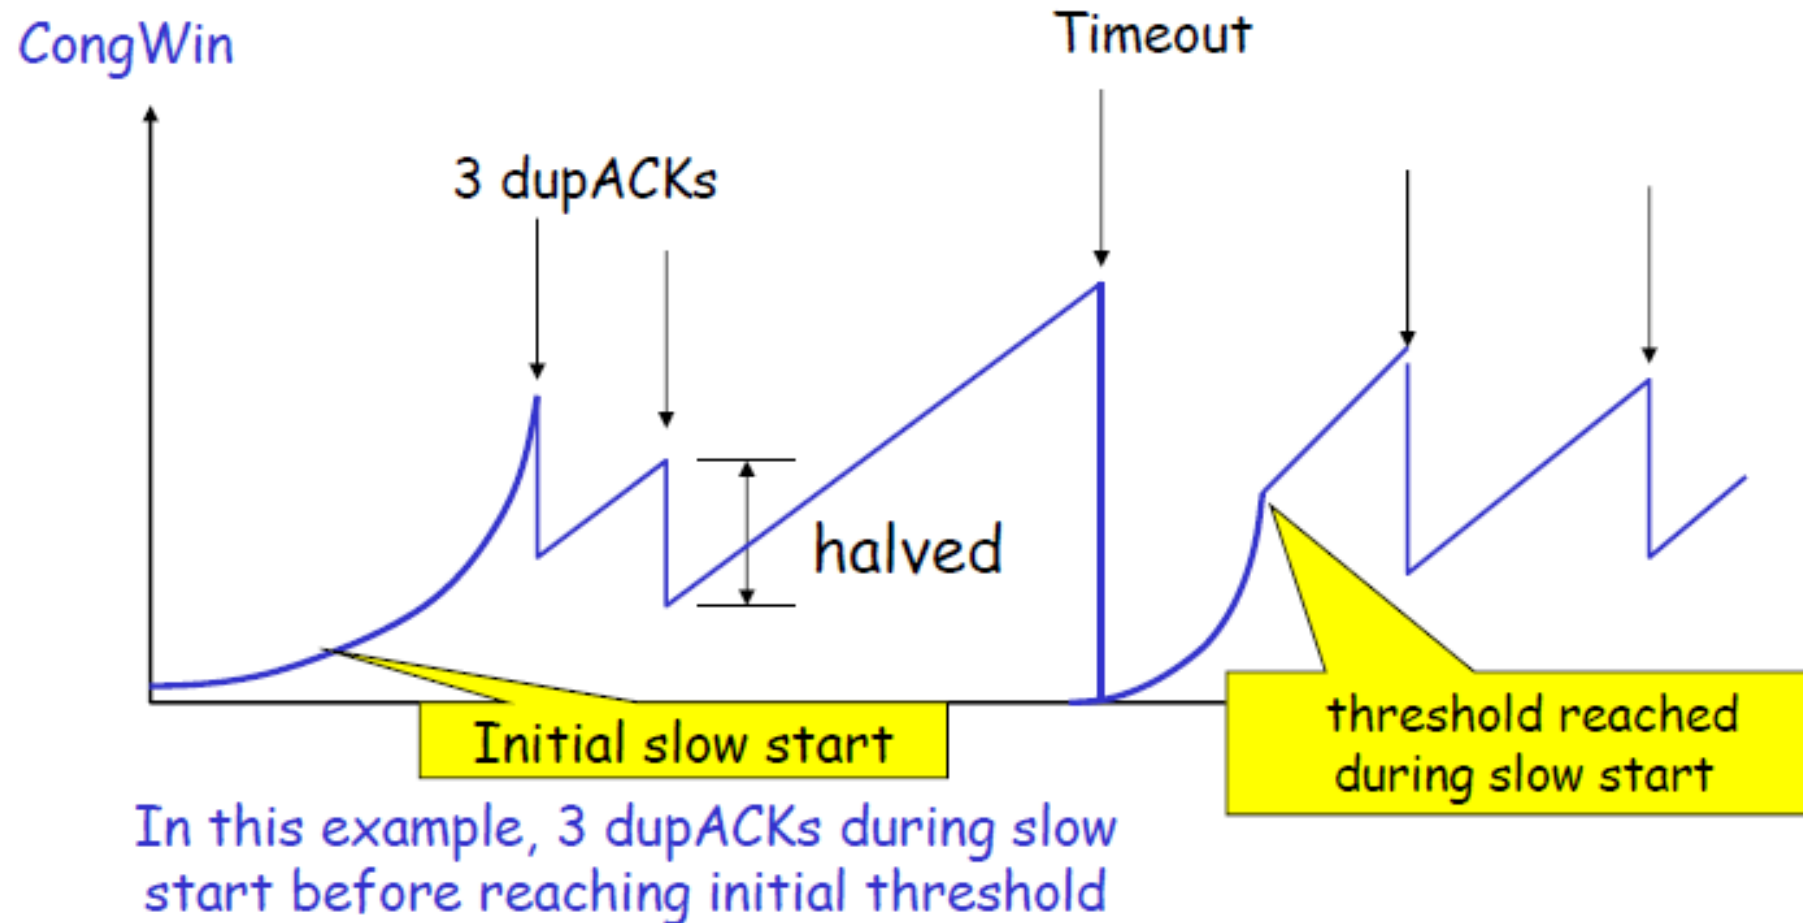
Note: For simplicity, CongWin is in number of segments in the above graph.

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Threshold

- for initial slow start, threshold is set to a large value
 - assume no threshold until the first loss event
- at a loss event, threshold is set to $1/2$ of **cwnd** just before loss event
- subsequently, threshold is variable

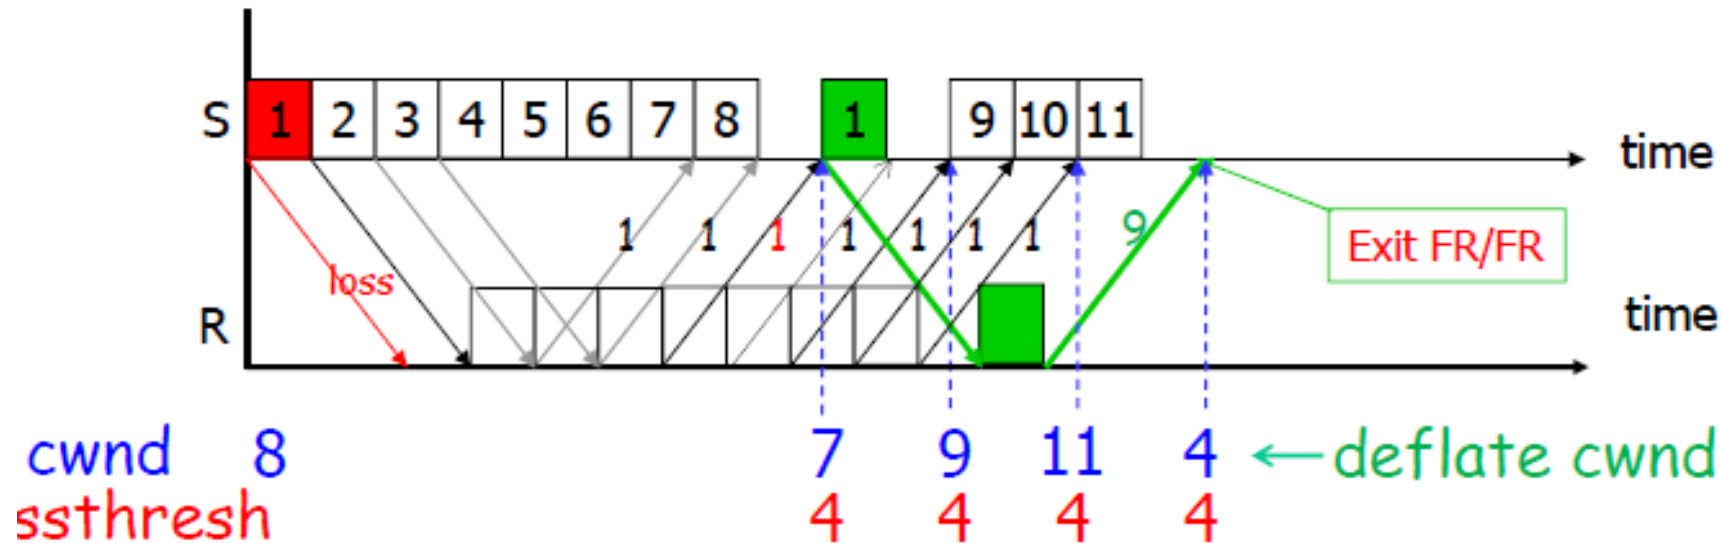
TCP Reno (example scenario)



Summary (TCP Reno)

- When cwnd is below Threshold, sender in slow-start phase, window grows exponentially (until loss event or exceeding threshold).
- When cwnd is above Threshold, sender is in congestion-avoidance phase, window grows linearly.
- When timeout occurs, Threshold set to $\text{cwnd}/2$ and cwnd is set to 1 MSS.
- When a triple duplicate ACK occurs, Threshold set to $\text{cwnd}/2$ and cwnd set to Threshold (also fast retransmit happens).

Fast Recovery entry and exit



- Above scenario: Packet 1 is lost, packets 2, 3, and 4 are received; 3 **dupACKs** with seq. no. 1 returned
- Fast retransmit
 - Retransmit **packet 1** upon 3 dupACKs
- Fast recovery (in steps)
 - **Inflate cwnd** with #dupACKs such that new packets 9, 10, and 11 can be sent while repairing loss

Summary: TCP Congestion Control

