

به نام خدا

حدیث غفوری 9825413

سوال 1

قبل از آپدیت

```
28  
29 select sum(customer_rate),count(rental_id)  
30 from rental inner join inventory using(inventory_id)  
31 inner join film f2 using(film_id)  
32 where film_id =2;  
33  
34  
35  
36
```

Data Output Explain Messages Notifications

	sum bigint	count bigint
1	350	7

```
34 select film_id,rental_id,customer_rate  
35 from rental inner join inventory using(inventory_id)  
36 inner join film f2 using(film_id)  
37 where film_id =2;  
38  
39
```

Data Output Explain Messages Notifications

	film_id integer	rental_id integer	customer_rate integer
1		2	10310
2		2	13421
3		2	4364
4		2	7733
5		2	15218
6		2	10992
7		2	11758

بعد از آپدیت

```

29 select sum(customer_rate),count(rental_id)
30 from rental inner join inventory using(inventory_id)
31 inner join film f2 using(film_id)
32 where film_id =2;
33
34 -- select film_id,rental_id,customer_rate
35 -- from rental inner join inventory using(inventory_id)
36 -- inner join film f2 using(film_id)

```

Data Output Explain Messages Notifications

	sum bigint	count bigint
1	320	7

```

39 -- update rental
40 -- set customer_rate = 20
41 -- where rental_id = 7733;
42 select * from film
43 where film_id=2;
44

```

Data Output Explain Messages Notifications

	score real
stound:4 'car':17 'china':20 'databas':8 'epistl':5 'explor':12 'find':15 'goldfing':2 'must':14	45

سوال 2

A:

If a trigger function executes SQL commands then these commands might fire triggers again. This is known as cascading triggers. There is no direct limitation on the number of cascade levels. It is possible for cascades to cause a recursive invocation of the same trigger; for example, an INSERT trigger might execute a command that inserts an additional row into the same table, causing the INSERT trigger to be fired again. It is the trigger programmer's responsibility to avoid infinite recursion in such scenarios.

Recursion occurs when the same code is executed again and again. It can lead to an infinite loop and which can result in governor limit sometime. Sometime it can also result in unexpected output.

It is very common to have recursion in the trigger which can result in unexpected output or some error. So we should write code in such a way that it does not result to recursion. But sometime we are left with no choice.

For example, we may come across a situation where in a trigger we update a field which in result invoke a workflow. The workflow contains one field update on the same object. So trigger will be executed two times. It can lead us to unexpected output.

Another example is our trigger fires on after update, and it updates some related object, and there is one more trigger on a related object which updates child object. So it can result from too infinite loop.

B:

In Postgres 9.2 or later, use the `function pg_trigger_depth()`.

so that the trigger function is not even executed when called from another trigger (including itself - so also preventing loops).

This typically performs better and is simpler and cleaner:

```
CREATE TRIGGER set_history
BEFORE UPDATE ON field_data
FOR EACH ROW
WHEN (pg_trigger_depth() < 1)
EXECUTE PROCEDURE gener_history();
```

The expression `pg_trigger_depth() < 1` is evaluated before the trigger function is entered. So it evaluates to 0 in the first call. When called from another trigger, the value is higher and the trigger function is not executed.

```

110 -- drop trigger recommendation
111 -- on rental;
112 insert into rental(rental_date,inventory_id,customer_id,staff_id)
113 values(now(),3 ,6,1);
114 -- select * from rental
115 -- where customer_id=6
116 -- order by rental_date desc;
117 -- select * from customer
118 -- where customer_id=6;
119 -- where rent_count >0;
120 -- select * from inventory;

```

Data Output Explain Messages Notifications

INSERT 0 1

Query returned successfully in 92 msec.

Query Editor Query History

```

108 -- EXECUTE PROCEDURE give_suggest();
109
110 -- drop trigger recommendation
111 -- on rental;
112 -- insert into rental(rental_date,inventory_id,customer_id,staff_id)
113 -- values(now(),3 ,6,1);
114 -- select * from rental
115 -- where customer_id=6
116 -- order by rental_date desc;
117 select * from customer
118 where customer_id=6;
119 -- where rent_count >0;
120 -- select * from inventory;

```

Data Output Explain Messages Notifications

	last_name character varying (45)	email character varying (50)	address_id smallint	activebool boolean	create_date date	last_update timestamp without time zone	active integer	rent_count integer
	Davis	jennifer.davis@sakilacustomer.org	10	true	2006-02-14	2013-05-26 14:49:45.738	1	0

```

110 -- drop trigger recommendation
111 -- on rental;
112 -- insert into rental(rental_date,inventory_id,customer_id,staff_id)
113 -- values(now(),3,6,1);
114 -- select * from rental
115 -- where customer_id=6
116 -- order by rental_date desc;
117 select * from customer
118 where customer_id=6;
119 -- where rent_count >0;
120 -- select * from inventory;

```

Data Output										Explain	Messages	Notifications
(45)	last_name character varying (45)	email character varying (50)	address_id smallint	activebool boolean	create_date date	last_update timestamp without time zone	active integer	rent_count integer				
	Davis	jennifer.davis@sakilacustomer.org	10	true	2006-02-14	2022-01-03 18:05:59.960236	1					

```

112 -- insert into rental(rental_date,inventory_id,customer_id,staff_id)
113 -- values(now(),3,6,1);
114 -- select * from rental
115 -- where customer_id=6
116 -- order by rental_date desc;
117 select * from customer
118 where customer_id=6;
119 -- where rent_count >0;
120 -- select * from inventory;

```

Data Output									Explain	Messages	Notifications
(45)	last_name character varying (45)	email character varying (50)	address_id smallint	activebool boolean	create_date date	last_update timestamp without time zone	active integer	rent_count integer			
	Davis	jennifer.davis@sakilacustomer.org	10	true	2006-02-14	2022-01-03 18:07:14.744296	1	4			

بعد از ۴ بار insert کردن دفعه ی ۵م تریگر فعال میشود

```

114 select * from rental
115 where customer_id=6
116 order by rental_date desc;
117 -- select * from customer
118 -- where customer_id=6;
119 -- where rent_count >0;
120 -- select * from inventory;

```

Data Output	Explain	Messages	Notifications			
rental_id [PK] integer	rental_date timestamp without time zone	inventory_id integer	customer_id smallint	return_date timestamp without time zone	staff_id smallint	last_update timestamp without time zone
16066	2022-01-03 18:07:33.120202	984	6	[null]	1	2022-01-03 18:07:33.120202
16065	2022-01-03 18:07:33.120202	3	6	[null]	1	2022-01-03 18:07:33.120202
16064	2022-01-03 18:07:14.744296	3	6	[null]	1	2022-01-03 18:07:14.744296
16063	2022-01-03 18:07:00.25322	3	6	[null]	1	2022-01-03 18:07:00.25322
16062	2022-01-03 18:06:31.071843	3	6	[null]	1	2022-01-03 18:06:31.071843
16061	2022-01-03 18:05:59.960236	3	6	[null]	1	2022-01-03 18:05:59.960236

```

1
2 WITH group_box AS (
3     SELECT
4         title,
5         rating,
6         SUM(amount) sum_amount
7     from film f2 inner join inventory using(film_id)
8     inner join rental using(inventory_id)
9     inner join payment using(rental_id)
10    group by film_id
11 )
12 SELECT
13     title,
14     rating,
15     sum_amount.
```

Data Output Explain Messages Notifications

	title character varying (255)	rating mpaa_rating	sum_amount numeric	rank_in_all bigint	rank_in_rating bigint	is_in_first_quarter text
1	Academy Dinosaur	PG	33.79	718	147	no
2	Ace Goldfinger	G	52.93	508	89	no
3	Adaptation Holes	NC-17	34.89	699	137	no
4	Affair Prejudice	G	83.79	261	47	no
5	African Egg	G	47.89	557	99	no
6	Agent Truman	PG	111.81	118	24	yes
7	Airplane Sierra	PG-13	82.85	263	64	no
8	Airport Pollock	R	86.85	239	46	yes
9	Alabama Devil	PG-13	71.88	354	80	no
10	Aladdin Calendar	NC-17	131.77	66	18	yes
11	Alamo Wildcat	G	30.70	761	122	no

```

28 with group_box as (
29     select extract (month from rental_date) month_d,
30     rating,
31     sum(amount) amount
32     from film
33     inner join inventory using(film_id)
34     inner join rental r2 using(inventory_id)
35     inner join payment using(rental_id)
36     where rating is not null
37     group by rating,extract (month from rental_date)
38     order by extract (month from rental_date)
39 )
40
41 SELECT
42     distinct month_d,

```

Data Output Explain Messages Notifications

	month_d double precision	rating mpaa_rating	amount numeric	previous_month_sales numeric	next_months_sales numeric
1		2 NC-17	113.56	[null]	1666.90
2		2 PG	94.69	[null]	1659.99
3		2 PG-13	118.59	[null]	1857.58
4		2 R	82.71	[null]	1746.78
5		2 G	104.63	[null]	1423.60
6		6 PG	1658.99	94.69	5696.52
7		6 R	1745.78	82.71	5436.85
8		6 PG-13	1856.58	118.59	6521.85
9		6 G	1422.60	104.63	5014.11
10		6 NC-17	1665.90	113.56	5713.54
11		7 NC	5605.52	1659.99	4700.45

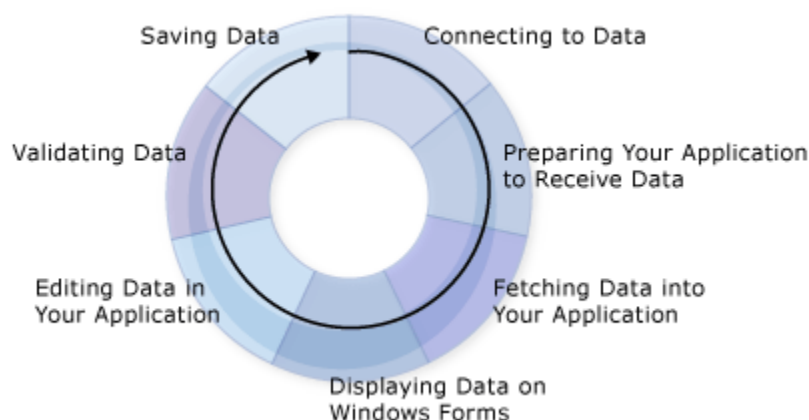
```

32 select film.film_id,city_id , s.store_id,f.staff_id,count(rental_id)
33 from film,rental r,staff f,store s,address a,inventory i
34 where f.store_id = s.store_id and
35 a.address_id = s.address_id and
36 r.staff_id = f.staff_id and
37 i.film_id = film.film_id and
38 i.inventory_id= r.inventory_id
39 group by
40 film.film_id,
41 cube(city_id , s.store_id,f.staff_id)
42 order by
43 count(rental_id) desc
44

```

Data Output Explain Messages Notifications

	film_id integer	city_id smallint	store_id integer	staff_id integer	count bigint	
1	103	[null]	[null]	[null]	34	
2	738	[null]	[null]	[null]	33	
3	489	[null]	[null]	[null]	32	
4	767	[null]	[null]	[null]	32	
5	382	[null]	[null]	[null]	32	
6	730	[null]	[null]	[null]	32	
7	331	[null]	[null]	[null]	32	
8	735	[null]	[null]	[null]	31	
9	891	[null]	[null]	[null]	31	
10	621	[null]	[null]	[null]	31	



Connecting to Data

Preparing Your Application to Receive Data

Fetching Data into Your Application

Displaying Data on Forms

Editing Data in Your Application

Validating Data

Saving Data

B:

definition of OLE DB is

a strategic system-level programming interface to data across the organization. OLE DB is an open specification designed to build on the success of ODBC by providing an open standard for accessing all kinds of data.

definition of ODBC is

an industry standard and a component of Microsoft® Windows® Open Services Architecture (WOSA). The ODBC interface makes it possible for applications to access data from a variety of database management systems (DBMSs). ODBC permits maximum interoperability—an application can access data in diverse DBMSs through a single interface. Furthermore, that application will be independent of any DBMS from which it accesses data. Users of the application can add software components called drivers, which create an interface between an application and a specific DBMS

1. OLE DB is a component based specification and ODBC is a procedural based specification
2. SQL is the core of accessing data using ODBC but just one of the means of data access through OLE DB
3. ODBC is constrained to relational data stores;

OLE DB supports all forms of data stores (relational, hierarchical, etc)

4. OLE DB provides Full access to ODBC data sources and ODBC drivers

C:

In short, an ORM is a layer between the server and the database.

The server talks with the ORM and the ORM talks to the database.

The ORM creates objects, that map to the relational data.

It handles your queries, so you don't have to write native SQL, you can query the database with your application language

An object-relational mapper (ORM) is a code library that automates the transfer of data stored in relational database tables into objects that are more commonly used in application code.

Pros

you don't have to learn/know/write SQL, because the ORM handles it

it will be easier to change your database dialect





your application is less vulnerable to SQL injections

ORMs provide a high-level abstraction upon a relational database that allows a developer to write Python code instead of SQL to create, read, update and delete data and schemas in their database. Developers can use the programming language they are comfortable with to work with a database instead of writing SQL statements or stored procedures.

The ability to write Python code instead of SQL can speed up web application development, especially at the beginning of a project

ORMs also make it theoretically possible to switch an application between various relational databases. For example, a developer could use SQLite for local development and MySQL in production. A production application could be switched from MySQL to PostgreSQL with minimal code modifications.

ORMS for python:

web framework	None	Flask	Flask	Django
ORM	SQLAlchemy	SQLAlchemy	SQLAlchemy	Django ORM
database connector	(built into Python stdlib)	MySQL-python	psycopg2	psycopg2
relational database	 SQLite	 MySQL	 PostgreSQL	 PostgreSQL

List of ORMs:

sequelize: Postgres, MySQL, MariaDB, SQLite, Microsoft SQL Server

TypeORM: Postgres, MySQL, MariaDB, SQLite, Microsoft SQL Server, Oracle, sql.js, CockroachDB

objection: Postgres, MySQL, MariaDB, SQLite, Microsoft SQL Server, Oracle, Amazon Redshift

سوال 8

A:

Rowstore

Page #1

ProductID	ProductName	Price	...
1	Product A	20	...
2	Product B	30	...
3	Product C	15	...

Page #2

ProductID	ProductName	Price	...
4	Product D	50	...
5	Product E	20	...
6	Product F	40	...

Columnstore

Page #1

ProductID
1
2
3
4
5
6

Page #2

ProductName
Product A
Product B
Product C
Product D
Product E
Product F

Page #3

Price
20
30
15
50
20
40

Page ...

...
...
...
...
...
...
...

Row oriented databases are databases that organize data **by record**, keeping all of the data associated with a record next to each other in memory. Row oriented databases are the traditional way of organizing data and still provide some key benefits for **storing data quickly**. They are optimized for reading and writing rows efficiently.

Common row oriented databases:

Postgres , MySQL

Column oriented databases are databases that organize data **by field**, keeping all of the data associated with a field next to each other in memory. Columnar databases have grown in popularity and provide performance advantages **to querying data**. They are optimized for **reading and computing** on columns efficiently.

Common column oriented databases:

Redshift , BigQuery,Snowflake

Row Oriented Databases

They are optimized **to read and write a single row** of data which lead to a series of design choices including having a row store architecture.

In a row store, or row oriented database, the data is stored row by row, such that the first column of a row will be next to the last column of the previous row.

For instance

Facebook_Friends

Name	City	Age
Matt	Los Angeles	27
Dave	San Francisco	30
Tim	Oakland	33

This data would be stored on a disk in a row oriented database in order row by row like this:

Matt	Los Angeles	27	Dave	San Francisco	30	Tim	Oakland	33
------	-------------	----	------	---------------	----	-----	---------	----

This allows the database **write a row quickly** because, all that needs to be done to write to it is to tack on another row to the end of the data.

Writing to Row Store Databases

Let's use the data stored in a database:

Matt	Los Angeles	27	Dave	San Francisco	30	Tim	Oakland	33
------	-------------	----	------	---------------	----	-----	---------	----

If we want to add a new record:

Jen	Vancouver	30
-----	-----------	----

We can just append it to the end of the current data:

Matt	Los Angeles	27	Dave	San Francisco	30	Tim	Oakland	33	Jen	Vancouver	30
------	-------------	----	------	---------------	----	-----	---------	----	-----	-----------	----

Reading from Row Store Databases

Say we want to get the **sum of ages** from the Facebook_Friends data. To do this we will need to load all nine of these pieces of data into memory to then pull out the relevant data to do the aggregation.

Matt	Los Angeles	27	Dave	San Francisco	30	Tim	Oakland	33
------	-------------	----	------	---------------	----	-----	---------	----

This is wasted computing time.

Let's assume a Disk can only hold enough bytes of data for three columns to be stored on each disk. In a row oriented database the table above would be stored as:

Disk 1		
Name	City	Age
Matt	Los Angeles	27

Disk 2		
Name	City	Age
Dave	San Francisco	30

Disk 3		
Name	City	Age
Tim	Oakland	33

To get **the sum of all the people's ages** the computer would need to look through all three disks and across all three columns in each disk in order to make this query.

So we can see that **while adding data to a row oriented database is quick** and easy, **getting data out of it can require extra memory** to be used and multiple disks to be accessed.

Row oriented databases are fast at retrieving a row or a set of rows but when performing an **aggregation** it brings extra data (columns) into memory which **is slower** than only selecting the columns that you are performing the aggregation on. In addition the **number of disks** the row oriented database might need to access is usually **larger**.

while adding data to a row oriented database is quick and easy, getting data out of it can require extra memory to be used and multiple disks to be accessed.

Column Oriented Databases

Data Warehouses were created in order to support analyzing data. These types of databases are **read optimized**.

the data is stored such that each row of a column will be next to other rows from that same column.

Facebook_Friends

Name	City	Age
Matt	Los Angeles	27
Dave	San Francisco	30
Tim	Oakland	33

A table is stored **one column at a time** in order row by row:

Matt	Dave	Tim	Los Angeles	San Francisco	Oakland	27	30	33
------	------	-----	-------------	---------------	---------	----	----	----

Writing to a Column Store Databases

If we want to add a new record:

Jen	Vancouver	30
-----	-----------	----

We have to navigate around the data to plug each column in to where it should be.

Matt	Dave	Tim	Jen	Los Angeles	San Francisco	Oakland	Vancouver	27	30	33	30
------	------	-----	-----	-------------	---------------	---------	-----------	----	----	----	----

If the data was stored on a single disk it would have the same extra memory problem as a row oriented database, since it would need to bring everything into memory. However, column oriented databases will have significant benefits when stored on separate disks.

If we placed the table above into the similarly restricted three columns of data disk they would be stored like this:

Disk 1		
Name		
Matt	Dave	Tim

Disk 2		
City		
Los Angeles	San Francisco	Oakland

Disk 3		
Age		
27	30	33

Reading from a Column store Database

To get the sum of the ages the computer only needs to go to one disk (Disk 3) and sum all the values inside of it. **No extra memory needs** to be pulled in, and it accesses a **minimal number of disks**.

by organizing data by column the number of disks that will need to be visited will be reduced and the amount of extra data that has to be held in memory is minimized. This greatly **increases the overall speed of the computation**.

B:

Row based tables have disadvantages in case of analytic applications where aggregation are used and faster search & processing are required. In row based tables all data in a row has to be read even though the requirement may be there to access data from a few columns.

These are not efficient in performing operations applicable to the entire datasets and hence aggregation in row-oriented is an expensive job or operations.

Typical compression mechanisms which provide less efficient result than what we achieve from column-oriented data stores.

C:

In case of analytic applications, where aggregations are used and faster search & processing are required, row-based storage are not good. In row based tables all data stored in a row has to be read even though the requirement may be there to access data from a few columns. Hence, these queries on huge amounts of data would take lots of times

These are not efficient in performing operations applicable to the entire datasets and hence aggregation in row-oriented is an expensive job or operations.

چون می‌خواهیم فروش را مقایسه کنیم و آنالیز کنیم طبق متن بالا از column based استفاده می‌کنیم. چون سرعت بیشتری در داده‌هایی با حجم بالا دارد.

D:

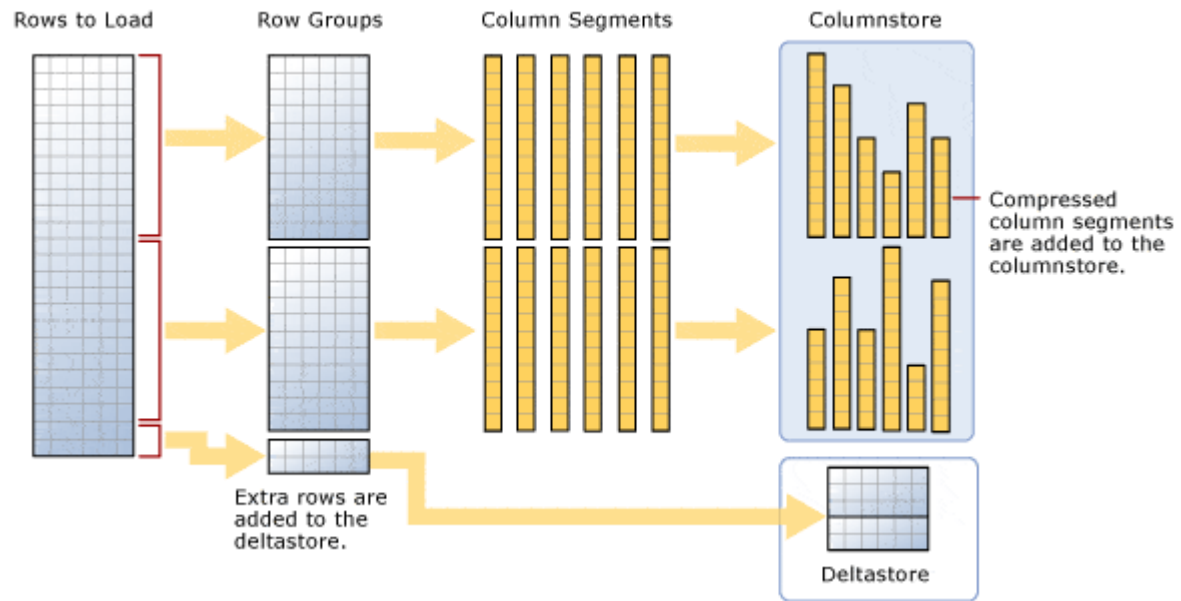
طبق توضیحات بالا و مثال زده شده:

Insert, update در روش rowstore بهینه‌تر هستند چون فقط نیاز است داده‌های جدید را به انتهای داده‌های قبلی اضافه کنیم و برای عملیات‌های transaction مناسب‌تر است.

E:

Columnstore simply means a new way to store the data in the index. Instead of the normal Rowstore or b-tree indexes where the data is logically and physically organized and stored as a table with rows and columns, the data in columnstore indexes are physically stored in columns and logically organized in rows and columns. Instead of storing an entire row or rows in a page, one column from many rows is stored in that page. It is this difference in architecture that gives the columnstore index a very high level of compression along with reducing the storage footprint and providing massive improvements in read performance.

The index works by slicing the data into compressible segments. It takes a group of rows, a minimum of 102,400 rows with a max of about 1 million rows, called a rowgroup and then changes that group of rows into Column segments. It's these segments that are the basic unit of storage for a columnstore index, as shown below



Columnstore indexes are designed for large data warehouse workloads, not normal OLTP workload tables.