

# Operating Systems

Isfahan University of Technology  
Electrical and Computer Engineering Department  
1400-1 semester

Zeinab Zali

**Session 4: Virtualization and Computing Environments**

# Virtualization

- **Process**: OS abstraction of the **processor**, main memory and I/O devices for a running program
  - Multiple processes can concurrently run, each thinking itself as the exclusive user of the hardware.
- **Virtual Memory**: OS abstraction of **Memory**
  - Each process perceives the same picture of memory used only by itself (its address space).
- **File**: OS abstraction of **I/O devices**
  - All input and output in the system is performed by reading and writing files

# CPU Virtualization

- **virtualizing the CPU:** Turning a single CPU (or small set of them) into a seemingly infinite number of CPUs and thus allowing many programs to seemingly run at once

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

# CPU Virtualization

- **virtualizing the CPU:** Turning a single CPU (or small set of them) into a seemingly infinite number of CPUs and thus allowing many programs to seemingly run at once

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

```
zeinabzali:./cpu A
A
A
A
A
A
```

# CPU Virtualization

- **virtualizing the CPU:** Turning a single CPU (or small set of them) into a seemingly infinite number of CPUs and thus allowing many programs to seemingly run at once

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

```
zeinabzali:./cpu A & ./cpu B & ./cpu C
[1] 27878
[2] 27879
A
B
C
A
B
C
A
B
C
A
B
C
A
B
C
A
B
C
A
```

# Memory Virtualization

- Memory is **just an array of bytes**
  - to read memory, one must specify an address to be able to access the data stored there
  - to write (or update) memory, one must also specify the data to be written to the given address
- A program keeps all of its data structures in memory, and accesses them through various instructions
  - **loads and stores**
- Each instruction of the program is in memory too
  - thus memory is accessed on each instruction fetch

# Memory Virtualization

- **virtualizing the Memory:** Each process accesses its own private virtual address space which the OS somehow maps onto the physical memory of the machine

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: mem <value>\n");
        exit(1);
    }
    int *p;
    p = malloc(sizeof(int));
    assert(p != NULL);
    printf("(%d) addr pointed to by p: %p\n", (int) getpid(), p);
    *p = atoi(argv[1]); // assign value to addr stored in p
    //p = atoi(argv[1]);
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(%d) value of p: %d\n", getpid(), *p);
    }
    return 0;
}
```

# Memory Virtualization

- **virtualizing the Memory:** Each process accesses its own private virtual address space which the OS somehow maps onto the physical memory of the machine

```
setarch $(uname --machine) --addr-no-randomize /bin/bash
```

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: mem <value>\n");
        exit(1);
    }
    int *p;
    p = malloc(sizeof(int));
    assert(p != NULL);
    printf("(pid) addr pointed to by p: %p\n", (int) getpid(), p);
    *p = atoi(argv[1]); // assign value to addr stored in p
    //p = atoi(argv[1]);
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(pid) value of p: %d\n", getpid(), *p);
    }
    return 0;
}
```

```
zeinabzali:./mem 10
(28692) addr pointed to by p: 0x555555756260
(28692) value of p: 11
(28692) value of p: 12
(28692) value of p: 13
(28692) value of p: 14
```



# Memory Virtualization

- **virtualizing the Memory:** Each process accesses its own private virtual address space which the OS somehow maps onto the physical memory of the machine

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: mem <value>\n");
        exit(1);
    }
    int *p;
    p = malloc(sizeof(int));
    assert(p != NULL);
    printf("(%d) addr pointed to by p: %p\n", (int) getpid(), p);
    *p = atoi(argv[1]); // assign value to addr stored in p
    //p = atoi(argv[1]);
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(%d) value of p: %d\n", getpid(), *p);
    }
    return 0;
}
```

```
zeinabzali:./mem 10 & ./mem 100 & ./mem 1000
[1] 28699
[2] 28700
(28699) addr pointed to by p: 0x555555756260
(28700) addr pointed to by p: 0x555555756260
(28701) addr pointed to by p: 0x555555756260
(28699) value of p: 11
(28700) value of p: 101
(28701) value of p: 1001
(28699) value of p: 12
(28700) value of p: 102
(28701) value of p: 1002
(28699) value of p: 13
(28700) value of p: 103
(28701) value of p: 1003
(28699) value of p: 14
(28700) value of p: 104
(28701) value of p: 1004
(28699) value of p: 15
(28700) value of p: 105
(28701) value of p: 1005
```

# Concurrency

- problems that arise, and must be addressed, when working on many things at once (i.e., concurrently)

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
#include "common_threads.h"

volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        counter++;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: threads <loops>\n");
        exit(1);
    }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);
    pthread_create(&p1, NULL, worker, NULL);
    pthread_create(&p2, NULL, worker, NULL);
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("Final value  : %d\n", counter);
    return 0;
}
```

# Concurrency

- problems that arise, and must be addressed, when working on many things at once (i.e., concurrently)

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
#include "common_threads.h"

volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        counter++;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: threads <loops>\n");
        exit(1);
    }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);
    pthread_create(&p1, NULL, worker, NULL);
    pthread_create(&p2, NULL, worker, NULL);
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("Final value : %d\n", counter);
    return 0;
}
```

```
zeinabzali:./threads 10
Initial value : 0
Final value : 20
zeinabzali:./threads 100
Initial value : 0
Final value : 200
zeinabzali:./threads 1000
Initial value : 0
Final value : 2000
zeinabzali:10000
10000: command not found
zeinabzali:./threads 10000
Initial value : 0
Final value : 14991
zeinabzali:█
```



# Virtualization

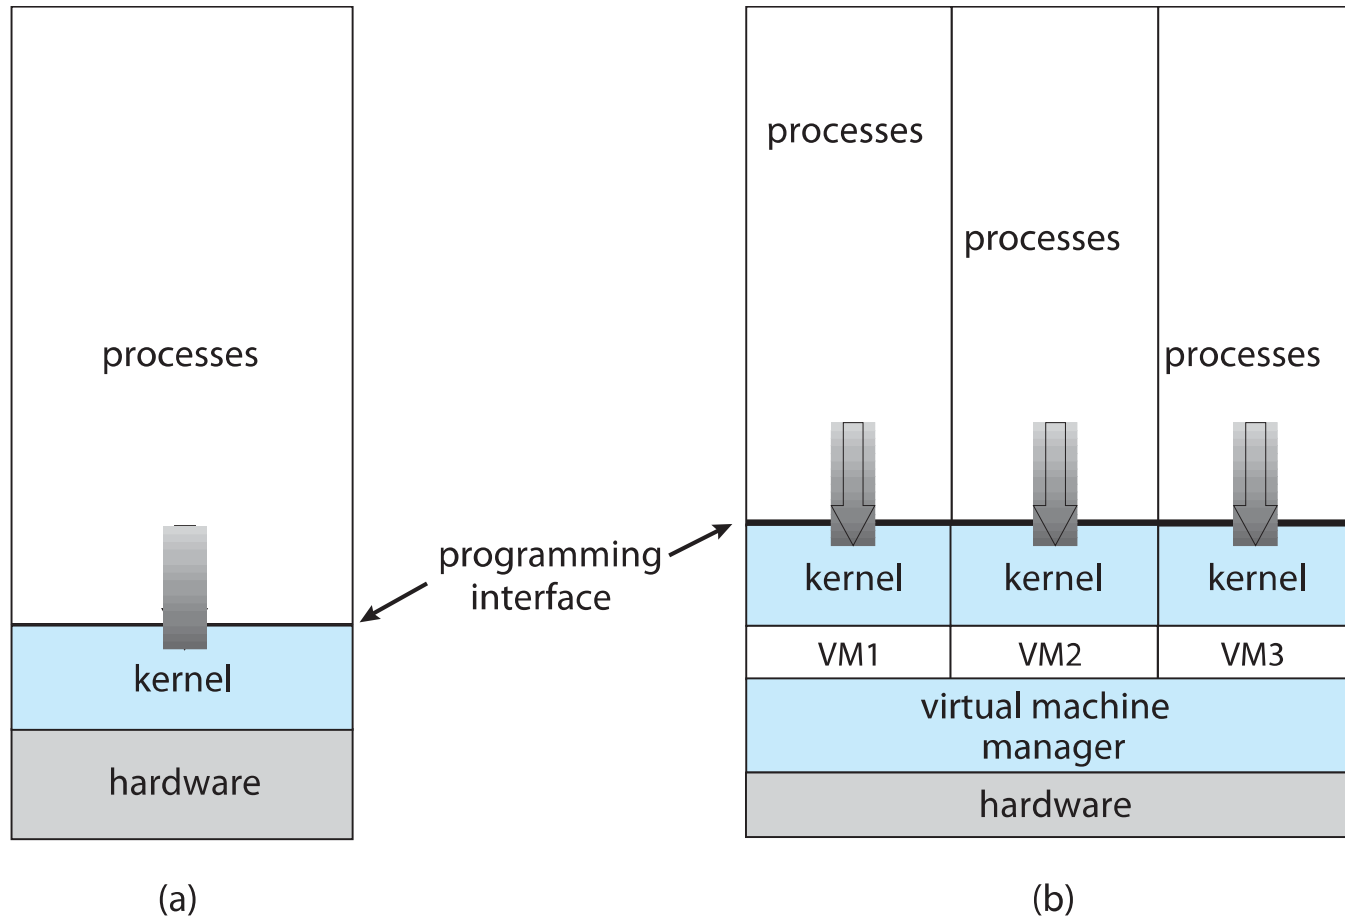
---

- **Virtualization** is a technology that allows us to **abstract** the hardware of a single computer (the CPU , memory, disk drives, network interface cards, and so forth) into **several different execution environments**, thereby creating the illusion that each separate environment is running on its own private computer.
- A user of a **virtual machine** can switch among the various operating systems in the same way a user can switch among the various processes running concurrently in a single operating system.
- **Emulation** is simulating computer hardware in software.
- Broadly speaking, virtualization software is one member of a class that also includes emulation.
  - Emulation is typically used when the source CPU type is different from the target CPU type.

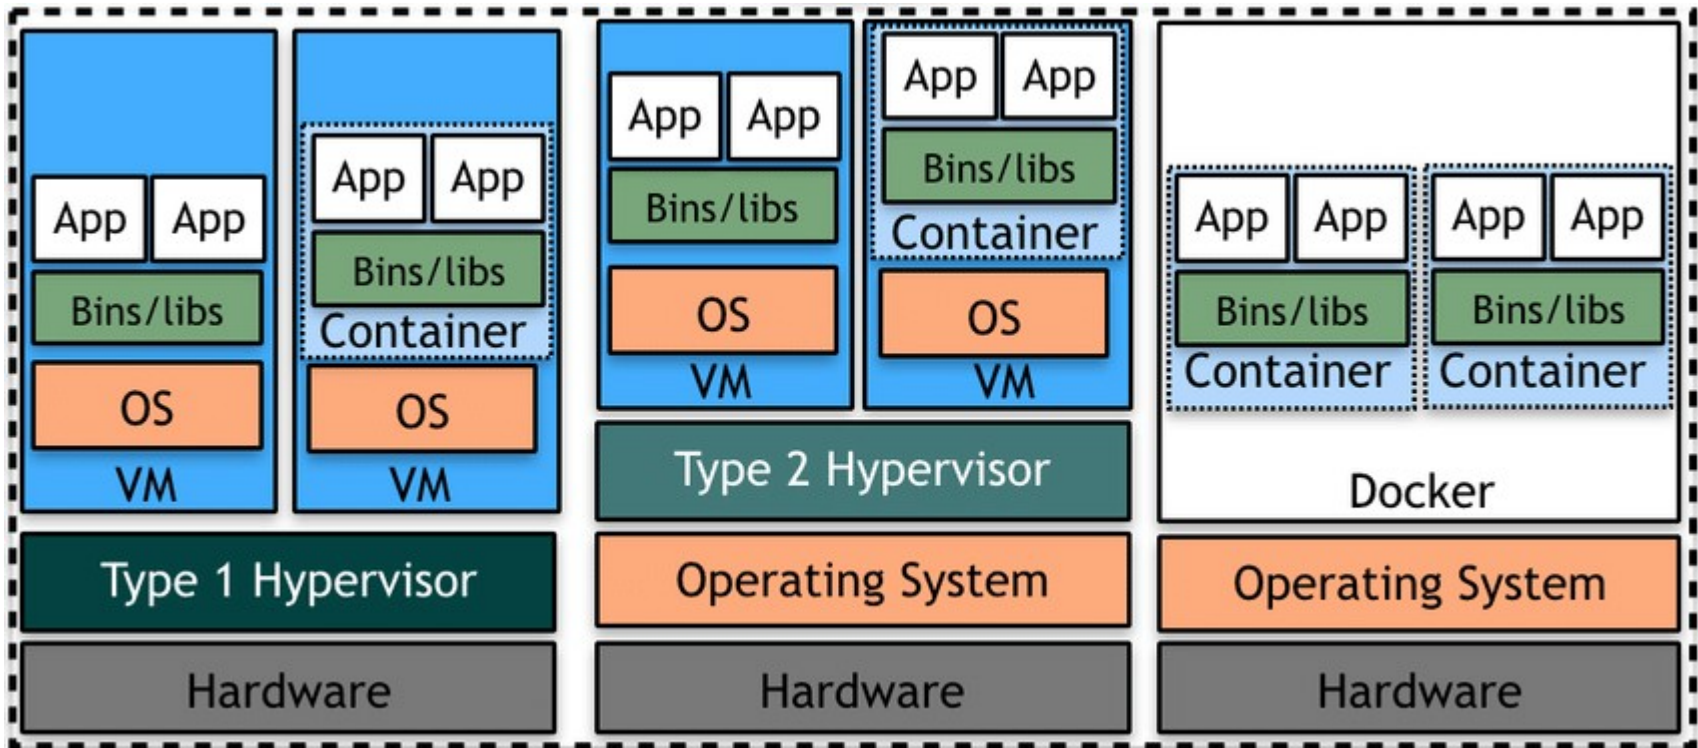




# Computing Environments - Virtualization



# Container-based Virtualization



# Computing Environments





# Distributed Systems

- Collection of separate, possibly heterogeneous, systems networked together
  - **Network** is a communications path, **TCP/IP** most common
    - ▶ **Local Area Network (LAN)**
    - ▶ **Wide Area Network (WAN)**
    - ▶ **Metropolitan Area Network (MAN)**
    - ▶ **Personal Area Network (PAN)**
- **Network Operating System** provides features between systems across network
  - Communication scheme allows systems to exchange messages
  - Illusion of a single system








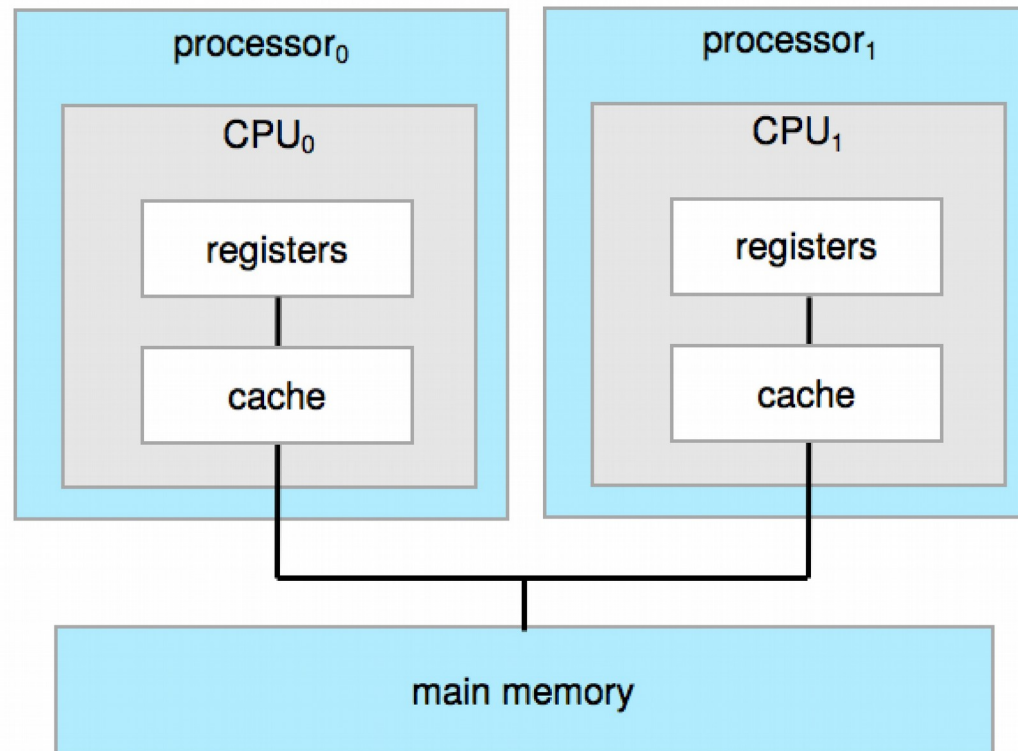
# Computer-System Architecture

- **Multiprocessors** systems growing in use and importance
  - Also known as **parallel systems**, **tightly-coupled systems**
  - Advantages include:
    1. **Increased throughput**
    2. **Economy of scale**
    3. **Increased reliability** – graceful degradation or fault tolerance
  - Two types:
    1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
    2. **Symmetric Multiprocessing** – each processor performs all tasks






# Symmetric Multiprocessing Architecture





# Dual-Core Design

- Multi-chip and **multicore**
- Systems containing all chips
  - Chassis containing multiple separate systems






# Clustered Systems

- Like multiprocessor systems, but multiple systems working together
  - Usually sharing storage via a **storage-area network (SAN)**
  - Provides a **high-availability** service which survives failures
    - ▶ **Asymmetric clustering** has one machine in hot-standby mode
    - ▶ **Symmetric clustering** has multiple nodes running applications, monitoring each other
  - Some clusters are for **high-performance computing (HPC)**
    - ▶ Applications must be written to use **parallelization**
  - Some have **distributed lock manager (DLM)** to avoid conflicting operations






# Clustered Systems





# Client Server


- Client-Server Computing
  - Dumb terminals supplanted by smart PCs
  - Many systems now **servers**, responding to requests generated by **clients**
    - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
    - ▶ **File-server system** provides interface for clients to store and retrieve files






# Peer-to-Peer

- Another model of distributed system
- **P2P does not distinguish clients and servers**
  - Instead all nodes are considered peers
  - May each act as client, server or both
  - Node must join P2P network
    - ▶ Registers its service with central lookup service on network, or
    - ▶ Broadcast request for service and respond to requests for service via **discovery protocol**
  - Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype





# Cloud Computing

- Delivers computing, storage, even apps **as a service** across a network
- Logical extension of **virtualization** because it uses virtualization as the base for its functionality.
  - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage










# Cloud Computing types



- **Public cloud** – available via Internet to anyone willing to pay
- **Private cloud** – run by a company for the company's own use
- **Hybrid cloud** – includes both public and private cloud components
- Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
- Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
- Infrastructure as a Service (**IaaS**) – servers or storage available over Internet





# Cloud Computing Environments

- Cloud computing environments composed of traditional OSeS, plus VMMs, plus cloud management tools
  - Internet connectivity requires security like firewalls
  - Load balancers spread traffic across multiple applications





# Real-Time Embedded Systems

---

- Real-time embedded systems most prevalent form of computers
  - Vary considerable, special purpose, limited purpose OS, **real-time OS**
- Real-time OS has well-defined fixed time constraints
  - Processing **must** be done within constraint
  - Correct operation only if constraints met
- Examples of RTOS (Real-Time OS)
  - VxWorks
  - LynxOS

