


# Operating Systems

Isfahan University of Technology  
Electrical and Computer Engineering Department  
1400-1 semester

Zeinab Zali

Session 8: Implicit Threading, Signals, ...




# Implicit Threading

---

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- some methods explored
  - Thread Pools
  - OpenMP
  - Fork-Join
  - Grand Central Dispatch
  - Intel Threading Building Blocks





# Client-Server Example


---

Remember our own example

What happens if the number of alive threads exceeds the maximum number of concurrent threads that the system can support?

How can we prevent this issue?






# Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool
  - Separating task to be performed from mechanics of creating task allows different strategies for running task
    - ▶ i.e. Tasks could be scheduled to run periodically
- Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```




# Sample thread pool API

```
void first_task() {...}  
void second_task() {...}  
void third_task() {...}  
main(){  
    // Create a thread pool.  
    pool tp(2);  
    //Add some tasks to the pool.  
    tp.schedule(&first_task);  
    tp.schedule(&second_task);  
    tp.schedule(&third_task);  
}
```



# Fork-Join Parallelism

- Multiple threads (tasks) are **forked**, and then **joined**.





# Fork-Join Parallelism


---

- General algorithm for fork-join strategy:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))


    result1 = join(subtask1)
    result2 = join(subtask2)

    return combined results
```





# Fork-Join Parallelism







# OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN provides support for parallel programming
- Identifies **parallel regions** – blocks of code that can run in parallel


```
#pragma omp parallel
```

Create as many threads as there are cores

```
#pragma omp parallel for  
for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];  
}
```

Run for loop in parallel






# Threading Issues

---

- Semantics of **fork()** and **exec()** system calls
- Signal handling
  - Synchronous and asynchronous
- Thread cancellation of target thread
  - Asynchronous or deferred






# Semantics of `fork()` and `exec()`

---

- Does `fork()` duplicate only the calling thread or all threads?
  - Some UNIX systems have two versions of `fork()`:
    - one that duplicates all threads
    - one that duplicates only the thread that invoked the `fork()` system call.
- `exec()` usually works as normal – replace the running process including all threads
  - So when `exec()` is called immediately after forking, forking only the calling thread is sufficient






# Signal Handling

- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
  1. Signal is generated by particular event
  2. Signal is delivered to a process
  3. Signal is handled by one of two signal handlers:
    1. default
    2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
  - **User-defined signal handler** can override default
  - For single-threaded, signal delivered to process





# Signal Handling (Cont.)


---

- Where should a signal be delivered for multi-threaded?
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the processes

**kill(pid\_t, signal)**

**pthread\_kill(thread\_t, signal)**






# Thread Cancellation

- Terminating a thread before it has finished
  - Suppose multiple threads searching for a record in a database and one of them find it
  - a web page loads using several threads—each image is loaded in a separate thread. When a user presses the stop button on the browser, all threads loading the page are canceled.
- Thread to be canceled is **target thread**
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check if it should be canceled





# Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
  - Cancellation only occurs when thread reaches **cancellation point**
    - ▶ Ex: `pthread_testcancel()`
    - ▶ Ex: `read` function
- On Linux systems, thread cancellation is handled through signals






# Thread-Local Storage

---

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
  - Local variables visible only during single function invocation
  - TLS visible across function invocations
- Similar to `static` data
  - TLS is unique to each thread








# Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
  - Appears to be a virtual processor on which process can schedule user thread to run
  - Each LWP attached to kernel thread
  - How many LWPs to create?
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- This communication allows an application to maintain the correct number kernel threads





# Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
  - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

getconf GNU\_LIBPTHREAD\_VERSION



# End of Chapter 4

---

