

# Operating Systems

Isfahan University of Technology  
Electrical and Computer Engineering Department  
1400-1 semester

Zeinab Zali

Session 15: Synchronization-Condition Variables-monitors

# Condition Variable

- there are many cases where a thread wishes to check whether a condition is true before continuing its execution
- To wait for a condition to become true, a thread can make use of what is known as a condition variable
- A condition variable is an **explicit queue** that threads can put themselves on when some state of execution (i.e., some condition) is not as desired (by waiting on the condition)
- some other thread, when it changes said state, can then **wake one (or more)** of those waiting threads and thus allow them to continue (by signaling on the condition)

# CV definitions and routines

- `pthread_cond_t c`
- `pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m):`  
The `wait()` call is executed when a thread wishes to put itself to sleep
  - The responsibility of `wait()` is to release the lock and put the calling thread to sleep (atomically); when the thread wakes up (after some other thread has signaled it), it must re-acquire the lock before returning to the caller.
- `pthread_cond_signal(pthread_cond_t *c):` the `signal()` call is executed when a thread has changed something in the program and thus wants to wake a sleeping thread waiting on this condition

# CV example 1

```
1 void *child(void *arg) {
2     printf("child\n");
3     // XXX how to indicate we are done?
4     return NULL;
5 }
6
7 int main(int argc, char *argv[]) {
8     printf("parent: begin\n");
9     pthread_t c;
10    Pthread_create(&c, NULL, child, NULL); // create child
11    // XXX how to wait for child?
12    printf("parent: end\n");
13    return 0;
14 }
```

**Desired Output:**

```
parent: begin
child
parent: end
```

```
1  int done  = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c  = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6
7
8
9
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14     thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19
20
21
22
23 }
24
25 int main(int argc, char *argv[]) {
26     printf("parent: begin\n");
27     pthread_t p;
28     Pthread_create(&p, NULL, child, NULL);
29     thr_join();
30     printf("parent: end\n");
31     return 0;
32 }
```

# What is the problem with this solution?

```
1  void thr_exit() {
2      Pthread_mutex_lock(&m);
3      Pthread_cond_signal(&c);
4      Pthread_mutex_unlock(&m);
5  }
6
7  void thr_join() {
8      Pthread_mutex_lock(&m);
9      Pthread_cond_wait(&c, &m);
10     Pthread_mutex_unlock(&m);
11 }
```

# What is the problem with this solution?

```
1  void thr_exit() {
2      done = 1;
3      Pthread_cond_signal(&c);
4  }
5
6  void thr_join() {
7      if (done == 0)
8          Pthread_cond_wait(&c);
9  }
```

# Producer/Consumer Problem

```
1  int loops; // must initialize somewhere...
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8
9
10
11
12
13
14      }
15  }
16
17  void *consumer(void *arg) {
18      int i;
19      for (i = 0; i < loops; i++) {
20
21
22
23
24
25
26
27      }
28  }
```

// p1  
// p2  
// p3  
// p4  
// p5  
// p6  
  
// c1  
// c2  
// c3  
// c4  
// c5  
// c6



# Producer/Consumer Problem

```
1  int buffer;  
2  int count = 0; // initially, empty  
3  
4  void put(int value) {  
5      assert(count == 0);  
6      count = 1;  
7      buffer = value;  
8  }  
9  
10 int get() {  
11     assert(count == 1);  
12     count = 0;  
13     return buffer;  
14 }
```

# Producer/Consumer Problem

```
1  int loops; // must initialize somewhere...
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);           // p1
9          if (count == 1)                       // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                             // p4
12             Pthread_cond_signal(&cond);        // p5
13             Pthread_mutex_unlock(&mutex);      // p6
14         }
15     }
16
17     void *consumer(void *arg) {
18         int i;
19         for (i = 0; i < loops; i++) {
20             Pthread_mutex_lock(&mutex);         // c1
21             if (count == 0)                     // c2
22                 Pthread_cond_wait(&cond, &mutex); // c3
23             int tmp = get();                    // c4
24             Pthread_cond_signal(&cond);         // c5
25             Pthread_mutex_unlock(&mutex);       // c6
26             printf("%d\n", tmp);
27         }
28     }
```

# Producer/Consumer broken solution with if

| $T_{c1}$ | State   | $T_{c2}$ | State   | $T_p$ | State   | Count | Comment                |
|----------|---------|----------|---------|-------|---------|-------|------------------------|
| c1       | Running |          | Ready   |       | Ready   | 0     |                        |
| c2       | Running |          | Ready   |       | Ready   | 0     |                        |
| c3       | Sleep   |          | Ready   |       | Ready   | 0     | Nothing to get         |
|          | Sleep   |          | Ready   | p1    | Running | 0     |                        |
|          | Sleep   |          | Ready   | p2    | Running | 0     |                        |
|          | Sleep   |          | Ready   | p4    | Running | 1     | Buffer now full        |
|          | Ready   |          | Ready   | p5    | Running | 1     | $T_{c1}$ awoken        |
|          | Ready   |          | Ready   | p6    | Running | 1     |                        |
|          | Ready   |          | Ready   | p1    | Running | 1     |                        |
|          | Ready   |          | Ready   | p2    | Running | 1     |                        |
|          | Ready   |          | Ready   | p3    | Sleep   | 1     | Buffer full; sleep     |
|          | Ready   | c1       | Running |       | Sleep   | 1     | $T_{c2}$ sneaks in ... |
|          | Ready   | c2       | Running |       | Sleep   | 1     |                        |
|          | Ready   | c4       | Running |       | Sleep   | 0     | ... and grabs data     |
|          | Ready   | c5       | Running |       | Ready   | 0     | $T_p$ awoken           |
|          | Ready   | c6       | Running |       | Ready   | 0     |                        |
| c4       | Running |          | Ready   |       | Ready   | 0     | Oh oh! No data         |

```

1  int loops;
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);          // p1
9          while (count == 1)                  // p2
10             Pthread_cond_wait(&cond, &mutex); // p3
11             put(i);                          // p4
12             Pthread_cond_signal(&cond);      // p5
13             Pthread_mutex_unlock(&mutex);    // p6
14         }
15     }
16
17     void *consumer(void *arg) {
18         int i;
19         for (i = 0; i < loops; i++) {
20             Pthread_mutex_lock(&mutex);      // c1
21             while (count == 0)                // c2
22                 Pthread_cond_wait(&cond, &mutex); // c3
23             int tmp = get();                  // c4
24             Pthread_cond_signal(&cond);      // c5
25             Pthread_mutex_unlock(&mutex);    // c6
26             printf("%d\n", tmp);
27         }
28     }

```

```

1  int loops;
2  cond_t  cond;
3  mutex_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          Pthread_mutex_lock(&mutex);          // p1
9          while (count == 1)                    // p2
10             Pthread_cond_wait(&cond, &mutex); // p3

```

### TIP: USE WHILE (NOT IF) FOR CONDITIONS

When checking for a condition in a multi-threaded program, using a `while` loop is always correct; using an `if` statement only might be, depending on the semantics of signaling. Thus, always use `while` and your code will behave as expected.

```

18      int i;
19      for (i = 0; i < loops; i++) {
20          Pthread_mutex_lock(&mutex);          // c1
21          while (count == 0)                    // c2
22             Pthread_cond_wait(&cond, &mutex); // c3
23          int tmp = get();                      // c4
24          Pthread_cond_signal(&cond);          // c5
25          Pthread_mutex_unlock(&mutex);        // c6
26          printf("%d\n", tmp);
27      }
28  }

```

# Producer/Consumer still a broken solution

| $T_{c1}$ | State   | $T_{c2}$ | State   | $T_p$ | State   | Count | Comment             |
|----------|---------|----------|---------|-------|---------|-------|---------------------|
| c1       | Running |          | Ready   |       | Ready   | 0     |                     |
| c2       | Running |          | Ready   |       | Ready   | 0     |                     |
| c3       | Sleep   |          | Ready   |       | Ready   | 0     | Nothing to get      |
|          | Sleep   | c1       | Running |       | Ready   | 0     |                     |
|          | Sleep   | c2       | Running |       | Ready   | 0     |                     |
|          | Sleep   | c3       | Sleep   |       | Ready   | 0     | Nothing to get      |
|          | Sleep   |          | Sleep   | p1    | Running | 0     |                     |
|          | Sleep   |          | Sleep   | p2    | Running | 0     |                     |
|          | Sleep   |          | Sleep   | p4    | Running | 1     | Buffer now full     |
|          | Ready   |          | Sleep   | p5    | Running | 1     | $T_{c1}$ awoken     |
|          | Ready   |          | Sleep   | p6    | Running | 1     |                     |
|          | Ready   |          | Sleep   | p1    | Running | 1     |                     |
|          | Ready   |          | Sleep   | p2    | Running | 1     |                     |
|          | Ready   |          | Sleep   | p3    | Sleep   | 1     | Must sleep (full)   |
| c2       | Running |          | Sleep   |       | Sleep   | 1     | Recheck condition   |
| c4       | Running |          | Sleep   |       | Sleep   | 0     | $T_{c1}$ grabs data |
| c5       | Running |          | Ready   |       | Sleep   | 0     | Oops! Woke $T_{c2}$ |
| c6       | Running |          | Ready   |       | Sleep   | 0     |                     |
| c1       | Running |          | Ready   |       | Sleep   | 0     |                     |
| c2       | Running |          | Ready   |       | Sleep   | 0     |                     |
| c3       | Sleep   |          | Ready   |       | Sleep   | 0     | Nothing to get      |
|          | Sleep   | c2       | Running |       | Sleep   | 0     |                     |
|          | Sleep   | c3       | Sleep   |       | Sleep   | 0     | Everyone asleep...  |

# Producer/Consumer with buffer size MAX

```
1  cond_t  empty, fill;
2  mutex_t mutex;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          Pthread_mutex_lock(&mutex);
8          while (count == 1)
9              Pthread_cond_wait(&empty, &mutex);
10         put(i);
11         Pthread_cond_signal(&fill);
12         Pthread_mutex_unlock(&mutex);
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         Pthread_mutex_lock(&mutex);
20         while (count == 0)
21             Pthread_cond_wait(&fill, &mutex);
22         int tmp = get();
23         Pthread_cond_signal(&empty);
24         Pthread_mutex_unlock(&mutex);
25         printf("%d\n", tmp);
26     }
27 }
```

# Producer/Consumer with buffer size MAX

```
1  int buffer[MAX];
2  int fill_ptr = 0;
3  int use_ptr  = 0;
4  int count    = 0;
5
6  void put(int value) {
7      buffer[fill_ptr] = value;
8      fill_ptr = (fill_ptr + 1) % MAX;
9      count++;
10 }
11
12 int get() {
13     int tmp = buffer[use_ptr];
14     use_ptr = (use_ptr + 1) % MAX;
15     count--;
16     return tmp;
17 }
```





# Monitors

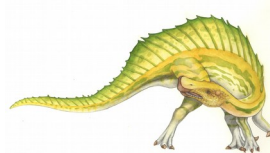
- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- Pseudocode syntax of a monitor:

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure P2 (...) { ... }


    procedure Pn (...) {.....}

    initialization code (...) { ... }
}
```





# Schematic view of a Monitor





# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex  
mutex = 1
```

- Each procedure ***P*** is replaced by


```
wait(mutex) ;  
    ...  
    body of P ;  
    ...  
signal(mutex) ;
```

- Mutual exclusion within a monitor is ensured





# Monitor with Condition Variables





# Usage of Condition Variable Example

- Consider  $P_1$  and  $P_2$  that need to execute two statements  $S_1$  and  $S_2$  and the requirement that  $S_1$  to happen before  $S_2$ 
  - Create a monitor with two procedures  $F_1$  and  $F_2$  that are invoked by  $P_1$  and  $P_2$  respectively
  - One condition variable “x” initialized to 0
  - One Boolean variable “done”
  - **F1 :**  
 $S_1;$   
`done = true;`  
`x.signal();`
  - **F2 :**  
`if done = false`  
`x.wait();`  
 $S_2;$






# Liveness

---

- Processes may have to wait indefinitely while trying to acquire a synchronization tool such as a mutex lock or semaphore.
- Waiting indefinitely violates the progress and bounded-waiting criteria discussed at the beginning of this chapter.
- **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress.
- Indefinite waiting is an example of a liveness failure.





# Liveness

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$

`wait(S) ;`

`wait(Q) ;`

`...`

`signal(S) ;`

`signal(Q) ;`

$P_1$

`wait(Q) ;`

`wait(S) ;`

`...`

`signal(Q) ;`

`signal(S) ;`

- Consider if  $P_0$  executes `wait(S)` and  $P_1$  `wait(Q)`. When  $P_0$  executes `wait(Q)`, it must wait until  $P_1$  executes `signal(Q)`
- However,  $P_1$  is waiting until  $P_0$  execute `signal(S)`.
- Since these `signal()` operations will never be executed,  $P_0$  and  $P_1$  are **deadlocked**.






# Liveness

---

- Other forms of deadlock:
- **Starvation** – indefinite blocking
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Solved via **priority-inheritance protocol**







# End of Chapter 6

